

independIT Integrative Technologies GmbH
Bergstraße 6
D-86529 Schrobenhausen



schedulix!Web

User Guide Release 2.7

Dieter Stubler Ronald Jeninga

March 30, 2016

Copyright © 2016 independIT GmbH

Legal notice

This work is copyright protected

Copyright © 2016 independIT Integrative Technologies GmbH

All rights reserved. No part of this work may be reproduced or transmitted
in any form or by any means, electronically or mechanically, including photo-
copying, recording, or by any information storage or retrieval system, without
the prior written permission of the copyright owner.

Contents

Contents	i
List of tables	iii
List of illustrations	v
1 General	1
1.1 Overview	1
1.2 Login	2
1.3 Main Desktop	2
1.4 Layout of the object window	4
1.4.1 Title bar	4
1.4.2 Navigator	6
1.4.2.1 Standard buttons	7
1.4.3 Editor	9
1.4.3.1 Standard buttons	10
1.5 Selecting values	14
1.5.1 Selecting values from drop-down lists	14
1.5.2 Choosing values using the Select button	15
1.6 Handling standard lists	15
1.6.1 Adding a line	17
1.6.2 Deleting a line	18
1.6.3 Changing line values	19
1.7 Copy & paste and the clipboard	19
1.7.1 Moving objects in the hierarchy	20
1.7.2 Linking objects	23
1.8 Graphic Legend	25
1.8.1 How schedulix objects are represented	25
1.8.2 Representation of parent-child relationships in workflows	25
1.8.3 Representation of dependencies between Scheduling Entities	26
1.9 Comments	26
1.10 Standard fields	27
2 Exit State Definitions	29
2.1 View	29
2.2 Concept	29
2.2.1 In short	29

2.2.2	Detailed description	29
2.3	Editor	29
3	Exit State Mappings	31
3.1	View	31
3.2	Concept	31
3.2.1	In short	31
3.2.2	Detailed description	31
3.3	Editor	32
4	Exit State Profiles	35
4.1	View	35
4.2	Concept	35
4.2.1	In short	35
4.2.2	Detailed description	35
4.3	Editor	36
5	Resource State Definitions	39
5.1	View	39
5.2	Concept	39
5.2.1	In short	39
5.2.2	Detailed description	39
5.3	Editor	39
6	Resource State Profiles	41
6.1	View	41
6.2	Concept	41
6.2.1	In short	41
6.2.2	Detailed description	41
6.3	Editor	41
6.3.1	Example	42
7	Resource State Mappings	45
7.1	View	45
7.2	Concept	45
7.2.1	In short	45
7.2.2	Detailed description	45
7.3	Editor	46
8	Named Resources	49
8.1	View	49
8.2	Concept	49
8.2.1	In short	49
8.2.2	Detailed description	49

8.3	Editor	50
8.3.1	Properties tab for categories	50
8.3.2	Content tab	51
8.3.3	Properties tab for Named Resource Definitions	51
8.3.4	Parameters tab	52
8.3.4.1	Parameter Details tab	53
8.3.4.2	Standard parameters	54
8.3.5	Resources tab	54
8.3.6	Job Definitions tab	55
8.4	Named Resource selector	56
9	Environments	59
9.1	View	59
9.2	Concept	59
9.2.1	In short	59
9.2.2	Detailed description	59
9.3	Navigator	60
9.4	Editor	60
9.4.1	Properties tab	60
9.4.2	References tab	61
10	Footprints	63
10.1	View	63
10.2	Concept	63
10.2.1	In short	63
10.2.2	Detailed description	63
10.3	Editor	64
10.3.1	Properties tab	64
10.3.2	References tab	65
11	Job servers and resources	67
11.1	View	67
11.2	Concept	67
11.2.1	In short	67
11.2.2	Detailed description	67
11.3	Navigator	68
11.4	Editor	68
11.4.1	Properties tab	68
11.4.2	Resources tab	71
11.4.2.1	Resource Details tab	72
11.4.2.2	Parameters tab	74
11.4.2.3	Allocations tab	75
11.4.3	Parameters tab	79

11.4.4	Config tab	80
11.4.4.1	Standard configuration parameters	81
11.4.5	Env.Map tab	82
11.4.6	LogFile Pattern tab	83
12	Batches and Jobs	85
12.1	View	85
12.2	Concept	85
12.2.1	In short	85
12.2.2	Detailed description	86
12.2.3	Recommended convention for batch objects	86
12.3	Navigator	87
12.3.1	Pinning	87
12.3.2	Folder bookmarks	88
12.3.3	Folder search	89
12.4	Folder editor	90
12.4.1	Properties tab	90
12.4.2	Content tab	91
12.4.3	Parameters tab	91
12.4.4	Resources tab	92
12.5	Editor for Job Definitions	93
12.5.1	Properties tab	93
12.5.2	Run tab	98
12.5.3	Restart tab	102
12.5.4	Children tab	103
12.5.4.1	Child Details tab	103
12.5.5	Parents tab	108
12.5.6	Dependencies tab	109
12.5.6.1	Dependency Details tab	110
12.5.7	Dependents tab	113
12.5.8	Required Resources tab	114
12.5.8.1	Resource Details tab	115
12.5.9	Parameters tab	119
12.5.9.1	Parameter Details tab	120
12.5.9.2	Predefined default parameters for the runtime system	124
12.5.10	References tab	127
12.5.11	Triggers tab	128
12.5.11.1	Trigger Details tab	128
12.5.12	Triggered by tab	132
12.6	Job Hierarchy navigation	133
13	schedulix!Web Users	137
13.1	View	137

13.2 Concept	137
13.2.1 In short	137
13.2.2 Detailed description	137
13.3 Navigator	138
13.4 Editor	138
13.4.1 schedulix!Web Connection	138
13.4.2 schedulix!Server Connections	140
14 schedulix Server Users	141
14.1 View	141
14.2 Concept	141
14.2.1 In short	141
14.2.2 Detailed description	141
14.3 Navigator	141
14.4 Editor	142
15 Groups	145
15.1 View	145
15.2 Concept	145
15.2.1 Description	145
15.3 Navigator	145
15.4 Editor	146
15.4.1 Properties tab	146
15.4.2 Grants	146
16 Time Scheduling	149
16.1 View	149
16.2 Concept	149
16.2.1 In short	149
16.2.2 Detailed description	150
16.3 Navigator	150
16.4 Editor	150
16.4.1 Sub Schedule Details sub-area	154
17 Submit Batches and Jobs	161
17.1 View	161
17.2 Concept	161
17.2.1 In short	161
17.2.2 Detailed description	161
17.3 Navigator	162
17.4 Editor	162
18 Bookmarks	165
18.1 View	165

18.2 Concept	165
18.2.1 In short	165
18.2.2 Detailed description	165
18.3 Navigation	167
19 Running Master Jobs	169
19.1 View	169
19.2 Concept	169
19.2.1 In short	169
19.2.2 Detailed description	169
19.3 Master Navigator	169
19.3.1 Master Navigator Query mask	175
19.4 Detail Navigation	180
19.4.1 Tree view	181
19.4.2 Search results	181
19.4.3 Detail navigation query mask	181
19.5 Details mask for jobs	182
19.5.1 Buttons	182
19.5.2 Properties tab	186
19.5.3 Run tab	188
19.5.4 Timestamps tab	192
19.5.5 Dependencies tab	193
19.5.6 Dependents tab	197
19.5.7 Resource(Req) tab	200
19.5.8 Parameters tab	201
20 Search Running Jobs	203
20.1 View	203
20.2 Concept	203
20.2.1 In short	203
20.2.2 Detailed description	203
20.3 Navigator	203
20.4 Navigator Details query mask	204
21 Calendar	207
21.1 View	207
21.2 Concept	207
21.2.1 In short	207
21.3 Detailed description	207
21.3.1 Query tab	208
21.3.2 List tab	208
21.3.3 Day tab	209
21.3.4 Week tab	210
21.3.5 Month tab	210

22 System Information	211
22.1 View	211
22.2 Concept	212
22.2.1 In short	212
22.2.2 Config tab	212
22.2.3 System tab	220
22.2.4 Worker tab	222
22.2.5 Sessions tab	223
22.2.6 SME/Q tab	225
Index	232

List of Tables

1.1 Description of the icons	4
11.1 Lock mode matrix	78
11.2 Description of the job server configuration parameters	82
22.1 ParameterHandling options	216
22.2 Overview of the Trace Levels	219
22.3 Session status	225

List of Figures

1.1	Login window	1
1.2	schedulixLauncher	2
1.3	Main Desktop	3
1.4	Example of a standard dialog	5
1.5	Busy message	5
1.6	About window	6
1.7	Navigator with a hierarchical view	7
1.8	Searching in the Navigator	8
1.9	Search results	9
1.10	Editor with tabs	9
1.11	Show Folder Path	11
1.12	Hide Folder Path	12
1.13	Show Hierarchy Path	12
1.14	Hide Hierarchy Path	13
1.15	Reverse Selection; Initial State	13
1.16	Reverse Selection; Result	14
1.17	Choosing a Resource State Profile; initial state	15
1.18	Choosing a Resource State Profile; selection list	16
1.19	Choosing a Resource State Profile; result	16
1.20	Example of how standard lists are handled	16
1.21	Handling lists; adding a line	17
1.22	Handling lists; choosing a value	17
1.23	Handling lists; result	18
1.24	Handling lists; deleting a line	18
1.25	Handling lists; editing	19
1.26	Moving objects; initial state	20
1.27	Moving objects; Content tab in the Editor	21
1.28	Moving objects; selection	21
1.29	Moving objects; cut	21
1.30	Moving objects; choosing the target folder	22
1.31	Moving objects; paste	22
1.32	Moving objects; result	23
1.33	Linking objects; initial situation	23
1.34	Linking objects; selection	24
1.35	Linking objects; paste	24
1.36	Legend for the graphical representation of schedulix objects	25

1.37	Legend for the graphical representation of parent-child relationships	25
1.38	Legend for the graphical representation of dependency relationships	26
1.39	Recording comments	27
2.1	Exit State Definitions	29
3.1	Exit State Mappings	31
3.2	Exit State Mapping Editor	32
4.1	Exit State Profiles	35
4.2	Exit State Profile Editor	36
5.1	Resource State Definitions	39
6.1	Resource State Profiles	41
6.2	Resource State Profile Editor	42
6.3	State chart for a resource	43
7.1	Resource State Mappings	45
7.2	Example for a Resource State Mapping	46
7.3	Resource State Mapping Editor	46
8.1	Named Resources	49
8.2	Categories; Properties tab	50
8.3	Categories; Content tab	51
8.4	Named Resources; Properties tab	51
8.5	Named Resources; Parameters tab	53
8.6	Named Resources; Parameter Details	53
8.7	Named Resources; instances	54
8.8	Named Resources; Job Definitions tab	55
8.9	Named Resources; selector	57
9.1	Environments	59
9.2	Environment navigator	60
9.3	Environment properties	61
9.4	Environment references	61
10.1	Footprints	63
10.2	Footprint properties	64
10.3	Footprint references	65
11.1	Job servers and resources	67
11.2	Job servers and Scope Navigator	69
11.3	Scope properties	69
11.4	Job server properties	70
11.5	Job server and scope resources	72

11.6	Resource Details	73
11.7	Resource parameters	75
11.8	Resource allocations	75
11.9	Job server and scope parameters	79
11.10	Job server and scope configuration	80
11.11	Job server environment mapping	83
11.12	Job server log file name patterns	84
12.1	Batches and Jobs	85
12.2	Navigator for batches and jobs	88
12.3	Batches and Jobs; view without a pin	88
12.4	Batches and Jobs; view with a pin	89
12.5	Folder properties	90
12.6	Folder content	92
12.7	Folder parameters	92
12.8	Folder resources	92
12.9	Job Definition properties	94
12.10	Examples of batches	95
12.11	Run Job tab	98
12.12	Job Restart tab	102
12.13	Batches and Jobs; Children tab	103
12.14	Batches and Jobs; Child Details tab	104
12.15	Static and dynamic children	105
12.16	Example of Merge Mode	107
12.17	Example of Ignored Dependencies	108
12.18	Batches and Jobs; Parents tab	109
12.19	Batches and Jobs; Dependencies tab	109
12.20	Example of Dependency Modes	110
12.21	Batches and Jobs; Dependency Details	111
12.22	Example for Unresolved Handling	112
12.23	Batches and Jobs; Dependents tab	114
12.24	Jobs; Required Resources tab	114
12.25	Batches and Jobs; Resource Requirement details	116
12.26	Example of Sticky Handling	117
12.27	Example of load distribution with Sticky Handling	118
12.28	Batches and Jobs; Parameters tab	119
12.29	Batches and Jobs; Parameter Details	120
12.30	Example using Expression parameters	121
12.31	Batches and Jobs; References tab	127
12.32	Batches and Jobs; Triggers tab	128
12.33	Batches and Jobs; Trigger Details tab	129
12.34	Example of an Immediate Merge trigger	130
12.35	Example of an Immediate Local trigger	131

12.36	Example of a Before Final Trigger	131
12.37	Jobs and batches; Triggered By tab	133
12.38	Batches and Jobs; hierarchy view	133
12.39	Hierarchy view with dependencies	134
12.40	Hierarchy view with multiple dependencies	134
12.41	Cyclic dependencies	135
13.1	User management for the Web front end	137
14.1	schedulix Server user management window	141
15.1	Group management	145
15.2	Group properties	146
15.3	Object privileges for a group	147
16.1	Time Scheduling	149
16.2	Time Scheduling Navigator	151
16.3	Time Scheduling Editor	152
16.4	Repeat Driver	155
16.5	Time Of Day Driver	155
16.6	Range of Day Filter	155
16.7	Day of Week Filter	156
16.8	Day of Month Filter	156
16.9	Week of Month Filter	157
16.10	ISO Week of Month Filter	157
16.11	Week of Year Filter	158
16.12	Month of Year Filter	158
16.13	Calendar Driver	159
16.14	Calendar Filter	160
17.1	Submit Batches and Jobs	161
17.2	Submit navigation	162
17.3	Submit editor	162
18.1	Bookmarks	165
19.1	Running Master Jobs	169
19.2	Running Master Jobs Navigator	170
19.3	State chart for batches and jobs	175
19.4	Running Master Jobs Query mask	176
19.5	Running Master Jobs "Details" window	180
19.6	Navigation Details query mask	181
19.7	Details mask for jobs	183
19.8	Confirm mask	183
19.9	Job and batch properties	187

19.10	Batch und Job Run information	189
19.11	Batch and job timestamps	192
19.12	Batch und job dependencies	193
19.13	Batches and Jobs Ignore Dependencies confirmation mask	193
19.14	Example of a Recursive Ignore	194
19.15	Example of a Job Only Ignore	195
19.16	Batch and Job Dependents tab	197
19.17	Batch and Job Ignore Dependencies	197
19.18	Job Resource Requirements	200
19.19	Job Ignore Resource Requirement	201
19.20	Batch and job parameters	202
20.1	Search Running Jobs	203
20.2	Search Running Jobs query mask	204
21.1	Calendar	207
21.2	Calendar Query tab	208
21.3	Calendar List tab	209
21.4	Calendar Day tab	209
21.5	Calendar Week tab	210
21.6	Calendar Month tab	210
22.1	System Information	211
22.2	System configuration	212
22.3	System runtime environment	221
22.4	Worker thread information	223
22.5	Session information	223

1 General

1.1 Overview

This chapter will introduce you to the schedulix!web user interface. General actions and procedures that are not just to be found in a particular dialog are explained and illustrated here in more detail.

The schedulix!web user interface is a browser-based dialog system. There is no need to install an additional client software.

To start the schedulix!web interface you have to open a web browser window (e.g. Microsoft Internet Explorer). Please enter the URL you have been given to get to the schedulix!web interface. We recommend saving this URL in your browser as a bookmark so you can start the schedulix!web interface with one click.




Figure 1.1: Login window

Main Desktop

1.2 Login

After entering the URL, a login window opens for verifying your password. All users have to identify themselves with a user name and password to work with schedulix.

The input mask for the password prompt is shown in Figure 1.1, whereby the exact appearance may vary slightly depending on the browser.

Here you need to enter your personal user name and matching password and confirm them by clicking *OK*. Your login details can be obtained from your system administrator.

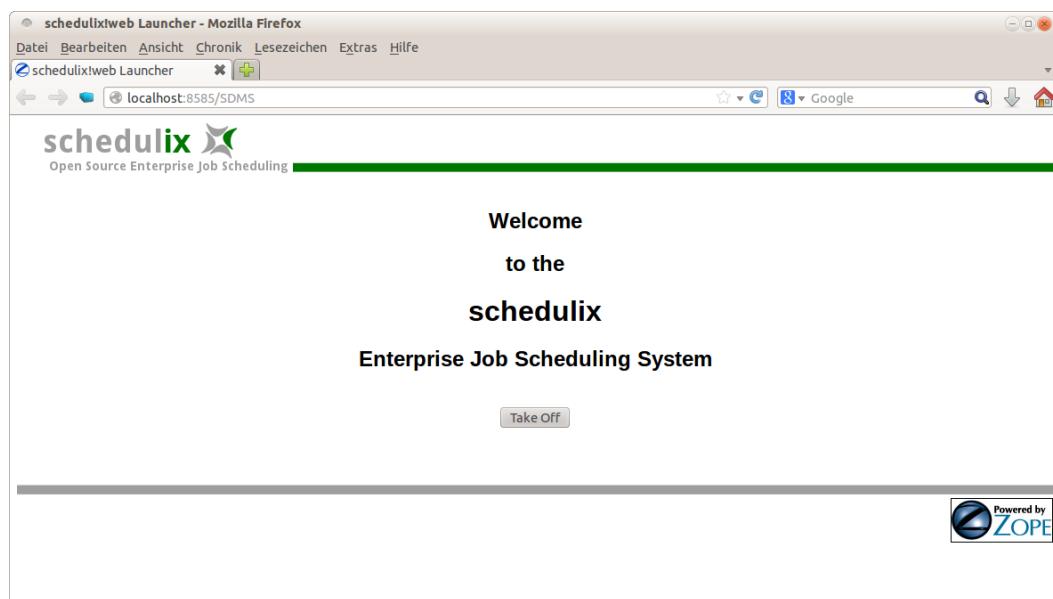


Figure 1.2: schedulixLauncher

After you have successfully logged in, the schedulixLauncher opens. Now open the main menu by clicking the *Take Off* button.

The launcher may be closed automatically depending on your browser. Your browser's security settings may prevent this or trigger a warning message.

After clicking the *Take Off* button appears the main menu called the *Main Desktop*.

1.3 Main Desktop

The *Main Desktop* is the main window of the schedulix!web user interface. You can launch all the function dialogs in schedulix directly from this window. From the "Connection" options box in the header you can select which server connection you want to use for the next function dialog that is to be opened.

Main Desktop

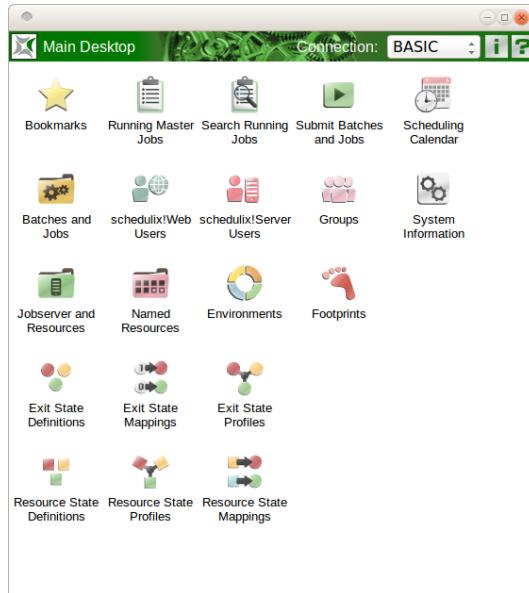


Figure 1.3: Main Desktop

The *Main Desktop* is shown in Figure 1.3.

The dialog boxes for maintaining or monitoring different objects can be opened by clicking the icons shown below:

Icon	Description
	The <i>Bookmarks</i> dialog is used to create, edit and delete views of currently running jobs.
	The <i>Running Master Jobs</i> dialog gives a clear view of the current Master Jobs.
	The <i>Search Running Jobs</i> dialog is used to search for running jobs according to specific criteria.
	The <i>Submit Jobs</i> dialog is used for starting workflows.
	The <i>Batches and Jobs</i> dialog is used to create and manage folders, batches and jobs.
	The <i>schedulix!web Users</i> dialog is used for the administration of schedulix!web users.
	The <i>schedulix!server Users</i> dialog is used for the administration of schedulix!server users.
	The <i>Groups</i> dialog is used for the administration of schedulixuser groups.

Continues on next page

Layout of the object window

Continued from previous page

Icon	Description
	The <i>Job Servers and Resources</i> dialog is used for the administration of job servers, scopes and associated resources.
	The <i>Named Resources</i> dialog is used for creating, modifying and deleting Named Resources.
	The <i>Environments</i> dialog is used to define runtime environments for jobs.
	The <i>Footprint</i> dialog is used for the administration of footprints (resource profiles) in the system.
	The <i>Exit State Definition</i> dialog is used for defining and administering Exit State names.
	The <i>Exit State Mappings</i> dialog is used for the administration of translations between return values from the system processes and logical Exit States.
	The <i>Exit State Profile</i> dialog is used for the administration of profiles, i.e. a summary and prioritisation of Exit States.
	The <i>Resource State Definition</i> dialog is used for defining and administering Resource State Names.
	The <i>Resource State Profile</i> dialog is used for the administration of profiles, i.e. a summaries of Resource States.
	The <i>Resource State Translation</i> dialog is used for the administration of translations of Exit States after Resource States have been changed.
	The <i>Calendar</i> dialog shows an overview of the pending submits.
	The <i>SysInfo</i> dialog shows information about the configuration and current state of the system.

Table 1.1: Description of each icon in the main menu

1.4 Layout of the object window

When a dialog box is called from the *Main Desktop*, it is opened in a new window. Each window is independent of other dialog windows and any number of dialogs can be open at the same time.

A standard dialog looks like this:

1.4.1 Title bar

Each dialog has a title bar at the top. The title bar contains the elements described below.

Layout of the object window




Figure 1.4: Example of a standard dialog



independIT icon/hourglass

The independIT icon is located in the upper left corner of the window. Clicking this icon opens a new window with the website of independIT GmbH.

When you perform an action, the new or modified data has to be sent to the schedulix server or data has to be downloaded from it. This activity is indicated by the hourglass icon. No further actions can be performed in the window for as long as the hourglass icon is active. If you do try to perform an action, this will trigger an error




Figure 1.5: Busy message

message as shown in figure 1.5

In this case, the box has to be confirmed by clicking OK and you have to wait until all the data has been transferred to or from the server. You can continue entering data as soon as the hourglass disappears and the IndependIT icon is displayed again.

Layout of the object window

Name of the dialog box The name of the active dialog box is shown next to the independIT icon.



About icon

Version and system information can be displayed by clicking this button.




Figure 1.6: About window

The following information is shown in this window.

schedulix!web n.n.n: Current version of the schedulix!web user interface.

Authentication: Name of the user who is currently logged in.

schedulix Server Connection: Connection data for the schedulix server.



Help icon

The Help icon is only visible on the Main Desktop instead of the Home icon. Clicking the Help icon takes you straight to the online documentation.



Home icon

The Home icon is used to call the Main Desktop. If the Main Desktop window has been minimised or it is in the background, it is restored and brought into focus. If it has been closed, clicking the icon opens a new Main Desktop window.

1.4.2 Navigator

The left pane of a dialog window is the Navigator. It shows the objects that have already been defined. These are displayed either in a hierarchy (if the type of objects can be arranged hierarchically, for example **named resources**) or as a simple list if there is no hierarchical order. The hierarchical view is a tree structure that looks like in figure 1.7.

Layout of the object window




Figure 1.7: Navigator with a hierarchical view

If an object displayed in the hierarchical view is a container that can contain other objects, this is indicated by a folder icon like e.g.



These are different objects (folders, scopes, named resources, etc.) depending on their colour and how they are displayed. If a container holds sub-objects as well, a or is displayed. By clicking this icon, the contents of the container can be displayed or hidden .

1.4.2.1 Standard buttons

Clicking the buttons above the Navigator as described below controls the window view and whether contents are displayed or hidden:



Refresh button

The *Refresh* button is used to refresh the list of objects in the navigation. After clicking this button, all the available objects are read out from the server again.



Collapse All

With this button, an expanded hierarchy (i.e. a hierarchy in which the sub-folders are displayed and the parent folder has a) is completely closed again. This means that only the objects at the top level of the hierarchy are visible.



Expand All

This button is used to fully open a closed or partially open hierarchy. This means that all the top level containers are scanned for sub-containers or sub-objects and these are displayed. The folder icon is preceded by the character. If the sub-containers hold further sub-objects, these are also displayed until all the objects throughout the entire hierarchy are shown on the screen. With a large hierarchy containing a large number of objects and sub-objects, this can take a relatively long

Layout of the object window

time. For this reason, the button is not available for every hierarchical object type.



These two buttons allow non-container objects to be displayed and hidden in hierarchies in the Navigator. Only one button is alternately always visible.



These two buttons allow the names of objects without read privileges to be displayed and hidden. Only one of these two buttons is alternately always visible.



The *Save* button is used to save the current state of the hierarchy in the Navigator.



This button activates the search mask in the navigation. This enables you to very easily find an object with a known name even in a complex hierarchy. Clicking the *Search* button opens the search mask:




Figure 1.8: Searching in the Navigator



This starts the search. In the example shown here, a search is being run in the entire hierarchy for named resources that contain the search string.

Hits are shown as a list in the Navigator. Hierarchies are not taken into consideration. The result could look like this:

Here we were searching for the term "host", and three Named Resources were found that contain this term. The search is not case-sensitive with regard to the search term.



Clicking this button closes the search mask and the standard button bar is displayed above the Navigator again.

Layout of the object window

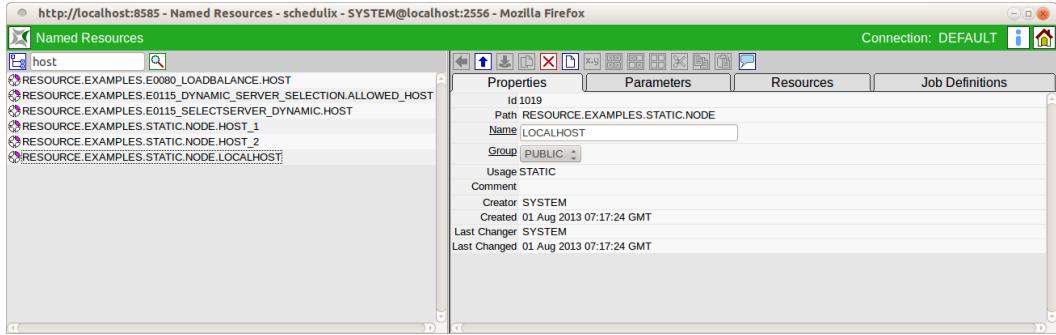


Figure 1.9: Search results

1.4.3 Editor

The Editor is the dialog window for modifying or creating objects. When an object is selected in the Navigator, its data is displayed in the Editor and can be modified here. A selected object can be saved, deleted or duplicated using the buttons at the top of the Editor. A new object can be created by clicking the *New* button. The data in the Editor is either displayed in full in one window or it is logically grouped in different tabs.




Figure 1.10: Editor with tabs

The number of tab sheets in a dialog varies depending upon the type of the object selected in the navigation. During an action (modifying or creating an object), you can switch between the single tabs as often as you need to and in any order. You

Layout of the object window

do not need to save your changes in between because all the data (including the data in hidden tabs) is kept in the memory. Only when you click the *Save* button is the data in all the tabs sent to the schedulix server and saved there. If any changes have been undone by clicking the *Cancel* button, all the modified data (including in any hidden tabs) is discarded.

A tab sheet is selected simply by clicking the respective tab on which the name of the tab sheet is shown. The data from the selected tab is then displayed in the mask and the tab is graphically moved into focus and highlighted in a light colour. The inactive tabs are dark in colour.

The Editor button bar shows different enabled and disabled buttons depending upon which tab is open.

1.4.3.1 Standard buttons

The following standard buttons are shown in many dialog boxes and are used in an identical or comparable way in each dialog.



Cancel

This button is used to undo actions. All changes that have not yet been saved can be revoked. This means that all changes made since the data was loaded or last saved are discarded.

The *Cancel* button is only enabled if a change has been made. Changes are recognised only when you leave the first field in which a change was made. If there is only one field in the mask, you have to press the keyboard tabulator key to accept your changes. This then enables the *Cancel* button.

To prevent entered data from being lost, a confirmation query is displayed after you have clicked the button. If you really want to discard your changes, you have to confirm this query by clicking *OK*. Otherwise, you can click *Abort* to return the previous dialog and preserve the changes you have made.



Up

Clicking the *Up* button takes you up one level in the hierarchy.



Save

A change that has been made is saved with this button. The *Save* button is only enabled if data has actually been changed. Changes are only recognised once you have left the changed input field.



Clone

This button is used to create a new object which possesses all the properties of the currently selected object. To enable the button, you have to enter a new name for the object. When you click the button, a new object is created with this name and displayed in the navigation window.

Important: If you click *Save* instead of the *Clone* button, only the name is changed, i.e. a new object is not created.

Layout of the object window



This button is used to delete an entire object. A confirmation query is displayed before the object is displayed.


If you really want to delete the object, confirm the query *OK*. If you want to cancel the operation, press the *Abort* button.

The object is deleted from the server after the action has been confirmed. There is now no way to undo the deletion. If the deletion was unintentional, you will have to create the deleted object again.

In the lists, you can also delete selections. The procedure here is as follows. First, you have to select the rows you want to delete either individually or by clicking the *Selection* button. Now click *Drop*. All the selected entries are deleted after you have confirmed the prompt.



This button is used to create a new object. If different objects can be created in a dialog, first you have to select the type for the new object. The input fields and tabs for the corresponding object type are then displayed. If a selection is not necessary, the Editor opens immediately after you click the button. If an Editor window still contains any unsaved changes, the button is disabled to prevent the changed values from being accidentally deleted.




These buttons are switches. When you click the *Show Folder Path* button, the scopes, submitted entities, scheduling entities and folders are displayed together with the complete parent hierarchy. The levels are each separated by a full stop. The view looks this after you have clicked the button:

#	Job	JobId	MasterId	Type	Amount	Keep	Sticky	Lockmode	Mapping	P	EP
1	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[1]	9024	9020	ALLOCATION	1	NONEKEEP	false	N	NONE	50	45
2	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[2]	9029	9020	ALLOCATION	1	NONEKEEP	false	N	NONE	50	45
3	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[3]	9034	9020	ALLOCATION	1	NONEKEEP	false	N	NONE	50	45
4	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[4]	9039	9020	ALLOCATION	1	NONEKEEP	false	N	NONE	50	45
5	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[5]	9044	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45
6	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[6]	9049	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45
7	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[7]	9054	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45
8	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[8]	9059	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45
9	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[9]	9064	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45
10	SYSTEM.EXAMPLES.E0075_LOADCONTROL.CHILD[10]	9069	9020	BLOCKED	1	NONEKEEP	false	N	NONE	50	45

Figure 1.11: Show Folder Path

When you click *Hide Folder Path*, the previously visible parent folders are not displayed. The window looks like this:



Layout of the object window

#	Job	Jobid	MasterId	Type	Amount	Keep	Sticky	Lockmode	Mapping	P	EP
1	CHILD[1]	9024	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
2	CHILD[2]	9029	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
3	CHILD[3]	9034	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
4	CHILD[4]	9039	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
5	CHILD[5]	9044	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
6	CHILD[6]	9049	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
7	CHILD[7]	9054	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
8	CHILD[8]	9059	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
9	CHILD[9]	9064	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
10	CHILD[10]	9069	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45

Figure 1.12: Hide Folder Path

These buttons are switches. Clicking the *Show Hierarchy Path* button displays the task with the complete superordinate parent-child hierarchy. If a task has a parent, this is shown in the name and so on until the top level has been reached. The levels are each separated by a colon.

#	Job	Jobid	MasterId	Type	Amount	Keep	Sticky	Lockmode	Mapping	P	EP
1	LOADCONTROL:CHILD[1]	9024	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
2	LOADCONTROL:CHILD[2]	9029	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
3	LOADCONTROL:CHILD[3]	9034	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
4	LOADCONTROL:CHILD[4]	9039	9020	ALLOCATION	1	NOKEEP	false	N	NONE	50	45
5	LOADCONTROL:CHILD[5]	9044	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
6	LOADCONTROL:CHILD[6]	9049	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
7	LOADCONTROL:CHILD[7]	9054	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
8	LOADCONTROL:CHILD[8]	9059	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
9	LOADCONTROL:CHILD[9]	9064	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45
10	LOADCONTROL:CHILD[10]	9069	9020	BLOCKED	1	NOKEEP	false	N	NONE	50	45

Figure 1.13: Show Hierarchy Path

When you click *Hide Hierarchy Path*, the previously visible hierarchy is no longer displayed. The view looks this after you have clicked the button:



Select All

This button is for selecting all the objects. This means that all the objects currently displayed in the Navigator or the Editor window (depending on the button bar where the button is) are selected. All the objects are then marked with a cross in front of their name.

Once all the objects have been selected, they can then be cut, copied or deleted.



Toggle Selection

This button is for reversing the current selection. This means that all the selected objects are deselected and all the objects that have not yet been selected are selected.

Layout of the object window

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "Jobserver and Resources". The main content area has tabs for "Resource Detail", "Parameters", and "Allocations". Under "Resource Detail", there is a table titled "GLOBAL EXAMPLES.LOCALHOST.SERVER" with columns: #, Job, JobId, MasterId, Type, Amount, Keep, Sticky, Lockmode, Mapping, P, EP. The table lists 10 entries from 1 to 10, all of which are selected (indicated by a blue border around each row). The "Allocations" tab is visible but empty.

Figure 1.14: Hide Hierarchy Path

The selected objects can now be deleted, cut or copied.
This behaviour is illustrated in the following two screenshots.

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "Job, Batch and Milestone Definitions". The main content area has tabs for "Properties", "Content", and "Parameters". On the left is a tree view of definitions under "SYSTEM.EXAMPLES.E050_COMPLEX_CONDITION": E050_COMPLEX_CONDITION, COMPLEX_CONDITION, LOAD_A, LOAD_B, LOAD_C, LOAD_COMPLETED, MAIN_LOAD, REPORT, SELECT_LOAD. The "Content" tab shows a list of these same items, with LOAD_COMPLETED and SELECT_LOAD checked (indicated by a blue border).

Figure 1.15: Reverse Selection; Initial State

Your selection is reversed when you click the Toggle Selection button.



Deselect All

This button reverses your selection and the markings on the selected entries are removed. This allows you to undo a selection that you have made.



Cut

This button is used to cut out objects out of a list or a tree. The objects are copied to the schedulix clipboard. The objects are only marked for cutting, but they disappear from the current view. They can be inserted somewhere else by clicking the *Paste* button.



Copy

With the *Copy* button, objects can be copied to the schedulix clipboard for inserting them somewhere else by clicking *Paste*.

Selecting values




Figure 1.16: Reverse Selection; Result



Paste

This button is used to paste previously selected and suitable objects into a list. The objects have to be selected beforehand in another dialog or pop-up window.



Select

With this button you can open a window where you can select objects that are suitable for a list and copy them to the clipboard by clicking the *Copy* button. The objects in the clipboard can then be inserted into the list by clicking the *Paste* button.



Time Scheduling

This button opens the **Time Scheduling** dialog window. It is only displayed if you have selected a Master Submittable Object in the navigation window.



Grants

This button opens the dialog for administering the access privileges.



Edit


The *Edit* button activates an edit mode. What data is edited depends upon the context.

1.5 Selecting values

1.5.1 Selecting values from drop-down lists

With some input fields it is necessary to make a selection from a range of default values. In these cases, a *drop-down* box is displayed containing a list of all the possible values.

A *drop-down* box looks like this:



A single value is selected from the drop-down list by expanding it

Handling standard lists

by clicking the button next to the box or using the keyboard shortcut ALT-DOWN ARROW.

The expanded drop-down list can look like this:



Clicking the required value places it into the field. Alternatively, you can select the values using the DOWN ARROW or UP ARROW keys.

1.5.2 Choosing values using the Select button

If the list of available values is only read from the server at runtime you have to use the *Select* button. In the next screenshot, a Resource State Profile can be chosen by clicking *Select*:




Figure 1.17: Choosing a Resource State Profile; initial state

When you click the *Select* button,



a list of possible objects is displayed in the Editor.

Clicking one of these objects places the element into the field. You can jump back from the selection mask to the editing window without making a selection by clicking the *Cancel* button.

1.6 Handling standard lists

In many dialog windows it is possible to not only enter values into the fields, but you can also create lists with links to additional objects. This way, for example, a list of Resource States belonging to a Resource State Profile is administered in that profile. Here is an example:

Here, to the standard fields *Name* and *Default Initial State* we want to append a list of states.

The following actions can then be performed in such a list.

Handling standard lists

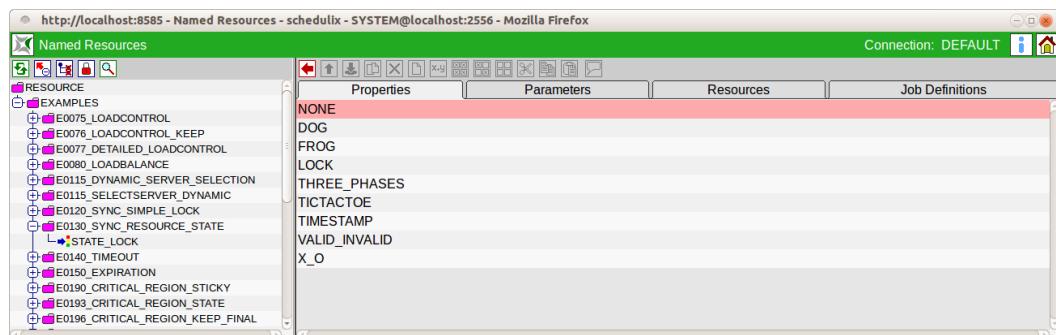


Figure 1.18: Choosing a Resource State Profile; selection list

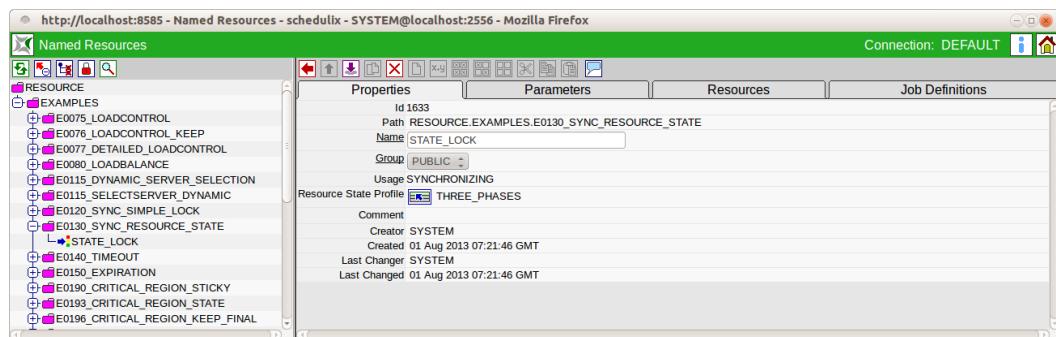


Figure 1.19: Choosing a Resource State Profile; result




Figure 1.20: Example of how standard lists are handled

Handling standard lists

1.6.1 Adding a line



A new line is added to the open list by clicking the *Add* button. An empty line appears after you have clicked the button.




Figure 1.21: Handling lists; adding a line

In our example, the dialog looks like in figure 1.21.

A value can then be chosen using the Select button.




Figure 1.22: Handling lists; choosing a value

In this example, the value "Valid" was chosen.

You now have to save the change with the Save button.

You can also modify an existing line by choosing a dependent object using the *Select* button. In this case, the link to the old object is deleted and a new link is created to the selected object.

Handling standard lists

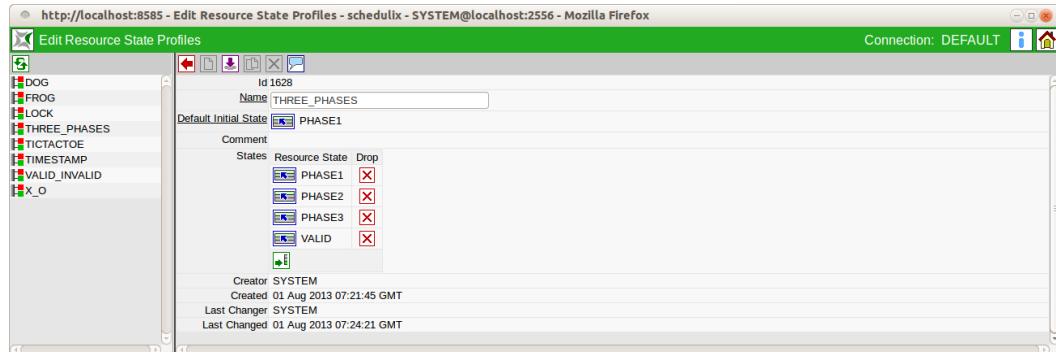


Figure 1.23: Handling lists; result

1.6.2 Deleting a line



Drop (delete)

When you click the *Drop* button, this removes a line and the link between the object loaded in the Editor and the selected object in the line is deleted after you have confirmed the prompt. In our example, clicking the *Drop* button will delete the line "Phase1".




Figure 1.24: Handling lists; deleting a line

The line with Phase1 in it no longer exists once you have clicked *OK*. The object in the line (i.e. the Resource State "Phase1") still exists, but it is no longer allocated to this profile.

The change now has to be saved on the server by clicking the *Save* button.

1.6.3 Changing line values

In many dialogs, not only does a line hold additional properties from other objects, but additional fields also have to be maintained for this link. These are shown in the list as a field string after the linked object. In the Footprint dialog, for example, not only does a Named Resource have to be selected in a line, but the other fields *Amount* and *Keep* need to be maintained as well.

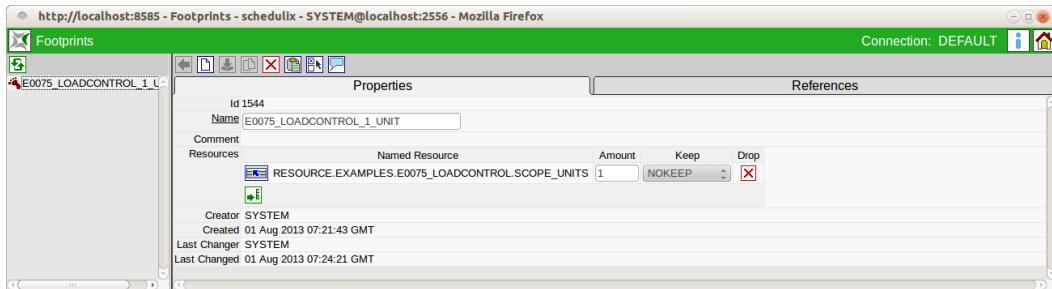


Figure 1.25: Handling lists; editing

This is done analog to how standard fields are edited.

The changes now have to be saved by clicking the *Save* button.

1.7 Copy & paste and the clipboard

The schedulix!web user interface features an intelligent copy & paste functionality and an object-sensitive clipboard. This allows objects to be selected, cut, inserted and deleted between different dialog boxes or in the same dialog. So long as the action has not been performed, the data is located in the clipboard of the schedulix!web interface system and can be used by other suitable dialogs.

The term *object-sensitive* describes the logic routine executed when inserting objects from the clipboard. Before objects are inserted from the clipboard, the system checks whether they are of the same type as the required objects in the dialog into which they are to be inserted. It is therefore not possible to insert an object of type A if the recipient dialog is expecting an object of type B. If there are no objects of type B in the clipboard, the insert operation cannot be performed.

This mechanism is not to be confused with the standard clipboard system used by Windows operating systems. The schedulix clipboard functions totally independently of the Windows clipboard and vice versa. The data stored in the schedulix clipboard is available exclusively to schedulix dialogs and cannot be used or read out by other Windows applications.

How the clipboard is used is demonstrated in two examples in the next sections.

Copy & paste and the clipboard

1.7.1 Moving objects in the hierarchy

In a folder hierarchy, you can use the clipboard to move objects to another level in the hierarchy.

Example:

The navigation window for the *Batches and Jobs* dialog shows the following hierarchy:




Figure 1.26: Moving objects; initial state

Below the COMPLEX_CONDITION folder there are various jobs and batches. This folder needs to be reorganised. The batch COMPLEX_CONDITION is to be moved to the BATCHES folder, and then the jobs LOAD_A, LOAD_B and LOAD_C are to be moved to the JOBS folder.

This can be done as follows using the clipboard:

1. Select the objects to be moved

First you select the batch COMPLEX_CONDITION in the Navigator. The list of all the objects in the COMPLEX_CONDITION folder are now displayed in the CONTENT tab in the Editor.

The batch is now marked by clicking the check box in front of the name and the icon.

2. Cutting out the selected objects Using the *Cut* button, the selected batch is cut out and disappears from the screen view. The object has now been saved to the clipboard. The dialog looks like this after the batch has been cut out:

The important to remember here is that although the object that has been cut out is no longer visible, it is still in the folder. So if you close the window now, nothing has happened. It is only actually moved during the subsequent paste operation.

3. Selecting the target folder

Copy & paste and the clipboard

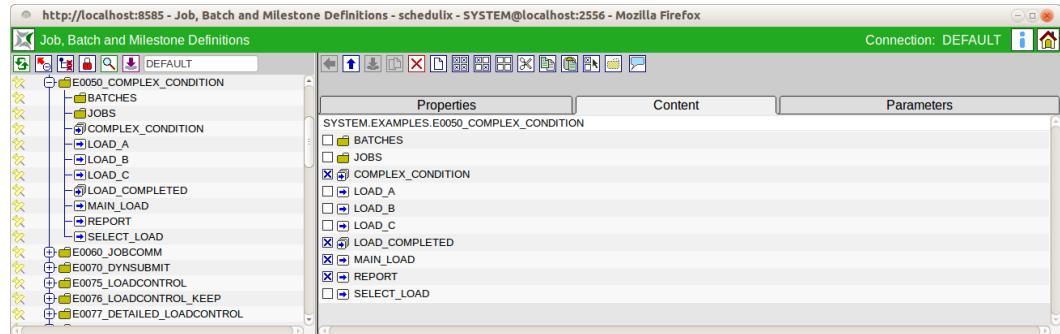


Figure 1.27: Moving objects; Content tab in the Editor

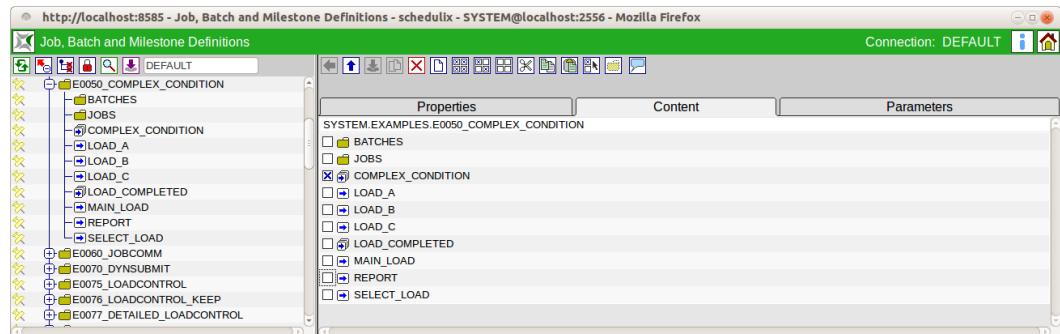


Figure 1.28: Moving objects; selection




Figure 1.29: Moving objects; cut

We now need to choose the target dialog and target folder for the move operation. This can be done simply by clicking the respective folder in the navi-

Copy & paste and the clipboard

gation pane of the target dialog. As the clipboard functions between different windows, the target folder can also be located in another dialog.

In our example, we have chosen the BATCHES folder and the Content tab opens.

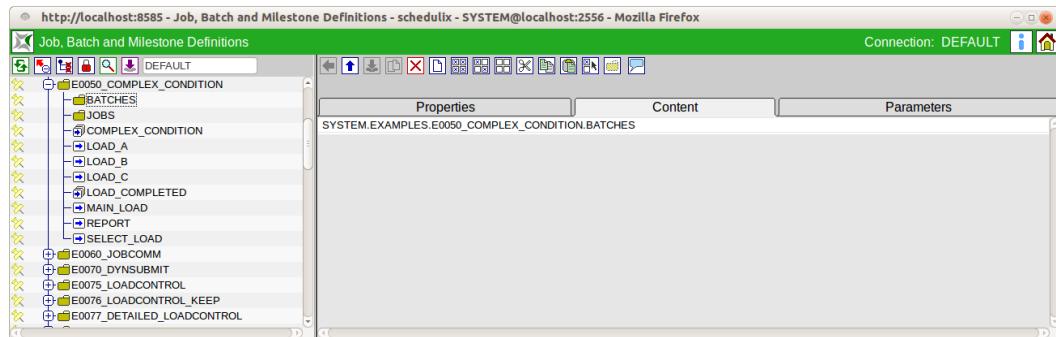


Figure 1.30: Moving objects; choosing the target folder

4. Inserting the selected objects

To complete the move, the object is inserted from the clipboard. This is done using the *Paste* button. Pasting the object now triggers the move action on the server. This action has now been completed and cannot be undone. The result looks like this:




Figure 1.31: Moving objects; paste

Performing the action with the aforementioned jobs and the JOBS folder produces the following hierarchy:

Copy & paste and the clipboard

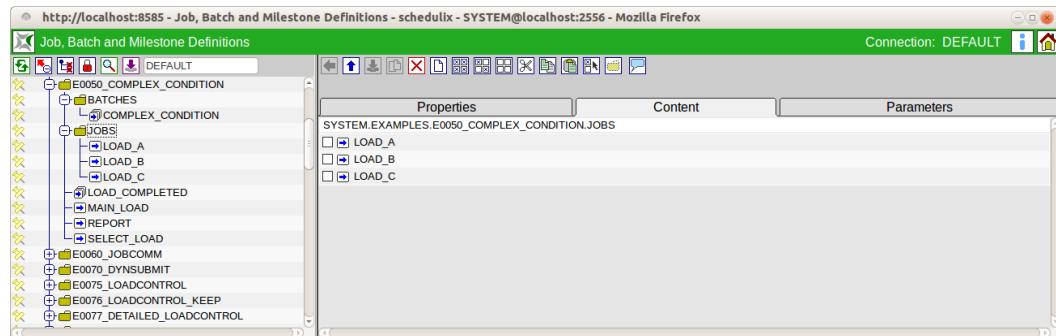


Figure 1.32: Moving objects; result

1.7.2 Linking objects

In some dialogs, objects are linked to other objects or lists of objects. An example of this is the definition of dependencies. The clipboard can also be used to fill such a list.

To illustrate this and demonstrate the procedure for such an action, we will define the dependencies for a job.

Here you can see the job END_PROCESSING, which does not yet have any dependencies:

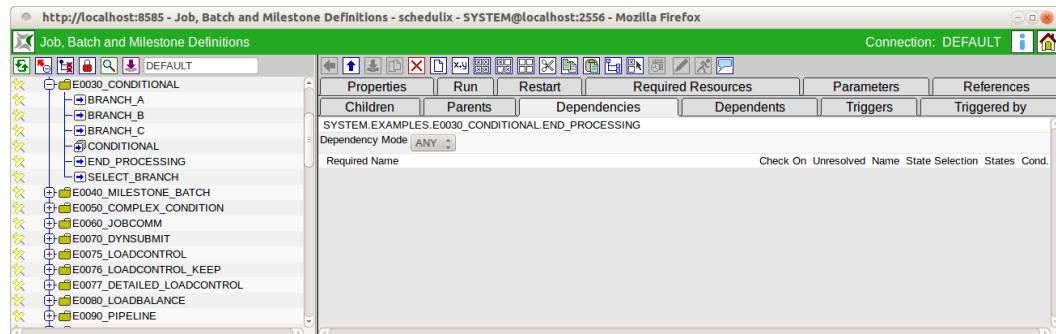


Figure 1.33: Linking objects; initial situation

Additional jobs have to be added at this point to make the execution of this job dependent upon other processing steps. These can be marked beforehand in other dialogs and copied to the clipboard. We now take the following steps:

1. Selecting the objects

Jobs can be marked in the Content tab of a folder, for example. In the next dialog, we have selected the three processing steps A, B and C.

Copy & paste and the clipboard




Figure 1.34: Linking objects; selection

2. Copying the objects to the clipboard

Clicking the *Copy* button copies the selected jobs to the clipboard.

3. Inserting the objects

In the Dependency Editor window of the END_PROCESSING job, we now click the *Paste* button and the cached jobs are added to the list. The list now looks like this:




Figure 1.35: Linking objects; paste

4. Saving the change

In contrast to the action for moving objects, which does not have to be additionally saved, copy actions in an editor window do need to be saved by clicking the *Save* button.

Important: The cached data remains in the clipboard. If the selected jobs are to be added to other dialogs, you just have to choose the appropriate target and then click the *Paste* button.

Graphic Legend

1.8 Graphic Legend

Different graphics are used in this documentation to illustrate correlations.

1.8.1 How schedulix objects are represented

The following symbols and their meanings are used in the graphics for schedulix objects:

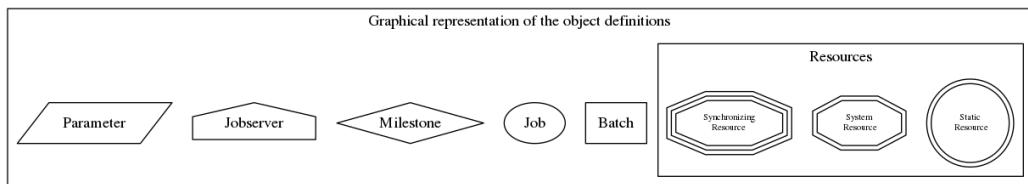


Figure 1.36: Legend for the graphical representation of schedulix objects

1.8.2 Representation of parent-child relationships in workflows

Relationships between parents and children are represented in the graphics as follows:




Figure 1.37: Legend for the graphical representation of parent-child relationships

The family tree is usually read from top to bottom from the highest ranking parent down to the respective children. The arrow always points from the parent to the children.

Comments

More details about parent-child relationships between tasks can be found in Chapter [Batches and Jobs](#).

1.8.3 Representation of dependencies between Scheduling Entities

Dependencies between individual Scheduling Entities are represented in the graphics as follows:

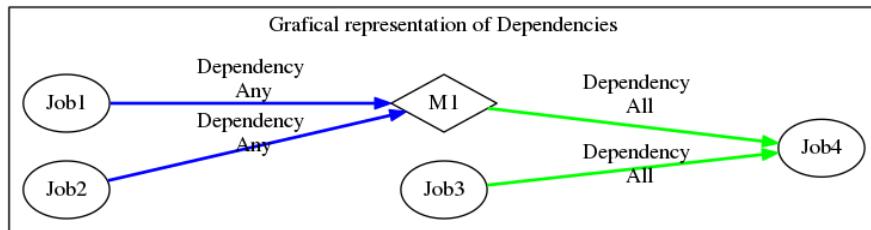


Figure 1.38: Legend for the graphical representation of dependency relationships

Dependencies are displayed from right to left. This means that the first job is shown on the far left, then comes the next one time-wise, etc.

The arrows always point from the required Scheduling Entity to the dependent Scheduling Entity. This thus shows the timeline for the Scheduling Entities. For example, A→B→C means that A is run first, and is followed by B first and then C. A is the required entity of B and B is the required entity of C. B is accordingly dependent upon A. C is dependent upon B (and so implicitly also dependent upon A).

The colour of the arrows corresponds to the interlinking of multiple dependencies. If the colour is green, all the required tasks are linked with "AND" (the dependency ALL).

If the colour is blue, all the required tasks are linked with "OR" (the dependency ANY).

If a task is only dependent upon precisely one other task, the arrow is green since in this case the links AND and OR are equivalent.

More details about dependency relationships between Scheduling Entities can be found in Chapter [Batches and Jobs](#).

1.9 Comments

Comments can be saved for most objects.



Clicking the *Comment* button opens the *Edit Comments* mask. Here you can enter any text comments relating to the object. Text comments consist of a tag and a

Standard fields

comment text. On display tags will be rendered as headlines. If the tag is left empty, the comment text will be concatenated with the preceding text comments comment text. Apart from entering the comment, no other action is performed on the object.

As an alternative to text comments, a link can be set to an intranet page instead. To do this, select "URL" as the type and enter the URL in the comment box. The links are also activated in the schedulix!Web interface. This allows any number of complex documents to be linked with the schedulix objects.

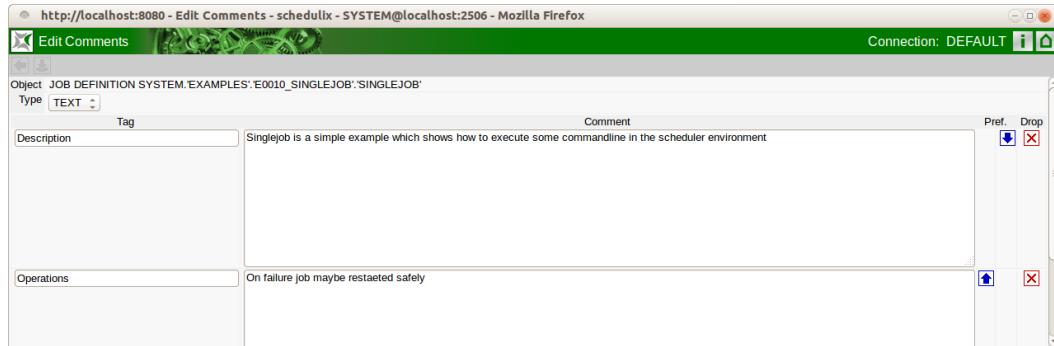


Figure 1.39: Recording comments

1.10 Standard fields

The following fields are defined for all objects:

Creator Name of the user who created the object.

Created Date and time at which this object was created.

Last Changer Name of the last user who changed the object.

Last Changed Date and time at which this object was last changed.

2 Exit State Definitions

2.1 View




Figure 2.1: Exit State Definitions

2.2 Concept

2.2.1 In short

The Exit State Definitions dialog is used to create logical names for exit states of a Submitted Entity.

2.2.2 Detailed description

Exit State Definitions are logical names for exit states of a Submitted Entity. These logical names do not have any other properties.

2.3 Editor

The properties of Exit State Definitions are maintained in this tab.

Exit State Definitions can only be edited by users who belong to the "ADMIN" group.

All the input fields are "read only" for all other users.

The fields have the following meanings:

Id System-wide unique number for identifying the object.

Editor

Name Unique name of the Exit State Definition. This name can be freely chosen.

3 Exit State Mappings

3.1 View




Figure 3.1: Exit State Mappings

3.2 Concept

3.2.1 In short

Exit State Mappings are used for translating numerical Exit Codes into logical Exit States.

3.2.2 Detailed description

Exit State Mappings are used to assign numerical value ranges of exit codes to logical Exit States. Each program that is run returns a numerical value to the calling program when it has finished. On UNIX systems, a value of 0 usually means that the program has been completed without any errors (SUCCESS). Values unequal to 0 indicate an error (FAILURE). However, deviating values with other meanings are also possible. For a better understanding of this logical meaning within schedulix Exit State Mappings can be used to assign value ranges to a logical name. The values range from -2^{31} to $2^{31} - 1$.

Editor

3.3 Editor




Figure 3.2: Exit State Mapping Editor

Names are assigned to numerical value ranges in this editor.

Exit State Mappings can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

Descriptions of all the standard buttons can be found in Chapter [1.4.3.1 Standard buttons](#).

The buttons in the list shown above have the following meanings:



Insert new value range

This button adds another row to the table of Exit Code Ranges. The row is inserted at the end of the table and allows a new start value to be entered. This new end value is determined automatically when you save or sort the data.



Resort list

An unsorted list that is created by adding new intervals can be arranged in the correct order with this button.



Select Exit State

This button is for selecting a previously defined **Exit State** from a list.



Delete value range

This button deletes a range of values together with the Exit State from the list. The start value of the next interval is replaced by the start value of the deleted interval, i.e. the next interval is enlarged automatically.

The fields have the following meanings:

Id System-wide unique number for identifying the object.

Name Unique name of the Exit State Mapping. This name can be freely chosen.

Editor

Exit Code Ranges The assignment of the respective value range is shown in a table. An Exit State can be assigned to a range of values in each row. The first row contains the start value -2^{31} , while the last row has the end value $2^{31} - 1$.

4 Exit State Profiles

4.1 View

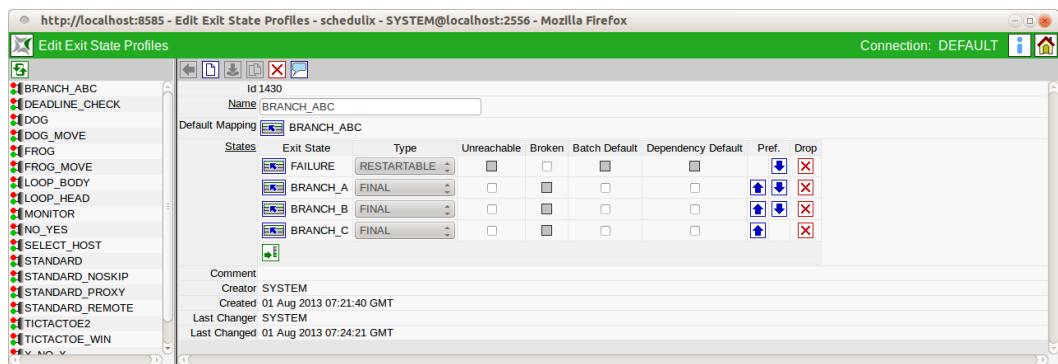


Figure 4.1: Exit State Profiles

4.2 Concept

4.2.1 In short

An Exit State Profile describes which Exit States can be achieved by a Submitted Entity. The Exit State Profile also defines whether a Submitted Entity can change its Exit State or if the Exit State is conclusive (FINAL).

4.2.2 Detailed description

Each Submitted Entity can finish with varying Exit States. The set of all valid Exit States for a Submitted Entity is described in an Exit State Profile. This description also defines the preferred Exit State that is crucial in a hierarchical workflow structure for determining the correct Exit State for a parent element. The Exit State with the highest preference is used.

Editor

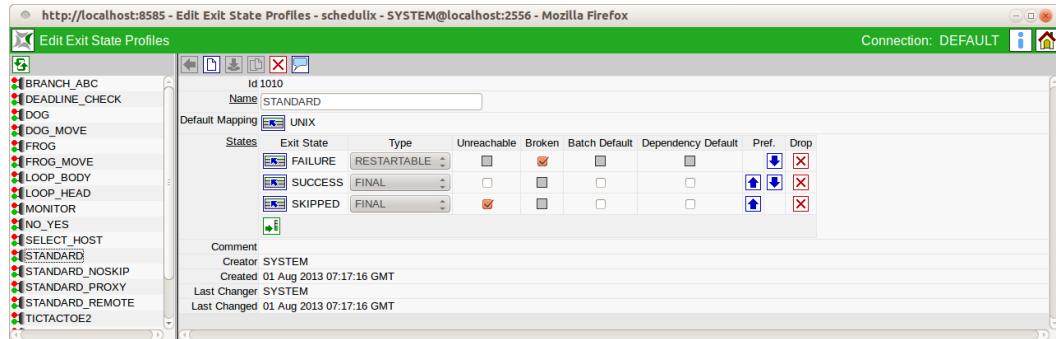


Figure 4.2: Exit State Profile Editor

4.3 Editor

In this editor, the Exit States belonging to a profile can be defined and the properties *Type*, *Unreachable*, *Broken* and *Preference* can be added to them.

Exit State Profiles can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

The fields shown above have the following meanings:

Name The *Name* field is used for assigning a unique name to an Exit State Profile.

Default Mapping The *Chooser* button is for selecting an **Exit State Mapping** from a list that is then used by default in conjunction with the Exit State Profiles provided that no other mapping has been defined in the job definition. An Exit State Mapping is used to translate numerical Exit Codes into Exit States.

The profile must contain an entry for each of these Exit States. However, not all Exit States in an Exit State Profile need to be accessible via the mapping.

Type One of three types can be assigned to each Exit State: FINAL, RESTARTABLE and PENDING.

The Exit State can no longer be changed once a task has attained its Final State. The result is final. Dependencies between jobs and batches can only be based on such FINAL Exit States.

If a task is to be capable of being restarted after it has been run, for example because an error has occurred, the Exit State must be of the type RESTARTABLE.

PENDING describes a state type that allows the Final State to be set externally via the API. A task in a PENDING state can be neither restarted nor are any dependencies fulfilled by it.

An example of how PENDING states are used:

A job sends an e-mail to an employee and requests the release of a result. After sending the e-mail, the job stops with a PENDING state. The employee can now manually set the state to FINAL. Alternatively, a process evaluating the response e-mail can change the status. Dependent tasks can only be started once the state is FINAL.

Unreachable A maximum of one Exit State in the list can be marked as being Unreachable. This Exit State is set when a job in a process can no longer be executed because the dependencies cannot be fulfilled.

The Unreachable State is not attained if one or more predecessors have been cancelled. Because the cancel operation is a manual intervention, the consequences of the intervention have to be handled manually as well.

The Unreachable State must be of the type FINAL.

Broken A maximum of one Exit State in the list can be marked as being Broken. This state is set when a job is put into an Error State due to a failure. Such an error occurs, for example, if the run program cannot be started. The Error State must be a Restartable State. The Broken flag allows a trigger to be used to automatically respond to such error situations.

Batch Default A final Exit State in the list can be labelled as the Batch Default. This state is set when a batch with this profile has no children and therefore also has no defined Exit State. If a Batch Default has not been set, the state with the lowest priority preference is used which (if present) has not been set as being Unreachable.

Dependency Default One or more final Exit States can be labelled as the Dependency Default. If a dependency between batches and jobs is created by selecting the DEFAULT state, the Exit States of a profile that are marked as being the Dependency Default fulfil the dependency condition regarding the Exit State.



Preference The *Preference* buttons are used to set the preference of the individual Exit States. The rows are shifted up or down by clicking the buttons. A higher position indicates a higher preference.

To determine the resultant Exit State of a Submitted Entity, the Exit State with the highest preference is selected from the list of Exit States of the children and its own Exit State. In doing so, an Exit State is only taken into account if it is present in the Exit State Profile of the parent .

Example: A batch has three subordinate child jobs. The batch is to display an error when at least one of its children returns an error. This means that the FAILURE state must precede the SUCCESS state.

The SUCCESS state will usually have the lowest priority, warnings have a higher priority, while Error states have the highest priority.

Editor

If an empty batch (one without any children) becomes FINAL, then there is no Exit State from which the FINAL Exit State of the batch can be determined. In this case, the Exit State from the Exit State Profile with the lowest preference is used. This is usually the SUCCESS state. The system attempts to use a FINAL state which is not the Unreachable State if one actually exists.

5 Resource State Definitions

5.1 View

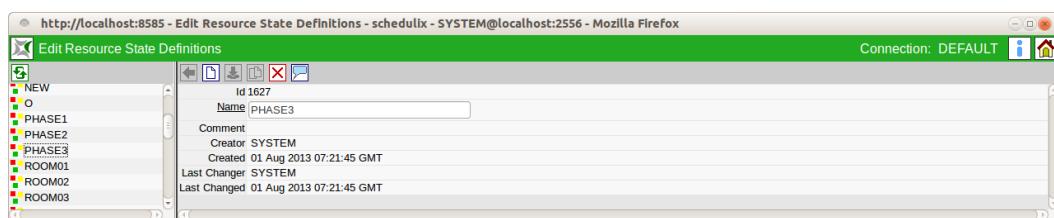


Figure 5.1: Resource State Definitions

5.2 Concept

5.2.1 In short

Resource State Definitions are used to create names for Resource States.

5.2.2 Detailed description

Resource State Definitions are logical names for Resource States. Each Synchronizing Resource can take on several states that can be set when a job has been completed.

5.3 Editor

The properties of Resource State Definitions are maintained in this tab.

Resource State Definitions can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

The fields have the following meanings:

Id System-wide unique number for identifying the object.

Name Unique name of the Resource State Definition. This name can be freely chosen.

6 Resource State Profiles

6.1 View

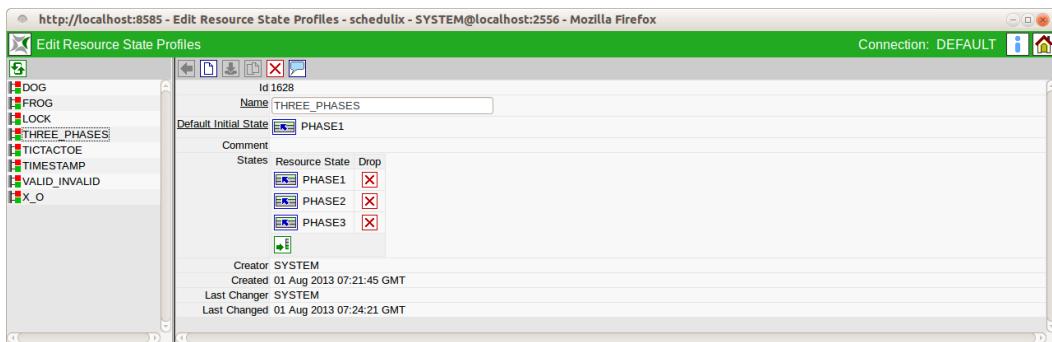


Figure 6.1: Resource State Profiles

6.2 Concept

6.2.1 In short

A Resource State Profile describes which states a resource can take on.

6.2.2 Detailed description

Each Synchronizing Resource can take on different Resource States. The set of all valid Resource States for a resource is described in a Resource State Profile. The initial state of a resource is also defined in this profile, whereby the state does not necessarily have to be present in the list of states.

6.3 Editor

The Resource States that belong to a Resource State Profile can be defined in this editor.

Resource State Profiles can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

Editor

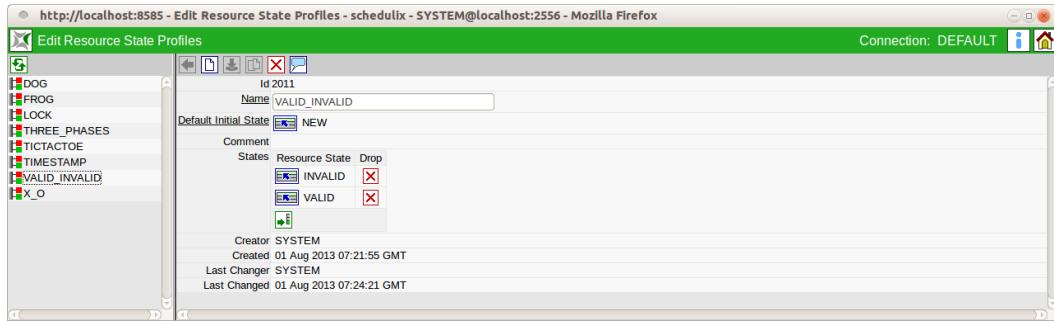


Figure 6.2: Resource State Profile Editor

The fields shown above have the following meanings:

Name The *Name* field is used for assigning a unique name to the Resource State Profile.

Default Initial State The *Default Initial State* field defines the initial state of the resource. This Resource State does not have to be present in the list of valid Resource States.

Resource State The valid Resource States are shown in the *Resource State* column in the *States* table.

6.3.1 Example

A database table for a data mart is shown in the scheduling system as a Synchronizing Resource with a State Model. The state chart looks like the diagram in Figure 6.3

The definition of a suitable Resource State Profile is shown in the screenshot above.

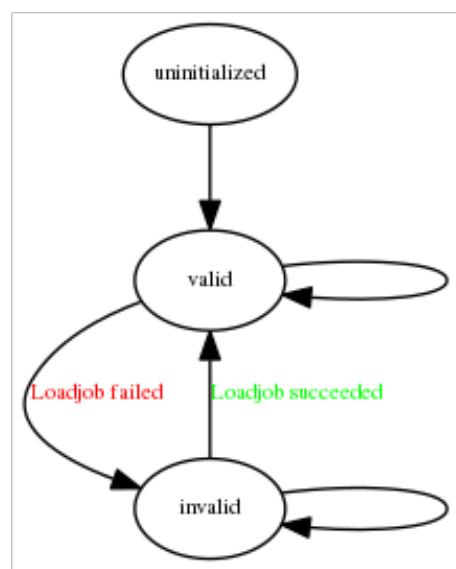


Figure 6.3: State chart for a resource

7 Resource State Mappings

7.1 View

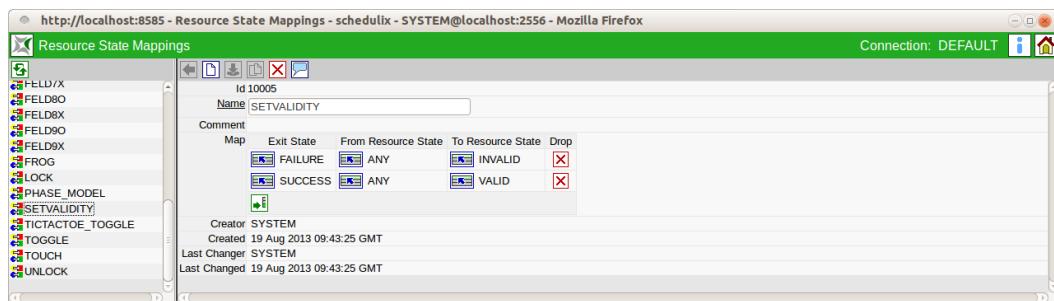


Figure 7.1: Resource State Mappings

7.2 Concept

7.2.1 In short

Resource State Mappings describe the state transition of a resource with certain Exit States.

7.2.2 Detailed description

Resources can change their Resource State after completion of a job depending upon the job's Exit State. Resource State Mappings describe these state transitions. Respectively one Exit State and one Resource State determine the new state of the resource.

For example, a "TABLE" resource can take on the state "VALID" or "INVALID". If the loading process belonging to the table is completed successfully (SUCCESS), the Resource State "VALID" is to be set. In the event of an error (FAILURE), the table is to be marked as being "INVALID". The corresponding definition of the Resource State Mapping would look like in Figure 7.2.

If no mapping exists for the combination of Exit State and From Resource State, the status of the resource remains unchanged. This means that it is not necessary to define a mapping from SUCCESS/VALID to VALID.

Editor

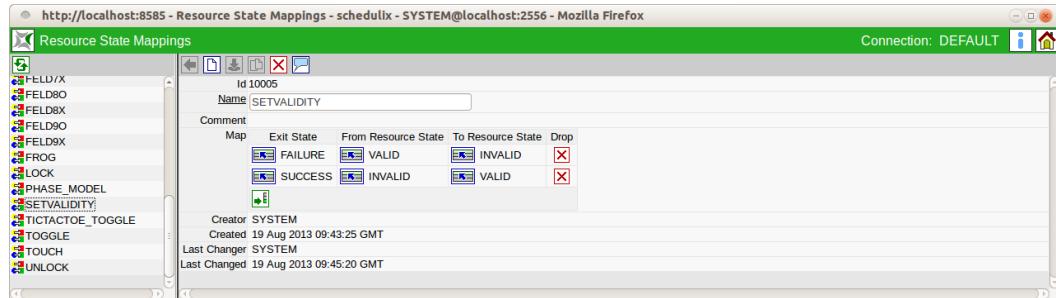


Figure 7.2: Example for a Resource State Mapping

If a state transition is to take place regardless of a Resource State, the value "ANY" needs to be selected in the line "From Resource State".

When implementing the status for a resource, the time of implementation is saved. Reference can be made to this time in the resource requirements, and so it may be practical to define the transition from "VALID" to "VALID".

7.3 Editor

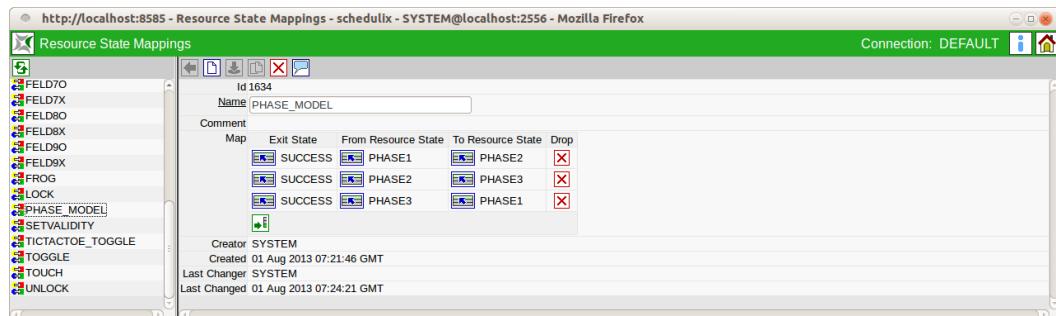


Figure 7.3: Resource State Mapping Editor

The Exit States are assigned to the resource state transitions in this editor.

Resource State Mappings can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

The fields have the following meanings:

- Id** System-wide unique number for identifying the object.

Name Unique name of the Resource State Mapping. This name can be freely chosen.

Exit State The Exit State of the job is given in the field *Exit State*.

From Resource State Either a specific Resource State or ANY is given in the field *From Resource State*. If the current state in combination with the Exit State is explicitly specified in the table, this rule is applied to determine the new state when implementing it. Otherwise, the ANY rule (if present) is used.

To Resource State The field *To Resource State* is the resultant state of the resource.

8 Named Resources

8.1 View




Figure 8.1: Named Resources

8.2 Concept

8.2.1 In short

This window is used to create and administer Named Resources. A Named Resource is the definition of a class of resources.

8.2.2 Detailed description

A Named Resource defines a class of resources. Named Resources can be instantiated as resources within a scope, folder or in a Submitted Entity.

Named Resources belong to a group. The right to use (i.e. instantiation and requirement), edit or delete the Named Resource is primarily reserved for users belonging to that group. These privileges can be granted to other groups.

Named Resources are saved in a hierarchy of categories. Categories can be nested as required, but Named Resources cannot be nested. This categorisation serves solely to facilitate the administration of Named Resources, and otherwise has no influence on how the system functions.

8.3 Editor

Named Resources and categories are maintained in the Editor menu. When an entry is selected in the Navigator, the details for this Named Resource or category are displayed here. New Named Resources and categories are created here as well.

8.3.1 Properties tab for categories

This tab is used for maintaining the category properties. It looks like this:




Figure 8.2: Categories; Properties tab

The Properties tab for a category and Named Resource has the following fields:

ID The ID is used to unequivocally identify the category or Named Resource Definition. It is automatically assigned by the system and is unique system-wide.

Path The path is the entire hierarchy of the viewed entry. The single hierarchy levels are separated by a period.

Name Any name can be chosen for the category or Named Resource. It must, however, be unique within the parent category.

Group The group to which the category is assigned can be selected from the drop-down list.

Cascade Set Group If this box is checked, the group is also set in all the Named Resources and categories below this category.

Comment The *Comment* gives a more detailed explanation about the object.

8.3.2 Content tab

The Content tab is only displayed for categories and looks like this:

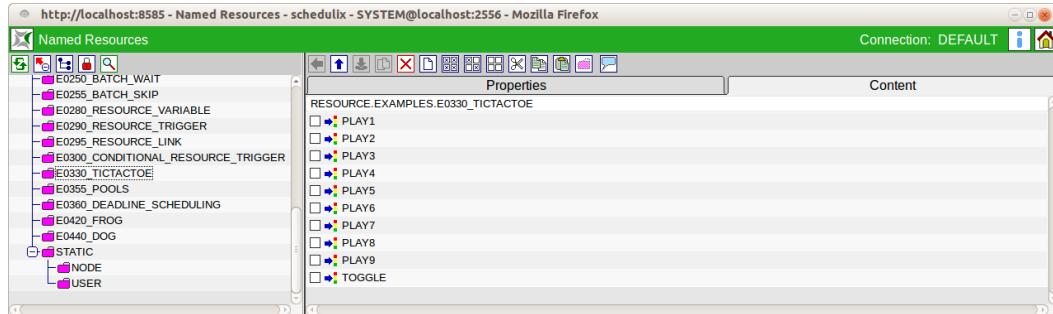


Figure 8.3: Categories; Content tab

This tab contains a list of all the Named Resources in the selected category. The following values are shown in the list:

Name of the Named Resource The name of the Named Resource is shown here. Selecting the Named Resource's name in this dialog opens the selection window and display of the "Details" data for this Named Resource. In this tab, you can move resources using the standard cut, copy and paste operations.

8.3.3 Properties tab for Named Resource Definitions

The Properties tab for a Named Resource Definition looks like this:




Figure 8.4: Named Resources; Properties tab

When you define a Named Resource, you have to fill in some other fields in addition to those mentioned above.

Usage The *Usage* field specifies the Named Resource type.

List of options for the usage parameters:

1. *Category*

Categories behave like folders and can be used to arrange the Named Resources in a clearly organised hierarchy.

2. *Static*

Static resources describe workflow environments such as those offered by installed software packages or user environments. An instance of this Named Resource must exist and be available in order to be able to start a Submitted Entity that conditionally requires this resource.

3. *System*

The Named Resource is a resource that maps a system parameter. This can be, for example, a CPU unit, main memory unit or a unit of disk space. This resource also has to be quantified when it is used, i.e. the creator of a job needs to qualify and quantify the requirements that have to be met by the system resources for this object. When defining a job he will have to specify, for example, that 3 CPU units, 3 units of 512 MB main memory and 10 units of 1 GB disk space are required.

4. *Synchronizing*

Synchronizing Resources are special resources that are used to synchronise concurrently running jobs. Whether a resource requirement can be fulfilled and the resource can be allocated can be made dependable upon the current status of the resource, the last status update time, and the contending allocation of the resource by other tasks. The requirement can be quantified analogue to System Resources.

Resource State Profile A Resource State Profile that defines the status that the Synchronizing Resource can take on can be assigned to that Synchronizing Resource. No status or timestamp-related requirements can be defined if a Resource State Profile has not been specified. Further information can be found in Chapters [5](#) and [6](#).

8.3.4 Parameters tab

Additional information about a resource that can be evaluated by jobs or Resource Triggers can be saved in the Parameters tab.

The Parameters tab looks like in Figure [8.5](#).

The "Parameter Details" tab is opened by clicking the parameter name.

Editor

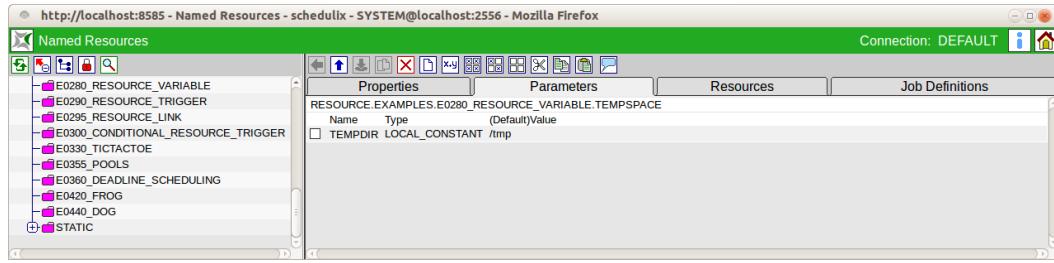


Figure 8.5: Named Resources; Parameters tab

8.3.4.1 Parameter Details tab




Figure 8.6: Named Resources; Parameter Details

The fields in the "Parameter Details" tab have the following meanings:

Parameter Name Name of the parameter.

Parameter Type Type of the parameter.

The following options are available:

- Constant: The Constant has a fixed value and applies for all resources.
- Local Constant: The Local Constant has a fixed value which may vary from resource to resource.

Default Value With the Default Value, we differentiate between Constants and Local Constants. It is the value of the parameter for Constants and the default value for Local Constants.

8.3.4.2 Standard parameters

Just as there are standard parameters for jobs and batches, standard parameters are also available for resources.

To be able to reference them, a slight trick has to be used analogue to the situation with batches and jobs. If you want to use the contents of the standard variable "STATE" for a processing operation, for instance, you have to create another variable (a constant) with the value "\$STATE". Due to the recursive parameter resolution, the constant is given the value of the parameter "STATE".

The following standard parameters are available:

STATE The state of the resource. This content is empty for non-Synchronizing Resources.

AMOUNT The total available amount of the resource. This value is empty for static resources.

FREE_AMOUNT The available free amount of the resource. This value is empty for static resources.

REQUESTABLE_AMOUNT The maximum amount of the resource that can be requested. This value is empty for static resources.

8.3.5 Resources tab

When you select a Named Resource in the navigation, its instances are displayed in the Resources tab. All the scopes, folders, Submitted Entities, Scheduling Entities or job servers that offer instances of this Named Resource are shown here.

The Resources tab looks like this:

Scope	State	Requestable Amount	Amount	Free Amount	Online
SERVER	1	2	2	2	true
SERVER	1	2	2	2	true

Figure 8.7: Named Resources; instances

The instantiations of the Named Resource are described by the following fields:

Editor

Scope The names of the scopes, Submitted Entities, Scheduling Entities or folders that offer instances (resources) of the respective Named Resource are shown here.

State If the selected Named Resource has a **Resource State Profile**, the current state of the resource is displayed here. This is only possible with Synchronizing Resources.

Requestable Amount The maximum amount of resources that can be requested by a job.

Free Amount The number of free instances of this resource is shown here. A bar graphically indicates the current allocation rate.

Amount The amount is the number of available units of the resource.

Online The availability status of the resource.

8.3.6 Job Definitions tab

The Job Definitions tab shows all the Job Definitions that request this Named Resource. This tab is for information purposes only; no changes can be made here.

The tab looks like this:

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "Named Resources - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area displays the "Job Definitions" tab for a named resource. On the left, there is a tree view under "RESOURCE EXAMPLES" showing various job definitions: E0075_LOADCONTROL, E0076_LOADCONTROL_KEEP, E0077_DETAILED_LOADCONT, E0080_LOADBALANCE, E0115_DYNAMIC_SERVER_SI, E0115_SELECTSERVER_DYN, E0120_SYNC_SIMPLE_LOCK, E0130_SYNC_RESOURCE_ST, and E0140_TIMEOUT. To the right of the tree view are four tabs: "Properties", "Parameters", "Resources", and "Job Definitions". The "Job Definitions" tab is active and contains a table with one row. The table has columns for Job Name, Amount, Keep, Sticky, Mapping, Expire, Unit, Lockmode, and States. The single entry is "CHILD" with "Amount" set to 1 and "Keep" set to "NOKEEP".

Figure 8.8: Named Resources; Job Definitions tab

Named Resource selector

The list of jobs is described in the following fields:

Job Name The names of the Submitted Entities that require the Named Resource are shown here.

Clicking the name opens a [Job Editor window](#).

Amount The amount of the resource that is required by the job.

Keep The value of the Keep parameter for the resource request from the job.

Sticky The value of the Sticky Flag for the resource request from the job.

Mapping If a Resource State Mapping was specified in the resource request, it is displayed here.

Expire The Expire value specifies the maximum length of time that may elapse since the last state transition of the resource for it to be regarded as being occupiable. A negative Expire value means that a resource must be at least as "old" as given here.

Lock mode The Lock mode describes the mode for accessing this resource (exclusive, shared, etc.).

States All the Resource States that are accepted by this job are shown here. Multiple states that are acceptable for this job are separated by commas.

8.4 Named Resource selector

This navigator is used for selecting single or multiple Named Resource Definitions and is called by other dialogs like [Footprint](#) as a search and selection mask. Which resources are displayed varies depending upon the call dialog. If the selector is called from the Footprint editor, only the system resources are displayed. If the selector is called from the Environment dialog, only static resources are displayed.

Named Resource selector




Figure 8.9: Named Resources; selector

9 Environments

9.1 View

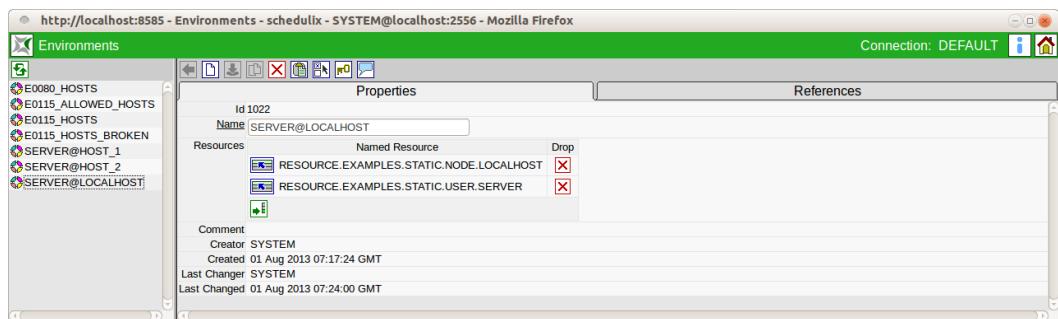


Figure 9.1: Environments

9.2 Concept

9.2.1 In short

The "Environment" dialog is used for administering environment definitions. An environment collates multiple requirements for static resources under one name.

9.2.2 Detailed description

Environments serve two purposes. On the one hand, they simplify the maintenance of resource requirements since they no longer have to be individually added in the Job Definition. All you have to do is to specify the respective environment. On the other hand, environments are used to determine which users or user groups are permitted to use specific runtime environments.

The Environment Definition therefore states which runtime environment is required by a job. For example, which host is to run a job or which users and programs must be available can be controlled with an environment.

The environment is a mandatory parameter in a Job Definition.

Editor

Environments can also be entered as a Folder Environment for a folder. This means that all the Job Definitions below such a folder "inherit" the resource requirements from the Folder Environment.

Where separate job servers are in use for Development and Production, for example, an environment can be used for the job which allows it to be run on both job servers. Creating two folders for Development and Testing (to which a "Development" and "Production" environment are respectively assigned whose resource requirements can only be fulfilled by development or production job servers) allows the Development and Production environments to be easily kept separate. If such a job is located somewhere below the "Development" folder, the combination of resource requirements for Job Environment and Folder Environment ensures that the correct Development job server is selected. If the job is moved or copied to the "Production" folder and then executed, the Production job server is selected automatically. The Job Definition does not have to be modified.

9.3 Navigator

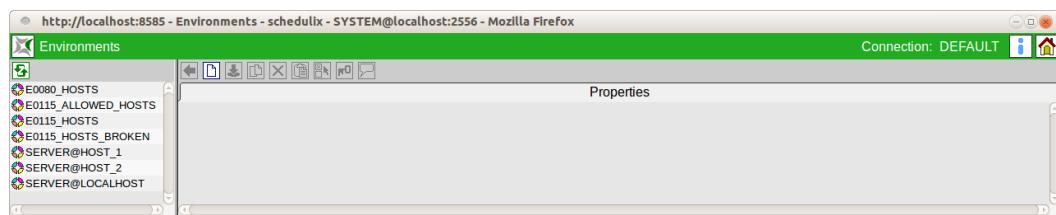


Figure 9.2: Environment navigator

The "Environment" dialog navigation window shows all the existing environments. If the logged in user is a member of the "ADMIN" group, all the environments are displayed. If this is not the case, then only those environments that can be used by the user who is logged in are shown. This means that the logged-in user must be a member of a group that holds the "use" privilege for that particular environment.

9.4 Editor

9.4.1 Properties tab

This tab is used for maintaining the environment properties. Environments can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

It looks like this:

The Properties tab for environments has the following fields:

Editor




Figure 9.3: Environment properties

ID The ID is assigned automatically and is used for unequivocal, system-wide identification of the object.

Name The name of the environment. Any name can be chosen, although it must be unique within the environment.

The "Resources" list shows all the resources that belong to the environment. These are exclusively static resources. The *Named Resource* fields have the following meanings:

Named Resource The name of the Named Resource is shown here.

Condition A condition is entered in the *Condition* field that must be fulfilled to ensure that the resource is recognised as being valid. The condition is evaluated in the context of the requesting job.

Drop The *Drop* button can be used to delete a specific row.

9.4.2 References tab




Figure 9.4: Environment references

Editor

This tab displays any references to the environment. It looks like in Figure 9.4.

The References tab shows a tree view that only contains the parent folder of objects that reference the environment.

Folder names are usually shown in italics unless the folder references the environment as a Folder Environment.

Clicking the name of a non-recursively displayed folder or a job opens an editor window for the folder or job.

10 Footprints

10.1 View




Figure 10.1: Footprints

10.2 Concept

10.2.1 In short

This dialog is used for creating a footprint definition. A footprint collates a group of System Resource requirements under one name.

10.2.2 Detailed description

Collating the resource requirements for footprints makes it easier to maintain the requirements for jobs in the runtime environment because they do not have to be individually specified.

The footprint can be overridden for each job. This is done by using an explicit requirement for a resource (contained in the footprint) with a deviating amount or by changing the setting for the **Keep parameter**. A resource requirement cannot be removed, i.e. all the resources defined in the footprint are requested. The amount can be reduced to 0, although resource must still be present.

Example:

A System Resource CPU_UNIT with the amount 4 (CPU units) is defined in the footprint. The footprint can be overridden by adding a requirement for the same

Editor

System Resource (CPU_UNIT) within the Job Definition (**Resources tab**) for an amount of 2 (CPU units). It is not possible to completely remove the resource CPU_UNIT, however.

10.3 Editor

10.3.1 Properties tab

This tab is used for maintaining the footprint properties.

Footprints can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

It looks like this:




Figure 10.2: Footprint properties

The Properties tab for footprints has the following fields:

ID The ID is assigned automatically and is used for unequivocal, system-wide identification of the footprint.

Name The name of the footprint. This can be freely chosen.

A list of the resources contained in the footprints can be entered in the "Resources" list.

The fields have the following meanings:

Named Resource The name of the Named Resource is shown here. It is selected with the *Chooser* button.

Amount This is the amount of the resource required by a job.

Keep The Keep parameter determines whether the Named Resource is retained after a job has been completed or if it can be released again. The following selection options are available:

1. *No Keep*

The Named Resource is released after completion of the job. Whether the job has been completed successfully or aborted due to an error is irrelevant as far as the release is concerned.

2. *Keep*

The resource is not released until the job has attained a final state. The resource is retained in the event of an error (Restartable State).

3. *Keep Final*

The resource is not released until the job and all its children have attained a final state.

10.3.2 References tab

This tab displays the references to the footprint. It looks like this:




Figure 10.3: Footprint references

The References tab shows a tree view that only contains the folder below which there are jobs that reference the footprint.

Folders are displayed in italics.

Clicking the name of a job opens a window where it can be edited.

11 Job servers and resources

11.1 View

The screenshot shows a web-based interface for managing job servers and resources. The left sidebar lists job servers under 'GLOBAL' and 'EXAMPLES'. Under 'GLOBAL', there are 'HOST_1' and 'HOST_2', each with a 'SERVER' node. Under 'EXAMPLES', there is also a 'SERVER' node. The main panel displays a table of resources. The columns are 'Usage', 'State', 'Timestamp', 'MaxReq', 'Amount', 'Load', 'Free', 'Online', 'Scope', 'Resource', 'Link', and 'Scope'. One resource entry is visible: 'E0076_LOADCONTROL' with a state of 'SYSTEM', amount of '4', and load of '1'. The 'Scope' column shows 'true' for this entry. The 'Resource' column lists various resource types like 'LOADCONTROL', 'SCOPE_UNITS', 'LOADCONTROL_KEEP', 'DETAILED_LOADCONTROL', 'SYNC_SIMPLE_LOCK', 'SIMPLE_LOCK', 'SYNC_RESOURCE_STATE', 'TIMEOUT', 'EXPIRATION', 'CRITICAL_REGION_STICKY', 'CRITICAL_REGION_STATE', 'CRITICAL_REGION_KEEP_FINAL', 'KEEP_FINAL', 'RESOURCE_VARIABLE', 'RESOURCE_TRIGGER', 'RESOURCE_LINK', 'CONDITIONAL_RESOURCE_TRIGGER', 'POOLS', 'DEADLINE_SCHEDULING', and 'STATIC'.

Figure 11.1: Job servers and resources

11.2 Concept

11.2.1 In short

A job server is a schedulix system process that has to run on all the computers that have been defined as executive units of the schedulix Scheduling System. If a job server is assigned to perform a task by the schedulix Server, the job server executes the script or program specified in the job definition in the required environment and with the necessary parameters. The job server monitors the program and sends the return state of the script or program back to the schedulix Server.

11.2.2 Detailed description

The physical distribution (i.e. which job servers run on which computers and which resources are made available by these job servers) can be defined and maintained

in the "Job Servers and Resources" dialog.

These job servers must be set up and started at system level.

A differentiation is made between two types of objects in this dialog. On the one hand there are so-called scopes, which can serve as a container for the respective job server or child scopes.

On the other hand, those job servers that reflect the executive system processes are administered in this dialog.

The hierarchy can be subdivided as required and is largely dependent upon the complexity of the system environment and administration structure (e.g. in the data centre).

For example, it is possible to define each physical host as a scope, and all the workflow environments required on this computer can be defined as job servers below this scope. If the system landscape is more complex, it is practical to use multiple levels for departments, system architectures (Unix, Windows) or similar entities.

The scopes and job servers can be assigned to resources at every level of this hierarchy. These resources can then be used by all the children job servers.

Example:

A value of 5 is entered for the System Resource CPU_UNITS at host system level. If 2 job servers have been configured on the computer, a total of 5 CPU units can be shared between them. If a job server occupies 3 CPU units when executing a job, only 2 CPU units remain available to the other job server. However, this behaviour is not additive. If the resource CPU_UNITS is also created on one of the aforementioned job servers (for example with a value of 3), only these 3 CPU_UNITS will be available to that job server. Accordingly, 5 CPU_UNITS from the parent scope will be available to the other job server.

11.3 Navigator

The Navigator provides a hierarchical structural view in which scopes are defined as folders. The job servers are located within the scopes.

11.4 Editor

11.4.1 Properties tab

The Properties tab is used for maintaining the properties of scopes and job servers.

"Properties" tab for scopes is shown in Figure 11.3.

"Properties" tab for job servers is shown in Figure 11.4.

The fields in the "Properties" tab have the following meanings:

ID The ID is used for unequivocal, system-wide identification of the object.

Path This is the complete path to the object within the hierarchy.

Editor




Figure 11.2: Job servers and Scope Navigator

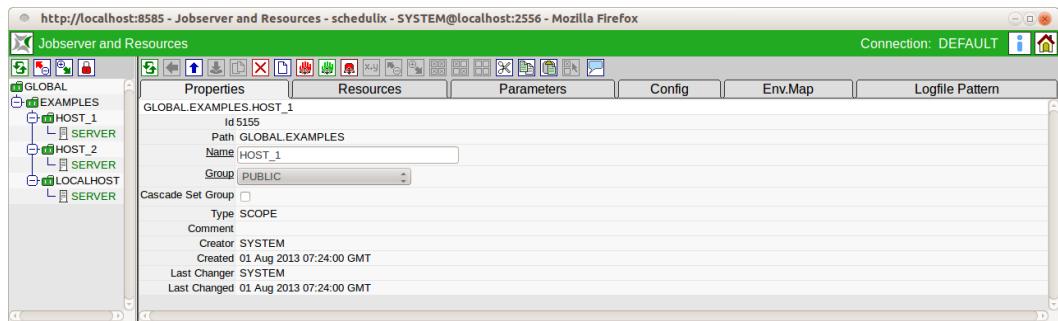


Figure 11.3: Scope properties

Name This is the name of the scope or the job server. The name can be freely chosen, but it must be unique within the hierarchy level.

Group The group can be selected from a drop-down list. It designates the scope's owner group.

Type The *Type* field shows the type of the selected object. The type is selected when creating a new server or scope and cannot subsequently be changed.

OPTIONS FOR TYPE

1. *Scope*

Editor

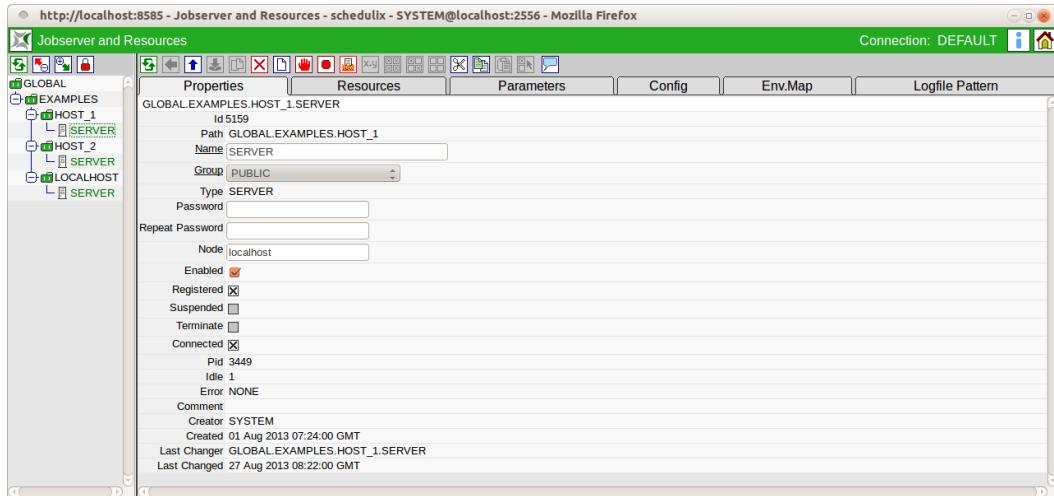


Figure 11.4: Job server properties

A scope is a category that can be grouped hierarchically organised below the job server. This could be a host computer, for example, or, at an even higher level, all the computers in a department or all the Windows-based PCs. The structure and depth of the hierarchy can be modified depending upon the system landscape and should reflect this.

2. Server

This object is a physical job server. The physical process has to log in to the schedulix system under the selected name and in doing so set up the connection to this job server definition.

Password The *password* used by the job server to log on to the schedulix Server. Each job server requires a name and password to log on to the schedulix Server. A job can only be handed over to the job server to be executed with a valid login. The password is hidden on the screen and must be identical to the field *Repeat password*.

Since the password is not displayed in the dialog, it is not clear whether a password has already been entered or not. The password does not have to be specified again when changing other properties. If the job server is logged in, it is automatically notified about any password changes.

Repeat Password The *Repeat Password* must be the same as the password you entered in the *Password* field. It is necessary to enter it twice to detect any accidental typing mistakes.

Node The *Node* specifies the computer on which the job server is running. This field has a purely documentary character.

Enabled If the job server has set the Enabled flag, the job server process can log onto the schedulix Server. The login will fail if the flag has not been set.

Suspended If the Suspended flag has been set, although the job server process can log on it will not be given any jobs to run. State changes for previously allocated jobs are accepted by the schedulix Server.

Terminate The Terminate flag indicates whether the job server is to terminate or not. If this flag has been set, the job server will receive this message at the next communication step and behave accordingly.

It may take a few seconds before the job server terminates. The termination can be verified by clicking the *Refresh button*. Once the job server has finished, the value displayed in the *Connected* field changes to "FALSE".

Connected The *Connected* field shows the connection status of an external job server process. If the job server process was able to successfully log on to the schedulix system, this value is set to "TRUE".

PID The PID is the process identification number of the job server process on the respective host system. It is communicated to the schedulix system during the registration.

The PID and the *Node* field allow the process to be located in the operating system running on the respective server.

Error An error message is displayed here if an error occurred while the job server process was running.

11.4.2 Resources tab

The Resources tab looks like this:

All the resources offered by the current scope or job server are shown in the Resources tab. If any resources (especially of the type *Synchronizing*) are to be awarded system-wide, this can be done in the "GLOBAL" scope.

If multiple instances of a resources are entered in different scopes or job servers, these are respectively a separate instance of that resource.

Example:

If a Synchronizing Resource A is entered in Scope X and Scope Y, these two instances of Resource A are mutually independent (in X and Y). This means that a possible synchronisation will only take place among all the jobs in Scope X and

Editor

The screenshot shows the 'Jobserver and Resources' editor interface. On the left, there is a tree view of resources under 'GLOBAL' and 'EXAMPLES'. Under 'GLOBAL', there are 'HOST_1' and 'HOST_2', each containing a 'SERVER' node. Under 'EXAMPLES', there is a 'LOCALHOST' node containing a 'SERVER' node. The right side of the screen displays a table for the selected resource, which is 'GLOBAL.EXAMPLES.HOST_1.SERVER'. The table has columns for 'Properties', 'Resources', 'Parameters', 'Config', 'Env.Map', and 'Logfile Pattern'. The 'Resources' tab is active, showing a list of resources with their details. One resource, 'E0115_DYNAMIC_SERVER_SELECTION', is expanded, showing its children: 'ALLOWED_HOST' (STATIC), 'SELECT' (STATIC), and 'LOCK' (SYNCHRONIZING). The 'LOCK' row has columns for 'Usage' (SYNCHRONIZING), 'State' (INFINITE), 'Timestamp' (INFINITE), 'MaxReq' (INFINITE), 'Amount' (INFINITE), 'Load' (true), 'Free' (true), 'Online' (true), 'Scope' (EXAMPLES), 'Resource' (true), and 'Link Scope' (true).

Figure 11.5: Job server and scope resources

Scope Y respectively. If a global synchronisation is to take place here, Resource A must be defined in the GLOBAL scope instead.

The resources are displayed as a map of the hierarchy from the navigation in the **Named Resources** dialog instead together with additional details. The first field in the list shows the name of the Named Resource or category. If this is a category, a folder icon stands here as well. The category can be opened and closed by clicking this icon. If it is a Named Resource, clicking the name opens the Resource Details tab.

The other fields are explained in the next section.

11.4.2.1 Resource Details tab

The Resource Details tab is displayed as soon as a resource is selected. The instance information for this Named Resource is entered and maintained in this tab. Here, an instance is the expression of the Named Resource in the selected scope or job server. All instance parameters can be entered and maintained here.

The Resource Details tab looks like this:

The fields in the "Resource Details" tab have the following meanings:

Usage The *Usage* field specifies the "Resource" type. More details about resource types can be found in the **Named Resource** dialog.

Resource State Profile This field is only visible if the resource usage has been set to "Synchronizing".

This is the **Resource State Profile** assigned to the resource.

Editor

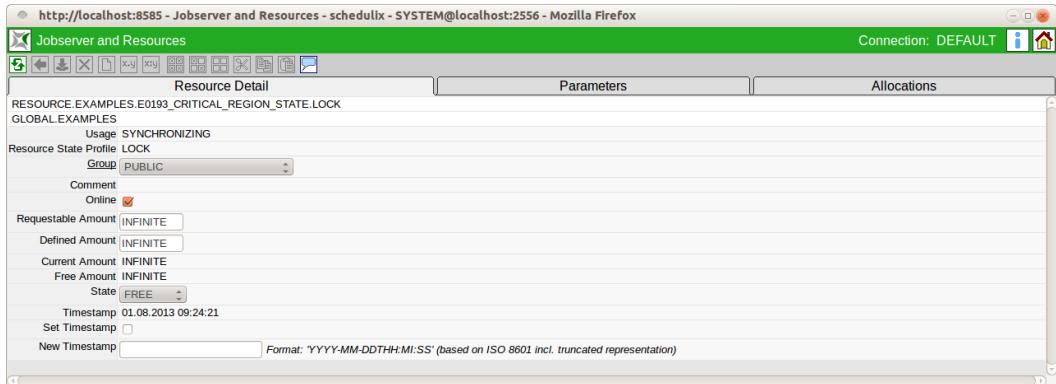


Figure 11.6: Resource Details

Online *Online* can be used to switch a resource in this scope or job server online or offline. An offline resource is temporarily unavailable. Example:

A scheduled computer or database downtime is pending. If the computer or database has been mapped as a resource in the system, the resource can be set as being offline as soon as it is physically no longer available. All jobs that require this resource are automatically no longer executed. If the resource is to be defined in another scope as well (as a separate instance), all the jobs that would require the resource would deviate to this scope or job server.

An offline resource does *not* trigger the error message "Job cannot run in any scope because of resource shortage".

Requestable Amount The *Requestable amount* is the maximum amount that can be requested.

Defined Amount This field is only visible if the resource usage has been set to "Synchronizing" or "System".

The *Defined Amount* states the current number of instances of the Named Resource for this scope or job server. If the number of instances for this scope changes, you can edit this field accordingly.

Example:

This may become necessary if the hardware has been modified. If the resource represents a CPU unit and 2 CPU units are installed in the current host, the amount is 2. If 2 more CPU units are then installed, this value must be increased to 4.

Current Amount The *Current Amount* and *Defined Amount* are identical.

Free Amount This field is only visible if the resource usage has been set to "Synchronizing" or "System".

The *Free Amount* designates the total number of instances of a resource in the selected scope or job server that have not yet been allocated to jobs.

State The *State* field is only visible if the resource usage has been set to "Synchronizing" and a **Resource State Profile** has been assigned. It designates the current status of the resource in this scope or job server. The state can be set or changed with this value. The drop-down list contains all the valid Resource States of the Resource State Profile, which can be selected here. This may be necessary when manually fixing an error to restore a resource to a state that allows further processing by follow-on jobs.

If you manually change this value, the value in the *State* field in the Resources tab is not automatically refreshed. Click the *Refresh* button if this is required and necessary.

Timestamp This field is only visible if the resource usage has been set to "Synchronizing" and a **Resource State Profile** has been assigned. The *Timestamp* indicates the time of the last change made to a Resource State.

Timestamps play a role if an expiration interval has been defined in an expiry object. If this is the case, the timestamp must not be older than the specified interval. More details about expiration times can be found in Chapter [12.5.8.1](#).

If the completion of a job causes the Resource State to change, the timestamp is automatically updated by the schedulix Server.

Set Timestamp This field is only visible if the resource usage has been set to "Synchronizing" and a **Resource State Profile** has been assigned. The *Set Timestamp* flag is a so-called action flag. If you set this flag and then click the *Record* button, the time shown in the *Timestamp* field is set to the current time.

Example:

This may be necessary when manual intervention is required to fix an error and the state has been changed manually in the *State* field .

11.4.2.2 Parameters tab

Additional information about a resource that can be evaluated by jobs or Resource Triggers can be saved in the Parameters tab. The parameters are defined in the Named Resources.

The columns in the list shown above have the following meanings:

Name The "Name" column shows the name of the parameter.

Editor

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "Jobserver and Resources - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area has tabs for "Resource Detail", "Parameters", and "Allocations". Under "Resource Detail", it shows a table for "GLOBAL.EXAMPLES.HOST_1.SERVER" with one row: "HOSTNAME LOCAL_CONSTANT" with value "HOST_1".

Figure 11.7: Resource parameters

Type This is the parameter type. The following options are available:

- Constant: The constant has a fixed value and applies for all resources.
- Local Constant: The Local Constant has a fixed value which may vary from resource to resource.

The values of the Local Constants can be changed here.

Default The *Default* field indicates whether the displayed value is the default value. The flag has to be set to restore the value to the default value.

Value This is the value of the parameter.

11.4.2.3 Allocations tab

The Allocations tab provides information about which tasks have been currently allocated the resource and which tasks require resources but cannot obtain them because they have the incorrect state or due to a lock.

The Sheet tab looks like this:

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "Jobserver and Resources - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area has tabs for "Resource Detail", "Parameters", and "Allocations". Under "Allocations", it shows a table for "GLOBAL.EXAMPLES.LOCALHOST.SERVER" with six rows of data:

#	Job	JobId	MasterId	Type	Amount	Keep	Sticky	Lockmode	Mapping	P	EP
1	CHILD[1]	21010	21005	ALLOCATION	1	NOKEEP	false	N	NONE	50	50
2	CHILD[2]	21015	21005	ALLOCATION	1	NOKEEP	false	N	NONE	50	50
3	CHILD[3]	21020	21005	ALLOCATION	1	NOKEEP	false	N	NONE	50	50
4	CHILD[4]	21025	21005	ALLOCATION	1	NOKEEP	false	N	NONE	50	50
5	CHILD[5]	21030	21005	BLOCKED	1	NOKEEP	false	N	NONE	50	50
6	CHILD[6]	21035	21005	BLOCKED	1	NOKEEP	false	N	NONE	50	50

Figure 11.8: Resource allocations

The columns in the list shown above have the following meanings:

Job Depending on the setting, either the name of the job or the full path of the folder hierarchy for this job is shown here.

Job Id This is the ID of the job instance that was started with either a direct submit in the **Submit Jobs** dialog or a submit of the Master Batch or job.

Master Id This is the ID of the job instance or batch that was started as a Master Job in the **Submit Jobs** dialog and contains the current job as a child. If the job is itself a Master Job, then its own Job Id is shown here.

Type The *Type* refers to the method used by the current job instance to access or attempt to access the resource. The following options are available:

1. Requested

The job asks the server if the resource is available.

2. Reserved

The resource is reserved by the job. This has the same effect as Allocated, except that reserved resources can be released again when the job is started. Allocated resources are released at the earliest when the job has finished.

3. Allocated

"Allocated" indicates that the job is currently running and accessing the resource. Other jobs that want to access this resource can, depending upon the type of access lock, be allocated this resource as well or be blocked.

4. Blocked

If a job is labelled as being "Blocked", it cannot be run at the moment because it is waiting for access to the selected resource but is not able to obtain it due to a non-matching lock mode or other criteria.

The reason why the resource is not available is indicated by the red lettering.

As soon as the resource becomes available again, the type changes from "Blocked" to "Requested".

5. Available

The resource is available.

6. Ignored

The user has told the system to ignore the resource request.

7. Master Reservation

A Master Reservation is a reservation of Sticky Resources.

Amount The *Amount* is the number of resource instances allocated to or requested by the current job. If the value of the amount exceeds the currently available number that can be maintained in the *Amount* field in the Resource Details tab and there is no alternative scope available with a sufficient number, the job cannot be executed.

Keep The *Keep* parameter defines the Keep state of the current resource. More information about the Keep state can be found in Chapter [10.3.1](#).

Sticky The *Sticky* parameter defines whether the job is given the current "Sticky" resource or not. More information about the Sticky Flag can be found in Chapter [12.5.8.1](#).

Lock mode The *Lock mode* defines the access mode being used in allocating the resource to the current job. The following access modes are available:

1. Exclusive (X)

No other job has access to this resource.

Example: When loading a database table.

2. Shared Exclusive (SX)

These access operations are mutually compatible, but they are not compatible with shared access operations.

Example: An application performs numerous small write operations on a database table.

3. Shared (S)

Shared access operations are mutually compatible.

Example: Evaluations of complete database tables or large parts of them.

4. Shared Compatible (SC)

These access operations are both mutually compatible as well as being compatible with S and SX access operations.

Example: An application performs numerous, brief read transactions on a database table. Such applications do not impede one another. Small write transactions or large-scale evaluations are not an impediment either.

5. NoLock (N)

The NoLock lock mode indicates that the resource is neither being locked nor is a lock to be heeded. A job that locks a resource with NoLock can always access it regardless of any other lock options.

Here is a compatibility matrix summarising all the lock options and their interactions.

(+) means that the lock modes are mutually compatible. Two jobs that request the two lock modes can run concurrently.

(-) means the lock modes exclude one another. A job stays blocked until the other job releases the resource again.

	Exclusive (X)	Shared (S)	Shared Exclusive (SX)	Shared Compatible (SC)	NoLock (N)
Exclusive	-	-	-	-	+
Shared	-	+	-	+	+
Shared Exclusive	-	-	+	+	+
Shared Compatible	-	+	+	+	+
NoLock	+	+	+	+	+

Table 11.1: Lock mode matrix

Mapping The *Mapping* indicates the current Resource State Mapping that is using the resource.

P (Priority) The *Priority* defines which start priority the process had at the time of the submit in the **Submit Jobs** dialog. The values for the priority range from 100 (lowest priority) to 0 (highest priority).

EP (Effective Priority) The longer a job waits to be executed in the system, the higher the effective priority of its execution in the schedulix system. This is to make sure that previously started jobs are taken into account before the later jobs (depending upon the settings in the **Priority field in the Batches and Jobs dialog** and the start time).

How the priority rises over time can be set in the server configuration. The default setting causes the priority to rise by one point for every half an hour that elapses since the job was submitted.

Example: If a job was started at 12 p.m. with a priority of 50, the effective priority at 12:30 p.m. is 49, at 1 p.m. it is 48 and so on until the highest level has been reached.

Blocked Processes If the current job is marked in the list as being blocked, the reason why the job cannot be executed is shown in red.

Example: If the resource is currently exclusively allocated to another job and the current job also wants that resource to be exclusively allocated to it, the *Lock mode* field is shown in red.

11.4.3 Parameters tab

All the parameters defined in the job server or scope are administered in the Parameters tab. These parameters can be used for the jobs and those programs that are executed in the jobs.

The parameters are inherited in scope hierarchies. A parameter can be defined at the highest level and is then automatically made available to every scope and job server located below this level.

The Parameters tab looks like this:




Figure 11.9: Job server and scope parameters

A list of all the parameters is displayed here. If a job server has been selected in the navigation, all the parameters inherited from the parent scope hierarchy are visible as well in addition to the current parameters. However, they cannot be changed at this level. To change an inherited parameter at this level, it has to be recreated under the same name. This value overwrites the value of the inherited part. The displayed value can now be changed.

The columns in the list shown above have the following meanings:

Name The name of the parameter is shown here.

Type This is the type used for the parameter. The following options are available:

1. Dynamic

This designates a dynamic parameter. The value of this parameter is derived from the value of an environment variable that has to be maintained in the system on which the current scope or job server is running.

This value is only determined when the job server is started. If the job server is currently running on the computer, the actual value of the identically named environment variables on the computer and the environment in which the job server is running is displayed in the *Value* field.

Important: The environment variable must have been created on the computer and in the job server's start environment.

Editor

2. Constant

If the parameter type is "Constant", it is a constant expression that needs to be entered in the *Value* field. If a job or its program uses this value, the parameter is replaced by this constant expression when the job or program is executed.

Example:

If a script is to be run independently of the current database, the database connection can be mapped as a parameter. If the database connection parameter is simultaneously defined with respectively different values in both Scope 1 (for example the test database) and Scope 2 (for example the production database), the script can run in both environments without having to modify it because the database connection is defined in the parameter.

Value This is the value of the parameter. If the parameter type is "Constant", the parameter value has to be entered. If it is a dynamic parameter, the current value of the environment variables in the system on which the job server is running is displayed here.

11.4.4 Config tab

The configuration of the job server is described in the "Config" tab. A table of key/value pairs is used for the configuration. Please refer to the job server documentation for the meanings of the individual values.

Value	Inherit	From	Value
/home/dieter/BICsuite/Versions/2.5.1/SDMS/s	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	/home/dieter/BICsuite/Versions/2.5.1/SDMS/sandbox/ENTERPRISE-linux
/home/dieter/BICsuite/Versions/2.5.1/SDMS/s	<input type="checkbox"/>		
/home/dieter/BICsuite/Versions/2.5.1/SDMS/s	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	/home/dieter/BICsuite/Versions/2.5.1/SDMS/sandbox/tmp
false	<input checked="" type="checkbox"/>	DEFAULT	false
true	<input checked="" type="checkbox"/>	DEFAULT	true
5	<input checked="" type="checkbox"/>	DEFAULT	5
30	<input checked="" type="checkbox"/>	DEFAULT	30
true	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	true
true	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	true
0	<input checked="" type="checkbox"/>	DEFAULT	0
45500	<input type="checkbox"/>		
localhost	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES.LOCALHOST	localhost
8900	<input type="checkbox"/>		
localhost	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	localhost
2556	<input checked="" type="checkbox"/>	GLOBAL.EXAMPLES	2556

Figure 11.10: Job server and scope configuration

The fields in the "Config" tab have the following meanings:

Editor

Value The *Value* field shows the respective key/value pair.

Inherit The *Inherit* field shows whether the key/value pair value is being inherited from a parent scope. The current entry is deleted by setting the flag in the *Inherit* field. In this case, the value of the key/value pair is determined via the inheritance mechanism. The entry "DEFAULT" means that this is preset by the system.

From The *From* field shows from which scope the value is being inherited. If *Default* is displayed for the scope, these are the default configuration values.

Value The inherited value of the key/value pair.

11.4.4.1 Standard configuration parameters

The following table shows the names of the configuration parameters for a job server and their meanings:

Name	Meaning
JOBEXECUTOR	The fully qualified path of the job executor.
JOBFILEPREFIX	Path and file prefix of the task files.
DEFAULTWORKDIR	Working directory of the job server and the default working directory for the jobs executed using the job server.
ONLINE_SERVER	Boolean value. If this is set to "False", the job server logs off when there is no work to be done.
NOPDELAY	Time between two requests for jobs sent to the scheduling server.
RECONNECTDELAY	Time between two attempts to log on to the scheduling server.
VERBOSELOGS	If this is set to "True", the start and end times of the jobs are written to their log files in addition to the process output.
USEPATH	This parameter defines whether the job server uses the path setting when starting processes. If "USEPATH" is set to "False", all the program calls have to be fully qualified.
TRACELEVEL	Defines how "communicative" the job server is. <ul style="list-style-type: none"> • 0: Errors are logged • 1: Errors and warnings are logged • 2: Errors, warnings and information are logged • 3: Debug level; all messages are logged
NOTIFYPORt	The Notify port is the port to which a UDP packet is sent if a job is ready for executing by the job server.
HTTPHOST	The job server can send log files using the HTTP protocol. The two parameters "HTTPHOST" and "HTTPPORT" have to be set for this.
HTTPPORT	Port for sending log files.
REPOHOST	Host name or IP address of the scheduling server.
REPOPORT	Port for the scheduling server.

Table 11.2: Description of the job server configuration parameters

11.4.5 Env.Map tab

When a job is started, a number of default parameters for the job are sent to the job server as key/value pairs. These parameters can be set as environment variables before the process is started. Whether and under what name these variables are made visible is configured in the "Env.Mapping" mask.

The columns in the list shown above have the following meanings:

Editor




Figure 11.11: Job server environment mapping

Env.Var. The name of the environment variable is shown in the *Env.Var.* field.

Sys.Var. The name of the parameter is shown in the *Sys.Var.* field.

Inherit The *Inherit* field shows whether the key/value pair value is being inherited from a parent scope. The current entry is deleted by setting the flag in the *Inherit* field. In this case, the value of the key/value pair is determined via the inheritance mechanism.

From The *From* field shows from which scope the value is being inherited. If *Default* is displayed for the scope, these are the default configuration values.

Sys.Var. The name of the parameter is shown in the *Sys.Var.* field.

11.4.6 Logfile Pattern tab

The job servers can send the contents of log files using the HTTP protocol. To prevent the uncontrolled reading of arbitrary files, only files with names conforming to certain patterns can be requested.

Permitted patterns have to be entered by the administrator. If he does not do so, no files can be requested.

The log file patterns are regular expressions with a twist. The character string "/ .. /" or "\ .. \\" (Directory up) cannot usually occur in the name of a requested file. This rule makes it easier for the administrator to define the log file patterns.

It is generally advisable to test the whole name. A pattern such as ". * \ .log" allows all files to be read which contain this string somewhere in their name. The pattern ". * \ .log \$", on the other hand, is an important limitation, since the name has to end with this string.

The pattern "^ /tmp / * \ .log \$" is being used in the screenshot. All files below the directory "/tmp" that end with ".log" can be requested.

Editor




Figure 11.12: Job server log file name patterns

Requests for log files that do not match any of the patterns are logged in the job server's log file as errors. The messages look something like this:

```
ERROR [Jobserver] 01-12-2009 16:43:43 CET [HttpThread] ERROR: Illegal file request: /etc/passwd
```

together with some additional information (which is logged for each file request):

```
...[HttpThread] Got Request from 1.2.3.4 : GET /?FNAME=/etc/passwd HTTP/1.1
...[HttpThread] Got Request from 1.2.3.4 : Host: localhost:8881
...[HttpThread] Got Request from 1.2.3.4 : User-Agent: Mozilla/5.0 (X11; U; Linux x86_64; en-US; rv:1.9.0.15)
...[HttpThread] Got Request from 1.2.3.4 : Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
...[HttpThread] Got Request from 1.2.3.4 : Accept-Language: en-us,en;q=0.5
...[HttpThread] Got Request from 1.2.3.4 : Accept-Encoding: gzip,deflate
...[HttpThread] Got Request from 1.2.3.4 : Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
...[HttpThread] Got Request from 1.2.3.4 : Keep-Alive: 300
...[HttpThread] Got Request from 1.2.3.4 : Connection: keep-alive
...[HttpThread] Got Request from 1.2.3.4 : Cookie: _ZopeId='55722729A4JSGHnrljM'; tree-s='...'
...[HttpThread] Got Request from 1.2.3.4 : Cache-Control: max-age=0
```

An error message is displayed to the user:

```
ERROR: The requested filename doesn't match any of the configured patterns
```

12 Batches and Jobs

12.1 View




Figure 12.1: Batches and Jobs

12.2 Concept

12.2.1 In short

The Batch and Jobs dialog is where all the Job Definitions in the schedulix Scheduling System are managed.

The objects "Batches" and "Jobs" are used to define the units to be executed in the Scheduling System.

A job is a container for an executable script or program in the Scheduling System. When a job is submitted and started, the job server runs the specified program or script and sends feedback after the process has been completed as to whether it has succeeded or failed (Exit State).

A batch is a container for other objects and contains so-called children. When a batch is started, all the children are started automatically as well. A batch does not have a program or script that is executed.

The batch (or job) that is the first object started in the Submit Batches and Jobs dialog is called the Master Job.

12.2.2 Detailed description

Each Scheduling Entity (batch or job) can belong to a batch run in a parent-child hierarchy. A job is started when its parent has been started. This means that when such an executable object hierarchy is submitted, all the jobs are submitted from top to bottom. However, this takes place within a transaction so that all the actions, as seen by an outsider, are performed simultaneously.

Each task (batch or job) can have dependencies. A dependency describes which executable object (required job) must have been executed beforehand. If only one specific Exit State is to be taken into account, the dependent job is not started until the required job has been completed in full and the correct Exit State has been reported back to the system.

This allows dependency models to be created which ensure that all the lead jobs really could be successfully executed in order to start a follow-on process (e.g. a report).

Each task can require different **resources**. This can be defined by entering the **environment** and a **footprint** and by specifying the required resources.

This ensures that has all the required resources (e.g. sufficient main memory and CPU units, the correct database and the requisite system programs, etc.) are available to the runtime environment being used to execute the job.

Synchronizing Resources can also be used to make sure that processes do not hinder one another or that the resources have specific states and are up to date time-wise. A load distribution can also be achieved using the resources.

Each Scheduling Entity can define or use certain parameters. This enables data to be transferred between the individual jobs or from the runtime environment. The parameters can be transferred in both directions within the parent-child relationship (from parent to child, from child to parent).

However, it is also possible to access the parameters of Scheduling Entities that do not belong to the parent-child relationship. This allows data to be transferred between all the Scheduling Entities of a Master Job. More information about the parameters can be found in Chapter [12.5.9](#).

Actions can be defined for a task to be performed when a particular event occurs (e.g. completion of the job) or when specific Exit States are attained. These actions are called triggers. A trigger defines which action is to be performed when a specific event occurs. The action that is triggered is again an executable object.

If the event occurs, this trigger is initiated and the job (or batch) is submitted. This functionality can be used to send a message to a system administrator, for example.

12.2.3 Recommended convention for batch objects

For the sake of clarity and to facilitate working with batch objects, we advise our users to adhere to the following convention.

A Batch Object should be located in a folder with the same name. This folder should contain the batch and its children (sub-batches and jobs). Sub-batches should, in

turn, be located in a separate subfolder. Job Objects that have children should be treated like batches, i.e. they should also be located in separate subfolders.

This convention is not mandatory, but it can significantly reduce the time and effort required to grant access privileges or deploy, copy and move sub-workflows, etc., and greatly improve the orientation.

The following function extensions have been implemented in the schedulix!web interface to make it easier to apply this convention:

- When you create a folder, a Batch Object can also be created at the same time in this folder.
- When you create a batch, you can also create a parent folder for the batch.
- When renaming batches or folders, the folder or batch with the same name can be renamed as well.
- When creating objects in a folder with a batch or job of the same name, at the same time the new object can be added to the batch or job as a child.
- When cloning jobs, the new job can be added as a child to the same batch or job in the folder.
- When cloning folders, the batch or job of the same name in the new folder can be automatically renamed.
- When cloning folders, the batch or job of the same name can be added as a child to the same batch or job in the parent folder.

12.3 Navigator

All objects of the types "batch" and "job" are hierarchically managed in the Navigator in containers (so-called folders). These folders are used for clearly organising objects.

In addition, parameters and a default **environment** can be defined in them. All the parameters for the folder are available to all the jobs in that folder. The folder environment behaves additively to the environment of a job.

12.3.1 Pinning

For greater clarity in the Navigator folder, objects can be "pinned" by clicking the *Pin* icon .

A pin can be removed again by clicking the *Unpin* icon .

Navigator

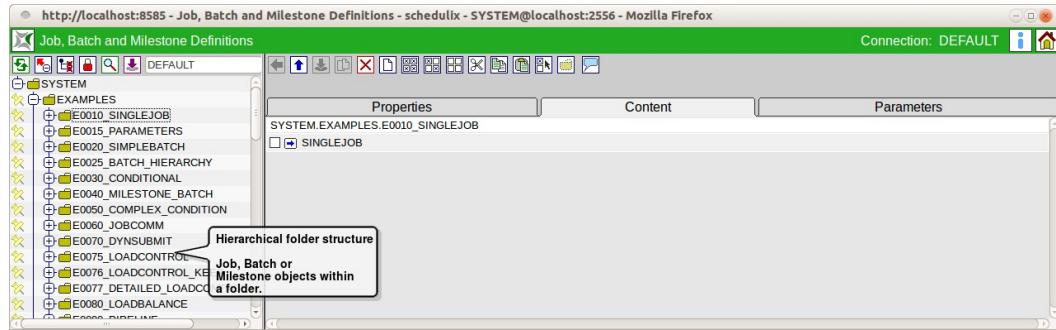


Figure 12.2: Navigator for batches and jobs

Rows with a pin are always displayed regardless of whether the parent folders are expanded or not. This allows you to scroll significantly more quickly in the Navigator.

The "Expand" and "Collapse" icons in the tree view are coloured yellow if there is a pinned object below a folder. This indicates that collapsing a folder will not close it completely, but that the path to the pinned object will always remain visible.




Figure 12.3: Batches and Jobs; view without a pin

12.3.2 Folder bookmarks



Navigator




Figure 12.4: Batches and Jobs; view with a pin

Clicking the *Store Bookmark* button creates a bookmark for the current Navigator setting.

The navigation settings contain information about how the jobs, batches and locked objects are displayed, the current expansion state of the tree and the pinned objects. The name for the saved bookmark is entered in the box to the right of the *Store Bookmark* button.

The "DEFAULT" bookmark is used when calling the "Batches and Jobs" dialog from the schedulix!web desktop.

Other bookmarks can be opened by clicking **Bookmark (Folder)**.

If the current user is a "Web GUI Admin", an options field is displayed next to the name of the bookmark.

Whether the bookmark is to be saved as a System Bookmark (visible to all users) or a User Bookmark (only visible to the current user) is set in this field.

12.3.3 Folder search



Clicking the *Find* button switches the Navigator to search mode.

In the input field, you can enter a search pattern corresponding to the syntax selected in the options field to the right of the input field. The SQL LIKE syntax and regular expressions are supported.

In the next options field, you can define whether this pattern is to apply to the entire path (FULL) or only the end of it (TAIL).

The hits are displayed in the Navigator by clicking the *Find* button in search mode. Clicking a hit opens the object in the editor pane, the object is pinned, and the search mode closes again.

12.4 Folder editor

12.4.1 Properties tab

All information pertaining to the general properties of the folder is recorded in the Properties tab.

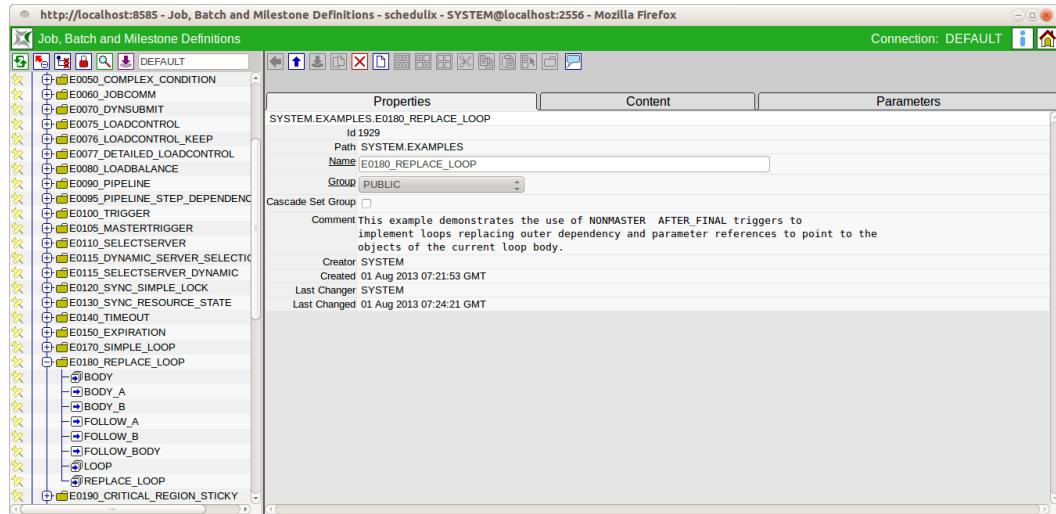


Figure 12.5: Folder properties

ID The *ID* is the unique identification number of the object.

Path The *Path* describes the parent folder hierarchy. All parent folders are separated by periods.

Name The name serves as a comprehensible and unambiguous identification of the object in the current context. It must be unique in the folder. Identical names in different folders are allowed.

Rename in Content If the name of a folder is changed and the folder conforms to the "Batch in Folder" convention (contains a batch or job object of the same name), the *Rename to Content* check box is displayed. If this option is selected and the folder is renamed or cloned, the batch or job object in the folder is renamed as well.

Group The *Group* field defines the schedulix user group that owns the object.

Cascade Set Group The *Cascade Set Group* is used to set the user group for the folder and all the subordinate objects.

Create Batch in Folder The *Create Batch in Folder* check box is displayed when creating a new folder and is used in support of the recommended convention for folders as well as batch and job objects. If a check mark is set here, a batch object of the same name is also created in the new folder.

Batch Exit State Profile The *Batch Exit State Profile* field is only displayed if the *Create Batch in Folder* check box has been enabled. It is used to select the Exit State Profile for the new batch.

Add as Child The *Add as Child* check box is displayed in the following cases. When creating a new folder, the *Create Batch in Folder* check box is enabled and the parent folder of the newly created folder conforms to the "Batch in Folder" convention (it contains a batch or job object of the same name). An existing folder is renamed or changed and both the folder itself as well as its parent folder conform to the "Batch in Folder" convention (they both contain a batch or job object of the same name). In this case, the *Add as Child* check box only has an effect when cloning the folder. The *Add as Child* check box is ignored when the folder is saved (renamed). The *Add as Child* check box can be used to select whether the newly created batch or a copied (cloned) batch or job is also to be added as a child of the parent batch or job folder.

Environment The *Environment* field is a selection box for choosing a valid **Environment** Object. This selection defines the additional requirements for all the jobs in the folder that have to be fulfilled by their execute environment.

Comment A more detailed explanation about the object can be entered in the *Comment* field.

12.4.2 Content tab

The content of the folder is shown in the Content tab.

12.4.3 Parameters tab

The parameters can be defined in the Parameters tab.

Name This is the name of the parameter. All the parameter values are exchanged between the individual jobs using this name.

Folder editor

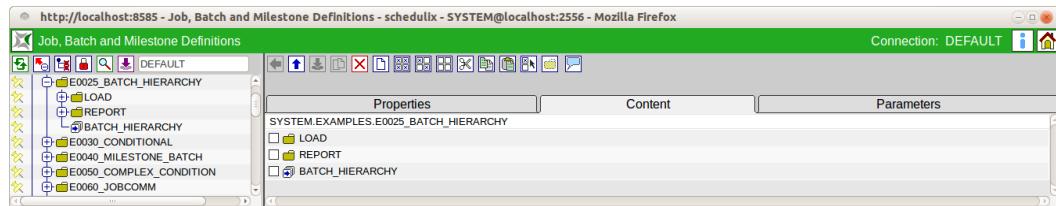


Figure 12.6: Folder content

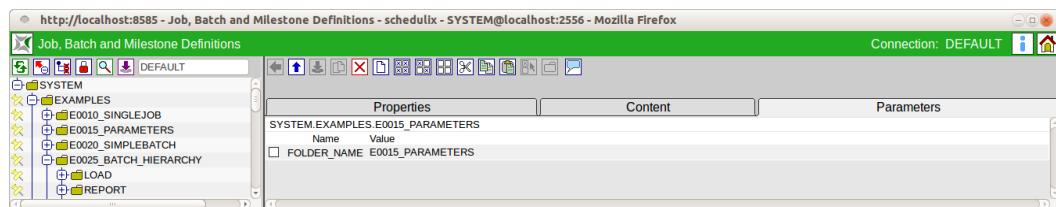


Figure 12.7: Folder parameters

Value The *Value* field shows the parameter value.

12.4.4 Resources tab

The resources are shown in the Resources tab. Clicking the name of the resource opens the **Resource Details** tab.

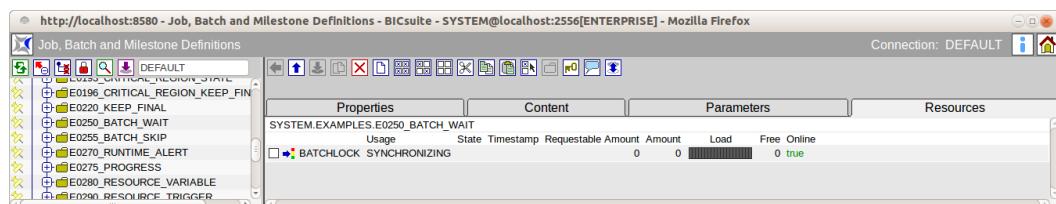


Figure 12.8: Folder resources

Usage The *Usage* field specifies the "Resource" type. More details about resource types can be found in the **Named Resource** dialog.

State The *State* field is only visible if the resource usage has been set to "Synchronizing" and a **Resource State Profile** has been assigned.

It designates the current status of the resource in this scope or job server. The state can be set or changed with this value. The drop-down list contains all the valid Resource States of the Resource State Profile, which can be selected here.

This may be necessary when manually fixing an error to restore a resource to a state that allows further processing by follow-on jobs.

If you manually change this value, the value in the *State* field in the Resources tab is not automatically refreshed. Click the *Refresh* button if this is required and necessary.

Timestamp This field is only visible if the resource usage has been set to "Synchronizing" and a **Resource State Profile** has been assigned. The *Timestamp* indicates the time of the last change made to a Resource State.

Timestamps play a role if an expiration interval has been defined in an expiry object. If this is the case, the timestamp must not be older than the specified interval. More details about expiration times can be found in Chapter [12.5.8.1](#).

If the completion of a job causes the Resource State to change, the timestamp is automatically updated by the schedulix Server.

Requestable Amount The maximum amount of resources that can be requested by a job.

Amount The *Amount* is the number of resource instances allocated to or requested by the current job. If the value of the amount exceeds the currently available number that can be maintained in the *Amount* field in the Resource Details tab and there is no alternative scope available with a sufficient number, the job cannot be executed.

Load The *Load* shows the graphical load of the resource.

Free This field is only visible if the resource usage has been set to "Synchronizing" or "System".

The *Free Amount* designates the total number of instances of a resource in the selected scope or job server that have not yet been allocated to jobs.

Online *Online* can be used to switch a resource in this scope or job server online or offline. An offline resource is temporarily unavailable.

12.5 Editor for Job Definitions

12.5.1 Properties tab

All information pertaining to the general properties of job or batch objects is recorded in the Properties tab. Only the respectively required information is displayed in the

Editor for Job Definitions

Properties tab depending on the type you selected in the navigation or while creating the new object.

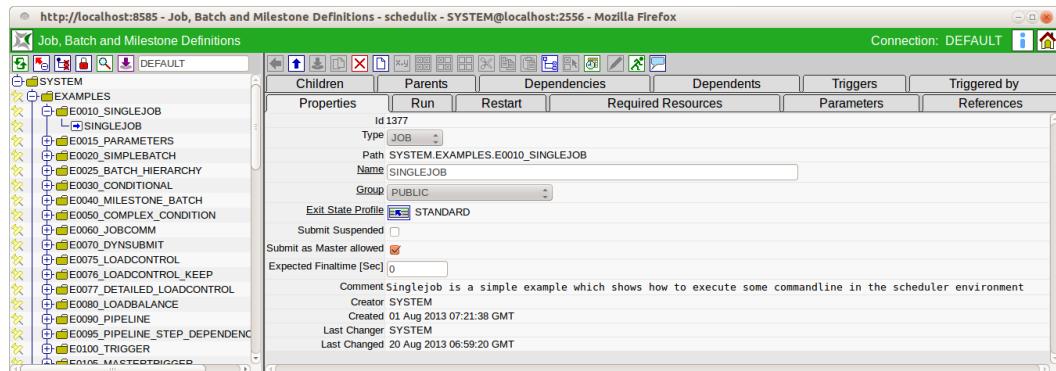


Figure 12.9: Job Definition properties

The fields in the "Properties" tab have the following meanings:

ID The *ID* is the unique identification number of the object.

Type This is the *type* of object. One of the following types can be selected here:

1. Batch

The object is a batch. A batch is used as a startable container for a number of jobs or batches. It does not itself have an executable program, but is only used as the start object for its child objects.

An example of modelling as a batch is a number of jobs that have to be executed on a daily basis. These can all be grouped under a batch called "Daily". All weekly jobs are likewise grouped under the batch "Weekly", which also includes the "Daily" batch as this obviously has to be executed as a batch at the same time. A "Monthly" batch can additionally be created that contains all the monthly executable jobs. This means that only these three batches have to be created in this system in the **Time Scheduling** and all the child jobs and batches are started automatically at these times.

The legend for the graphic can be found in Chapter 1.8.

2. Job

All executable objects (programs, scripts, etc.) that are to be executed in the Scheduling System on a **job server** must be defined as a job.

3. Folder

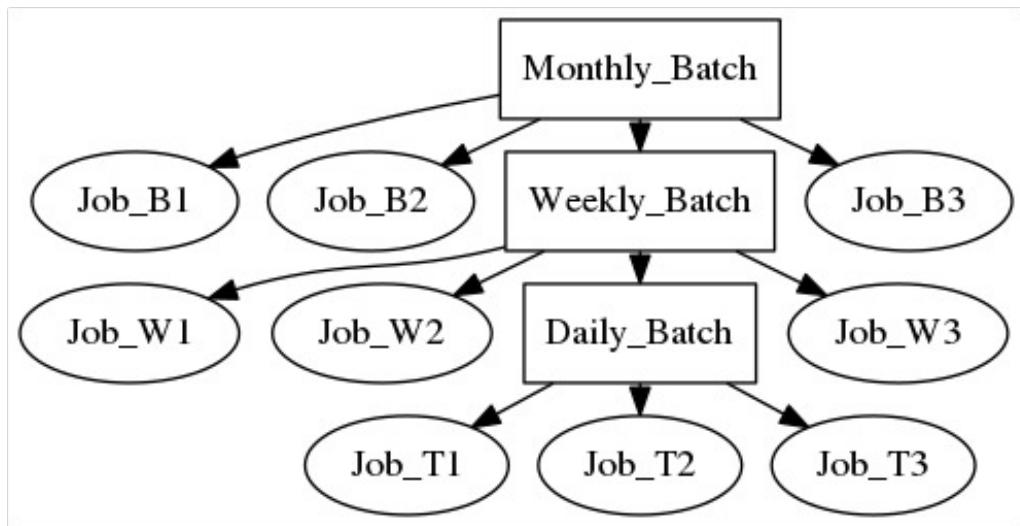


Figure 12.10: Examples of batches

A folder is a folder within the folder hierarchy for the executable objects. It is used for organising the other objects because in a normal scheduling system there are many different jobs and batches that have to be managed, and with a flat hierarchy it is very easy to lose track of them. How the hierarchy is structured is decided by each user. It could be logically organised (all objects on the same theme in one folder, etc.) or it could map the departmental or human resource structure (all objects relating to user/department X are placed in the same folder). The user should just be able to quickly find his way around the hierarchy and easily find the required objects.

Folders cannot be executed nor can they have dependencies. It is possible to define an **environment** for a folder. Furthermore, parameters can be defined which can be used by all the jobs located in the folder.

Path The *Path* describes the parent folder hierarchy. All parent folders are separated by periods.

Name The name serves as a comprehensible and unambiguous identification of the object in the current context. It must be unique in the folder. Identical names in different folders are allowed.

Decomposite The *Decomposite* check box is displayed if an object type is changed from "job" to "batch". If a check mark is set here, a so-called decomposition is

carried out. Not only is the object type changed to "batch", but a job is also created as a child of this batch which then takes over the job properties of the original job object. The batch object produced by the decomposition (with a job as a child) behaves, in the context of the Scheduling System, as the previously existing batch object. All the time schedules, dependencies, parameter references, etc. are retained. The result is an ideal starting point for decomposing jobs into smaller single steps (process decomposition).

Job Name The *Job Name* is only displayed if the *Composite* check box has been enabled. The name of the Job Object to be created during the decomposition can be entered here.

Rename Parent Folder If the name of a batch is changed and the batch conforms to the "Batch in Folder" convention (it is located in a folder of the same name), the *Rename Parent Folder* check box is displayed. If this check box is enabled, when you save (rename) the batch the parent folder of the same name is renamed as well. The *Rename Parent Folder* check box is ignored when cloning a batch.

Create Batch Folder The *Create Batch Folder* check box is displayed when a new batch is created or the type of an object is changed to "batch". It serves to support the recommended convention for folder, batch and job objects. If a check mark is set here, a folder with the same name as the batch is created for it and the new batch is created in this folder. The *Create Batch Folder* check box is also displayed if an object type is changed from "job" to "batch" (decomposition).

Add as Child The *Add as Child* check box is displayed in the following cases. When creating a new batch or job object where the parent folder of the new object to be created conforms to the "Batch in Folder" convention (contains a batch or job object of the same name).

In the case of an existing job where the name is changed and the parent folder conforms to the "Batch in Folder" convention. In this case, the *Add as Child* check box only has an effect when cloning the job. The *Add as Child* check box is ignored when the object is saved (renamed).

The *Add as Child* check box can be used to select whether the newly created object or a copied (cloned) job is also to be added as a child of the parent batch or job folder.

Group The *Group* field defines the schedulix user group that owns the object.

Exit State Profile An *Exit State Profile* can be selected in this field.

Submit Suspended The *Submit Suspended* switch can be used to delay the actual start of a workflow, for example where external authorisation by a third person is necessary for a job. Only when this authorisation has been given can the job be started.

Resume This field is also only displayed if the option *Submit Suspended* has been enabled. Here you can choose whether a workflow should be resumed automatically.

The following options are available:

- NO: disables this function.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Resume Time This field is only displayed if "AT" was selected for *Resume*. The required resume time is entered here in the format YYYY-MM-DDTHH:MI:SS. This format is based on the ISO standard 8601 and also permits an incomplete entry. If you enter 'T16:00', the job will be resumed at 16:00 hours (starting from the current set time).

Resume In This field is only displayed if "IN" was selected for *Resume*. Here you can specify how many time units (see *Unit*) the system is to wait for until the resume action is triggered.

Unit This field is only displayed if "IN" was selected for *Resume*. This field is used to define whether the entry in *Add Resume* is in minutes (MIN), hours (HOUR) or days (DAY).

Nice Value This field is only visible if the object is a batch.

The *Nice Value* indicates the priority to be applied for running children of this object. It is an offset value which is added to the priority of the children.

Example:

If the process has a Nice Value of -50 and the child has a priority of 100, then the priority of the child process of the batch is 50. If the child process were to be started on its own, the priority would be 100.

Editor for Job Definitions

Submit as Master Allowed This field is only visible if the object is a batch or a job.

This switch specifies whether it is possible to submit this job as a Master Job in the **Submit Batches and Jobs** dialog. If the switch is not enabled, this is not possible and the job or batch cannot be executed on its own, but only as a child of any parent job.

Comment A more detailed explanation about the object can be entered in the *Comment* field.

12.5.2 Run tab

The Run tab is only displayed if an object in the navigation has the type "Job" or if, in the case of a new object, "Job" was chosen in the selection box.

All the information about the executable command line for a job object is managed in the Run tab.




Figure 12.11: Run Job tab

The fields in the "Run" tab have the following meanings:

Priority The *Priority* field indicates the urgency with which the process, if it is to be started, is to be considered by the Scheduling System. The range of values for priorities ranges from 100 (very low) to 0 (very high). In the system, all the startable jobs with their priorities are taken into account and the job with the highest priority is started. If two jobs have an identical priority, the job with the lowest ID will be started.

The longer the waiting time for a job before it is started in the Scheduling System, the higher its priority. This means that if a process is started with a priority of 50, its priority increases per time unit. This is set in the Scheduling System configuration. For example, the priority of a job can increase by one point every half an hour. This means that the priority is 49 after half an hour, then 48 after one hour, etc. Consequently, older jobs are accorded a higher priority and taken into account before more recently submitted jobs.

Run Program A command line to be executed by a suitable job server is entered in the *Run Program* field. This command line is first interpreted by the Scheduling System and then decomposed into the program call and single parameters. The Bourne shell rules with regard to quoting are followed here.

If any environment variables or parameters are to be specified, these have to be quoted conformant to the respective execute environment (depending on the shell or command line interpreter of the environment; see [Environment](#)). A list of all the usable standard parameters can be found in the paragraph [12.5.9.2](#). Please refer to the documentation on your command line interpreter (shell) for more information about the quoting to be used.

If the target system is a Windows system, some other difficulties are very likely to arise. When a table of executables and parameters is made available for execution under Windows, first of all Windows compiles a command line from the table which it then interprets. The use of quoting under Windows also causes problems because the called programs either remove the quotes or don't remove them depending on the type of executable (.exe or .bat).

If the executable program (the first element of the command line) is a valid integer, the command line is not run by the job server. Instead, the job is treated as if it had completed itself with the integer as the Exit Code. Dummy jobs with 'true' or 'false' as the program can now be implemented as '0' instead of 'true' or '1' instead of 'false' and are therefore processed much more efficiently and quickly by the system.

Should it really be necessary to run an executable with a number as the name, this can be achieved by using a path prefix (e.g. './42' instead of '42').

Exit State Mapping An existing [Exit State Mapping](#) can be entered in the *Exit State Mapping*. This defines the mapping of the Exit Code for a job after an Exit State. In the absence of a specified mapping, the default mapping from the Exit State Profile is used. The mapping must be compatible with the profile.

Environment In the *Environment* field it is necessary to select an existing [Environment](#). All the Named Resources in this environment must be present on a job server so that the job can be run on that server.

Footprint An existing **footprint** can be selected in the *Footprint* field. All the Named Resources for this footprint must be present on a job server in a sufficient quantity so that the job can be run on that server.

Expected Runtime The *Expected Runtime* describes the anticipated time that will be required to execute a job. This is entered manually and can be used to monitor the runtime, for example.

Expected Finaltime The *Expected Finaltime* describes the anticipated time required to execute a job (from the submit until it attains a final state). This has to be entered manually and is used for the display in the calendar. There is obviously a standard parameter for this as well so that the expected value is made available to the job.

Rerun program The *Rerun Program* field specifies the command that is to be executed when repeating the job following an error (rerun). This can be a different command to Run program as it may be necessary to perform some cleanup work prior to a restart before the original command can be executed again.

If no value was entered in the Rerun program, the Run program is called when restarting the job.

See 'Run program' above for details regarding processing (quoting, treating integers as a program name).

Kill program The *Kill program* field determines which program is to be run to terminate a currently running job.

If no value is entered in this field, it will not be possible to prematurely terminate the job from the schedulix interface.

An example of a Kill program in a Unix environment is the "kill" command. The process ID of the current job can also be accessed in the Kill program using the predefined system variable "PID" (example: "kill -9 \$PID").

Working Directory The *Working Directory* specifies the directory in which the Run program is to be started. The job server switches to the specified directory before the Run program is started.

If no value is entered in this field, the default directory of the **job server** is used as the working directory.

A parameter or a folder, scope or job server parameter can also be specified as the working directory.

Logfile The *Logfile* field specifies the file in which all the normal outputs of the Run program are to be returned. These are usually all the outputs that use the standard output channel (stdout under UNIX).

Parameters can also be replaced in the *Logfile* field. For example, the parameter JOBID can be used to allocate a job a unique name or a Timestamp can be used for the date.

If no value is entered, all the outputs will be discarded.

If no path is specified in the log file, it is written to the working directory.

Logfile Write The *Logfile Write* field tells the system where a log file is to be created. The following options are available:

- Append

The Append option appends all the created log files to an existing log file. If no log file exists yet, a new one is created.

- Truncate

With the Truncate option, the log file is deleted and a new one is created for each new action.

Error Logfile The *Error Logfile* field specifies the file in which all the error outputs from the Run program are to be returned. Error outputs are all the outputs that use the error output channel (stderr under UNIX).

Parameters can also be replaced in the *Error Logfile* field. For example, the parameter JOBID can be used to allocate a job a unique name or a Timestamp can be used for the date.

If no value is entered, all the outputs will be discarded.

If no path is specified in the log file, it is written to the working directory.

If the same name as the one in the *Logfile* field is being used, these outputs will also be written to this file. The sequence of the outputs (error or normal) is random due to the buffering, however. No outputs from the normal output channel (stdout) are overwritten regardless of the content of the "Error Logfile Write" field.

Error Logfile Write The *Error Logfile Write* field tells the system where the error log file is to be created. The following options are available:

ERROR LOG FILE WRITE OPTIONS

- Append

The Append option appends all the created error log files to an existing error log file. If no error log file exists yet, a new one is created.

- Truncate

With the Truncate option, the error log file is deleted and a new one is created for each new action.

Editor for Job Definitions

12.5.3 Restart tab

The restart behaviour for a job can be defined in the Restart tab.

A job will only be restarted if its Job State is "FINISHED".

If a job gets a "RESTARTABLE" Exit State because of an "ERROR" Job State, it will NOT be restarted!




Figure 12.12: Job Restart tab

The fields in the "Restart" tab have the following meanings:

Restart This determines whether the job is to be automatically restarted if it attains a "RESTARTABLE" Exit State.

The following selections are possible:

- NO: No restart. No other fields are displayed in the "Restart" tab.
- IMMEDIATE: Immediate restart.
- AT: Restart at a specified time (e.g. "T16:00" at 16:00 hours).
- IN: Restart after a delay (e.g. 10 minutes)

Restart Time This field is only displayed if "AT" was selected under *Restart*.

Time when the job is to be restarted using the format "YYYY-MM-DDTHH:MI:SS". This format is based on the ISO standard 8601 and also permits an incomplete entry. If you enter 'T16:00' here, the job restart will be suspended and resumed at 16:00 hours (based on the submit time of the job).

Restart In This field is only displayed if "IN" was selected under *Restart*.

Number of time units (see *Unit*) for the delay until the job restarts.

Unit This field is only displayed if "IN" was selected under *Restart*.

Units for the delay interval (see *Restart In*).

Editor for Job Definitions

Restart On Table of Exit States that are to trigger a restart.

If the table is empty, jobs are only started for all "RESTARTABLE" Exit States.

Max. Restarts Maximum number of times that the job is restarted.

The system parameter "TriggerSoftLimit" applies if nothing is entered here. Entering 0 will restart the job any number of times.

Limit State Here an Exit State can be set, which will be the Exit State of the job, if the maximum number of restarts was reached.

12.5.4 Children tab

All the executable objects that have been defined as a child of the current job are displayed as a list in the Children tab. Other objects can be added as children using **copy and paste**.

The sequence in which the children are run is not fixed, but it can be defined in the dependencies.

The Children tab looks like this:

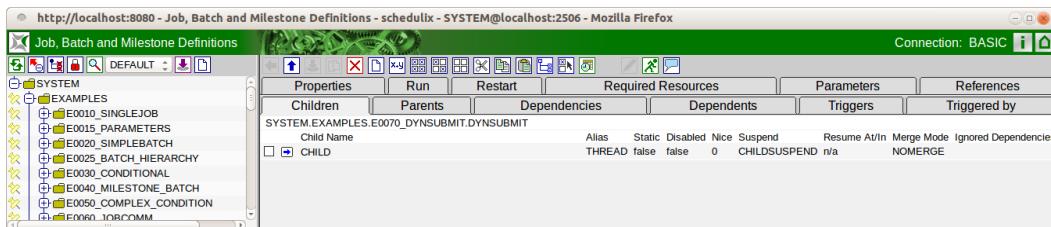


Figure 12.13: Batches and Jobs; Children tab

A list of all the job and batch objects that are children of the selected job is shown here. The columns in the list shown above have the following meanings:

Child Name This is the name of the child. Clicking the name opens the Child Details tab.

The other columns are explained in the next section.

12.5.4.1 Child Details tab

The Child Details tab is displayed after selecting the name of a child object in the Children tab. Here, all the data already displayed in the table as elements is listed again and can also be modified here. The tab looks like this:

The fields in the "Child Details" tab have the following meanings:

Editor for Job Definitions

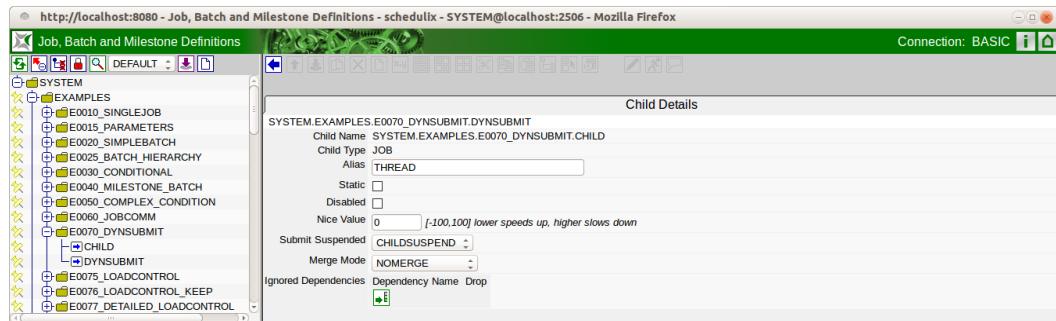


Figure 12.14: Batches and Jobs; Child Details tab

Child Name The *Child* name is always displayed with the full path to the parent folder hierarchy.

Child Type The *Child Type* specifies the type of the child (batch or Job).

Alias A child can be assigned a new logical name by entering it in the *Alias* field. This name must be unique among all the children of the selected job. The alias is only significant for Dynamic Children. The Parent Job can create additional instances of the children by stating the alias in the submit. This means that the parent does not need to know where the child has been saved. This abstraction makes it considerably easier to reorganise the folder structure.

Static The *Static* switch indicates whether a child is to be submitted statically or dynamically. In the case of a static child, the child instance is instantiated when the parent is submitted. If it is a dynamic Child (the *Static* switch has not been set), one or more instances of the child can be dynamically created at runtime.

Children are often used for dynamic parallelisation tasks. Here you only decide during the runtime of the Run program whether and how many dynamic children are to be created. These are then created by the schedulix API. This allows you to split up a process into a large number of parallel working sub-processes.

Disabled If the *Disabled* switch is set, instead of the real batch or job defined as child, a dummy job with the same name is submitted, which is treated by the system like an empty batch. The dependencies of the disabled job or batch will be handled correctly.

This way it is possible to skip a batch or job, which is part of a processing chain without loosing the order of processing of the chain.

The legend for the graphic can be found in Chapter 1.8.

Editor for Job Definitions

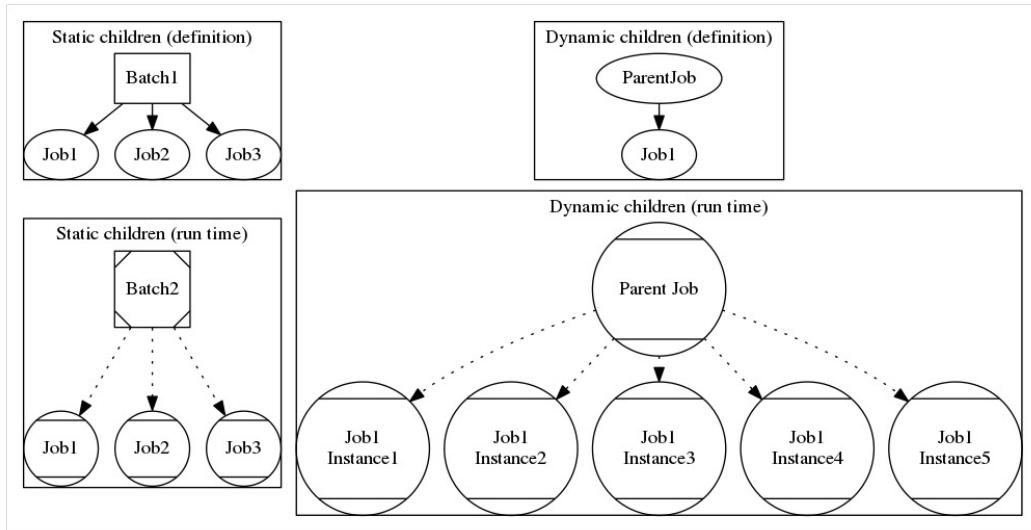


Figure 12.15: Static and dynamic children

Nice Value The *Nice Value* indicates the priority to be applied for running the selected child. It is an offset value which is added to the priority of the child.

Example: If the process has a Nice Value of -50 and the child has a priority of 100, the job instance started as a child of this job has a priority of 50. If the child process is started on its own, it has a priority of 100.

If the parent is a batch, the Nice Value of the child is added to that of the parent.

Submit Suspended The *Submit Suspended* parameter specifies the form in which the child object is delayed when being started or if it can be started immediately. The following options are available:

1. YES

The child is created as being suspended when it is submitted and it has to be started by being explicitly released.

2. NO

The child is not suspended when it is submitted and can be started immediately.

3. CHILDSUSPEND

Whether a delay takes place or not depends on the *Suspend* field for the child job. This means that the setting defined in the child job is applied.

Resume This field is also only displayed if the option *Submit Suspended* has been set to "YES". Here you can choose whether a workflow should be resumed automatically.

The following options are available:

- NO: disables this function.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Resume Time This field is only displayed if "AT" was selected for *Resume*. The required resume time is entered here in the format YYYY-MM-DDTHH:MI:SS. This format is based on the ISO standard 8601 and also permits an incomplete entry. If you enter 'T16:00', the job will be resumed at 16:00 hours (starting from the Job Submit time).

Resume In This field is only displayed if "IN" was selected for *Resume*. Here you can specify how many time units (see *Unit*) the system is to wait for until the resume action is triggered.

Unit This field is only displayed if "IN" was selected for *Resume*. This field is used to define whether the entry in *Add Resume* is in minutes (MIN), hours (HOUR) or days (DAY).

Merge Mode Merge Mode indicates whether a child object is to be started multiple times within a "Master Job" run or not.

The legend for the graphic can be found in Chapter [1.8](#).

The following options are available:

1. No Merge

Jobs are not merged. The child job is started without taking into consideration whether this job has already been run in the current Master Batch.

In the example shown in Figure [12.16](#), the job C3 is started three times if No Merge is set in each parent-child relationship.

2. Failure

When the job is submitted, the system checks whether the child job has already been started. If this is the case, an error message is returned and the start of the Master Job is aborted.

In the example shown in Figure [12.16](#), the job cannot be submitted.




Figure 12.16: Example of Merge Mode

3. Merge Local

A merge is performed in the child hierarchy of the parent object. This means that if the job in the child hierarchy has already been started once, it will not be started again. If the job has already been started within the Master Job but outside the child hierarchy of the parent, this fact is ignored and the job is submitted again.

The dependencies of this job are taken into account just as if this job had been started in this child relationship.

If in the example above Merge Local Mode has been set for the child relationship $C2 \rightarrow C6$ and $C2 \rightarrow C5$,

the job $C3$ is only started twice because the merge took place in the child hierarchy of $C2$. As the relationship $C8 \rightarrow C3$ is not in this child hierarchy, $C3$ is executed because no merge has been performed here.

4. Merge Global

The system verifies within the entire Master Job whether this job has already been started before. If this is the case, the job will not be restarted.

This may be necessary if a job should only run once, but it is needed as a requirement by several jobs. It can be defined as a child of multiple parent processes. The parent process that is run first starts the process. The other parent processes use Merge Global to check whether this job has already been run and they don't start it again.

The dependencies of this job are taken into account just as if this job had been started in this child relationship.

In the example above, the job C3 is only started once because a global merge takes into account all the relationships in the Master Job and therefore performs the merge for the relationships (C8→C3, C6→C3 and C5→C3).

Which relationship is used to actually start the job is immaterial.

List of Ignored Dependencies Here you can add a list of dependencies which are to be ignored by the child in this parent-child relationship. All of these dependencies are no longer taken into consideration for the child and it can be started without the dependency on the parent being fulfilled.

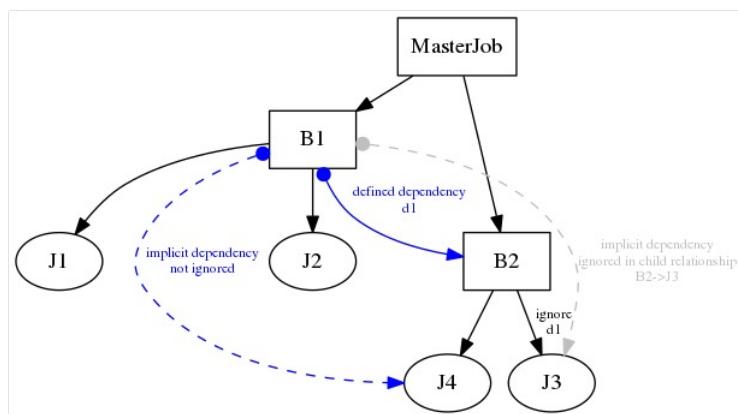


Figure 12.17: Example of Ignored Dependencies

An example of Ignored Dependencies ist shown in Figure 12.17. The legend for the graphic can be found in Chapter 1.8.

Only those dependencies that have a value entered in the *Dependency Name* field can be ignored. Unnamed dependencies cannot be chosen in the selection box and are therefore not ignored.

12.5.5 Parents tab

The Parents tab shows a list of higher-ranking parent objects which use the selected object as a child. This therefore visually represents the upward view, whereas the children represent the downward view since all the children of the currently selected object are displayed there. The tab looks like in Figure 12.18.

The tab is for information purposes only. Values cannot be changed nor can you add any new list entries. This is only possible in the direction parent-to-child.

The columns in the list above are explained under "Children tab".

Editor for Job Definitions




Figure 12.18: Batches and Jobs; Parents tab

12.5.6 Dependencies tab

The Dependencies tab shows a list of all the objects upon which the current object is dependent. A dependency exists if the current object is not allowed to start until all the objects upon which the job is dependent have been executed beforehand and completed with a specific **Exit State**.

Additional objects can be added as required Scheduling Entities using **copy and paste**.




Figure 12.19: Batches and Jobs; Dependencies tab

The tab looks like in Figure 12.19

The field in the list shown above has the following meaning:

Dependency Mode Dependency Mode states the context in which the list of dependencies has to be viewed. The following options are available:

1. ALL

All the dependencies must be jointly fulfilled before the object can start. This corresponds to linking all the conditions with "AND".

Example: Object C is dependent upon Job A and Job B (that is, it should only start once A and B have finished) and so the Dependency Mode ALL needs to be selected.

2. ANY

At least one of the dependencies must be fulfilled before the object can start. This corresponds to linking all the conditions with "OR".

Example: Object C is dependent upon Job A or Job B (that is, it should only start once A or B has finished) and so the Dependency Mode ANY needs to be selected.

The dependencies of a job are either logically linked by the ANY operator or by the ALL operator. A batch object is used where it is necessary to map complex dependency links (for example, dependent upon A or B and C or D like in Figure

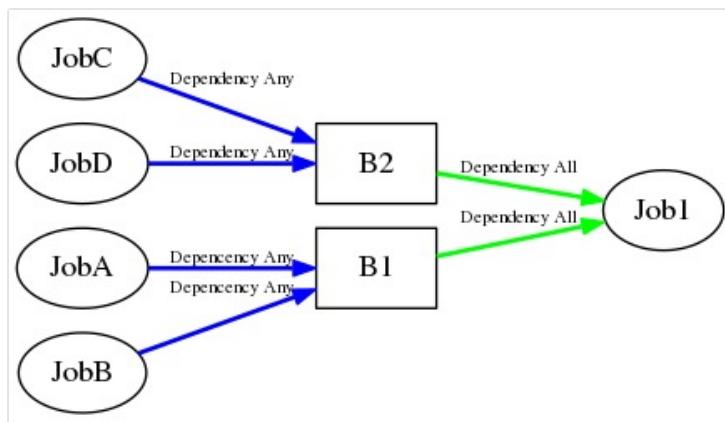


Figure 12.20: Example of Dependency Modes

12.20). The legend for the graphic can be found in Chapter 1.8.

In the example above, batches B1 and B2 were generated. B1 has linked the dependencies A and B with ANY and B2 has linked the dependencies C and D with ANY as well. Our object (Job1) can now link the two batches as a dependency using ALL, and in doing maps the complex condition (dependent upon A or B and C or D).

The "Dependencies" list shows all those objects upon which the currently selected object is dependent. All the elements in the list are linked using Dependency Mode. The list columns are explained in the next section.

12.5.6.1 Dependency Details tab

The Dependency Details tab opens when an entry has been selected in the Dependencies tab by clicking the Required Name.

All the aspects of a dependency between the two jobs can be configured here. The tab looks like this:

The fields in the "Dependency Details" tab have the following meanings:

Editor for Job Definitions




Figure 12.21: Batches and Jobs; Dependency Details

Dependency Name The *Dependency Name* is optional and is the prerequisite for ignoring dependencies.

Required Name This is the name of the object upon which the currently selected job is dependent. It corresponds to the full name with the parent folder hierarchy.

Required Type The type of the required object is displayed in the *Required Type* field.

Check On The type of integrity check to be used for the required object is described in the *Check On* field. The following options are available:

1. ALL_FINAL

If this option is selected, both the required job and all its children must have attained a Final State.

2. JOB_FINAL

If this option is selected, the system only checks whether the job has attained a Final State. The state of the child objects is not verified.

Unresolved Handling The *Unresolved Handling* selection field describes what to do if a dependent object instance is not present in the current Master Batch.

The legend for the graphic can be found in Chapter 1.8.

The following options are available:

1. IGNORE

The dependency is ignored if the required job does not exist. If the job is present, it is handled in the normal way.

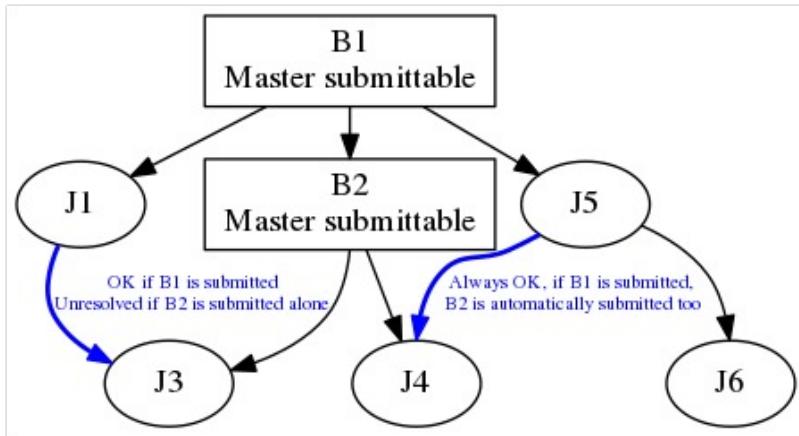


Figure 12.22: Example for Unresolved Handling

An Ignore may be necessary where smaller batches have been merged to form larger ones. If the small batch is running on its own, it can ignore all the dependencies that only have to be taken into account with larger batch runs. If it is running as part of the large batch, the required jobs are present and the dependencies have to be taken into consideration.

In the example above, an Ignore has to be entered in the dependency $J1 \rightarrow J3$ if the batch $B2$ can be started on its own as well.

2. ERROR

If the required job is not present, the submit action is terminated with an error.

3. SUSPEND

The current job switches to the Suspend State. This means that it has to be started either by an external program or manually as soon as the required job is available.

In the example above, Job $J3$ would switch to a Suspend State were $J1$ to be a dynamic child of $B1$ and it is not clear when the job is to be submitted.

The idea behind "Suspend" is that neither an unattended start nor an "ERROR" occur.

State Selection The *State Selection* field defines which Exit States for the required jobs fulfil the dependency and the dependent job can start. The following options are available:

1. FINAL

If this option is selected, the required job must have attained a Final State. This option allows the required Exit States to be selected in the following *Required States* table.

2. ALL REACHABLE

If this option is selected, all the Final Exit States that are not defined as being Unreachable in the Exit State Profile of the required job are valid.

3. UNREACHABLE

If this option is selected, only the Exit State that is defined as being Unreachable in the Exit State Profile of the required job is valid.

4. DEFAULT

If this option is selected, all the Final Exit States that are defined as being Dependency Default in the Exit State Profile of the required job are valid.

Required States This is a list of all the valid **Exit States** which the required object must have for the dependency to be fulfilled and so the dependent job can start. This list is only visible if the option FINAL has been selected in the *State Selection* field.

Example:

A Job A required by Job B can take on two Exit States (SUCCESS or WARNING). As Job B should only start if Job A has been completed with the SUCCESS Exit State (i.e. it has finished successfully), only a SUCCESS entry is allowed in the list of Required States. Once Job A has been completed and has received the SUCCESS Exit State, Job B can then start. If Job A finishes with WARNING, Job B cannot start. This list is necessary because a dependent job can accept multiple Exit States as a requirement. In our example, this means that Job B would have entered both Exit States (SUCCESS and WARNING) as Required States if it is to start in any case regardless of whether Job A was completely successful or it finished with a warning.

12.5.7 Dependents tab

The Dependents tab shows all the Scheduling Entities that are dependent upon the selected Scheduling Entity. This is the reverse view of the Dependencies tab as the dependent jobs are displayed here and not the required jobs.

The tab looks like this:

The list shows all the objects (jobs, batches), which require the selected Scheduling Entity. The list is purely informative and can neither be edited nor can the field values in this view be changed. To make any changes, you have to navigate to the list of displayed objects and open the Dependencies tab, where you can make the necessary changes.

The list columns are explained in the "Dependency Details" tab.

Editor for Job Definitions

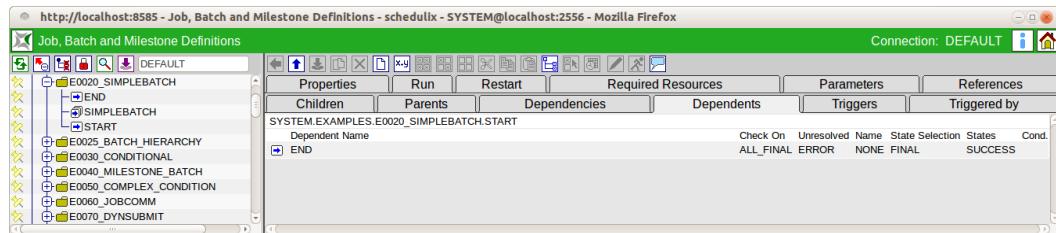


Figure 12.23: Batches and Jobs; Dependents tab

12.5.8 Required Resources tab

The Resources tab shows all the **Resources** that are required to execute this object. Additionally required resources can be added here.

All the resources that have already been defined by the **environment** or a specified **footprint** are displayed here. These can only be removed by changing the respective environment or footprint.

The Resources tab looks like this:

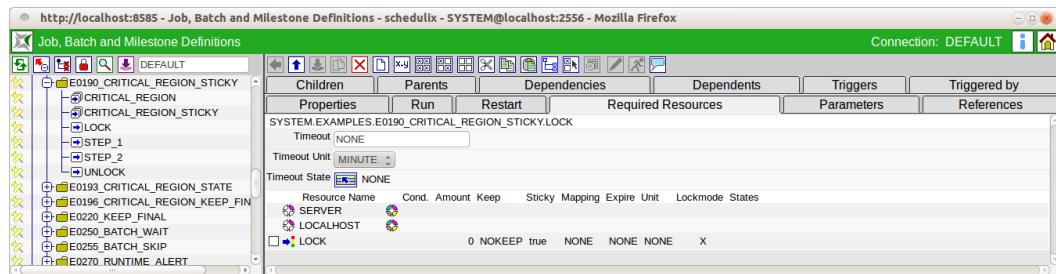


Figure 12.24: Jobs; Required Resources tab

The fields and columns in the "Required Resources" tab have the following meanings:

Timeout The *timeout* specifies how long the job is to wait for all the resources after it has been submitted. If no value is entered here, the job will wait for them indefinitely.

The value in this field must be regarded in conjunction with the *Timeout Unit* field. The value for the timeout is entered in the *Timeout* field. The unit is entered in the *Timeout Unit* field. Example: If the value 5 has been entered in the *Timeout* field and the timeout unit is **MINUTES**, a job will wait 5 minutes for a resource. If it does not become available within this time period, this job cannot be started and it terminates with the Exit State defined in the *Timeout State* field.

Timeout Unit The timeout unit is the unit that is to be used for the selected value in the *Timeout* field. If no value is entered in the *Timeout* field, the value in the *Timeout Unit* field is irrelevant.

Timeout State Here you can select the Exit State which the job is to take on should a timeout occur. By clicking the selection button you can choose those Exit States from the list that are to be made available to the *Exit State Profile*.

For example, you could define a separate Timeout State or use the Failure State. If no value is entered in the *Timeout* field, the value in the *Timeout State* field is irrelevant.

The "Resources" list shows all the resources that are currently required by the selected job for it to be executed. The list contains the following columns:

Resource Name This is the name of the resource that is needed to start the current job.

Next to the field name is an icon which indicates the context from which the resource originates. The following icons are used:

- Footprint



The resource originates from the *Footprint* field. An individual request can be created for each resource which overwrites the original request.

- Environment



The resource originates from the *Environment* field. This is only used for informational purposes and cannot be changed in this tab.

- No icon

If the resource name is not followed by an icon, the resource was added in this tab and can also be modified here. By clicking the name, you can switch to the Resource Details tab and change the values in the list. You can also click the *Selection* button in front of the name to mark this row and delete or move it.

12.5.8.1 Resource Details tab

The Resource Details tab opens when you select a list entry in the Required Resources tab.

All the details for an allocated resource can be entered and changed in this tab.

Editor for Job Definitions

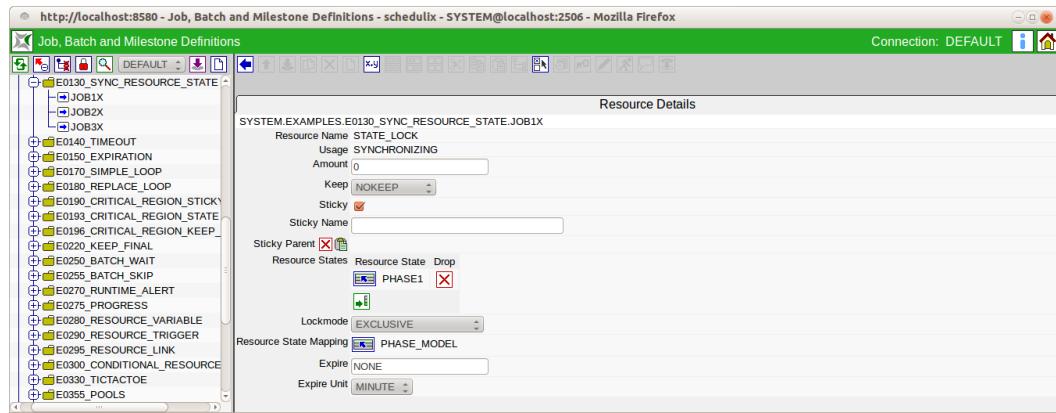


Figure 12.25: Batches and Jobs; Resource Requirement details

The tab looks like this:

The fields in the "Resource Details" tab have the following meanings:

Resource Name This is the name of the required resource. The full name together with the parent folder hierarchy is displayed here.

Usage This is the usage of the resource. More information about the usage can be found in Chapter 8.3.3.

Amount This field is only displayed if the usage type is either "SYSTEM" or "SYNCHRONIZING".

This is the amount that is required by the resource in this job. Maximum amounts of the resource are made available on the job server (or in the parent scopes, see [Job server](#)). Every job that requires an amount of this resource reduces the number of free amounts.

Example:

A job server makes the amount '5' available of Resource A (for example 5 CPU units on this server).

Job A starts now and requires the amount '3'. The job server has now occupied 3 CPU units, leaving 2 still available.

Job B wants to start, but it also requires an amount of '3'. The job cannot be handed over to the job server to be started because at this moment only 2 CPU units are available.

Job B can only start when Job A has finished and at least three CPU units are available.

Keep The value of the Keep parameter is set here. More information about the Keep parameter can be found in Chapter [10.3.1](#).

Sticky This switch is used to control how resources are returned. If this switch is set, the resource is retained within a Master Job until the last job that requires this resource with an activated Sticky Flag has been completed.

This functionality is for keeping reserved resources within a Master Job while preventing them from being used in the interim by other jobs.

Example:

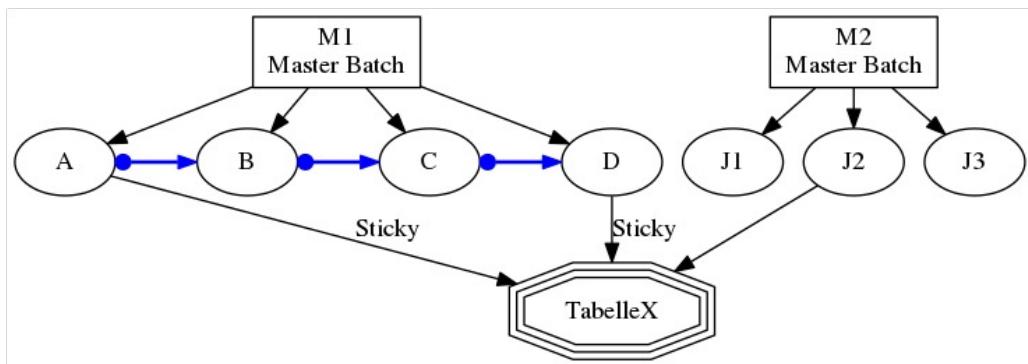


Figure 12.26: Example of Sticky Handling

Master Job M1 comprises the jobs A, B, C and D, which are executed in succession. Job A generates Table X, which is required by Job D. Jobs B and C don't need this Table X. Another Master Job (M2) also wants to access the table and modify it. However, this is to be prevented as long as Job D has not been executed because these actions would corrupt the results from Job A and make them unusable. Table X is now mapped as a Synchronizing Resource.

If a Sticky Flag is not set, the resource would be returned after Job A has been completed and the table could be modified by J2. With a set Sticky Flag (which has to be done in Job A and Job D), the resource is not returned following completion of Job A, but is retained in the Master Job as being allocated.

Other jobs cannot allocate this resource now. Only once Job D (the last job in this Master Job that requires the resource with a set Sticky Flag) has been completed is the resource released again and it can be used by other jobs.

With the Sticky Flag, all the necessary processes can be run on the same scope or job server using a Sticky Resource.

Here's another example to illustrate this:

In a system environment there are 2 UNIX workstations with an identical configu-

Editor for Job Definitions




Figure 12.27: Example of load distribution with Sticky Handling

ration and environment that are used for load distribution. A Master Job can run on one of these two workstations. However, since the intermediate results are to be transferred as data files between the individual jobs in the Master Job, the Master Job must run on the workstation after it has been selected.

This can be achieved by means of the Sticky Flag. Job servers (System A and System B) are installed respectively on these workstations and the environment is mapped as a Synchronizing Resource Definition (Resource Unix). This Resource Definition is now allocated within the two job servers (twice, once on each job server). Now all the children of the Master Job have to be defined so that they require the Resource Definition with the Sticky Flag.

The first job (Job A) in the Master Job now allocates one of the two available resources (on one or the other job server). Since this sticky has been allocated, it is not returned following completion of the first job. All the other jobs in the Master Job are defined as belonging exclusively to the selected job server, and therefore to its workstation, by virtue of this Sticky Resource.

In the diagram, one Master Job (Instance 1) is running entirely on job server System A and a second Master Job (Instance 2) is running entirely on System B.

Sticky name If the sticky flag described above is set, the Sticky Name field is displayed. If the sticky name is set, the resource is only available to those jobs in the same master run that want to allocate the resource with the same sticky name, but not to the other jobs.

Sticky parent If the sticky flag described above is set, the Sticky Parent field is displayed. If the sticky parent is set, the resource is only available to those jobs in the same master run that want to allocate the resource with the same sticky parent, but not to the other jobs.

Resource States list The Resource States list is only displayed if the resource has the type "Synchronizing" and a Resource State Profile has been allocated.

Editor for Job Definitions

A list of required Resource States can be entered here. The job can only be started if the required resource is in one of these states.

Lock mode The Lock mode indicates the access mode that is to be used to access the required resource. More information about the Lock mode can be found in Chapter [11.4.2.3](#).

Resource State Mapping This field is only displayed if the resource has the type "Synchronizing" and a Resource State Profile has been allocated.

The Resource State Mapping defines how and whether the state of the resource is to be changed after the job has finished. More information about Resource State Mappings can be found in Chapter [7](#).

Expire Gives the value of the Expiration interval. More information about Expire can be found in Chapter [8.3.6](#).

Expire Unit This is the unit used for the Expiration interval. More information about the Expire Unit can be found in Chapter [8.3.6](#).

12.5.9 Parameters tab

In the Parameters tab, you can define the parameters that can be used for communication and data transfer between jobs. Parameters that originate from the environment or from the job server's scope or folder environment and the job environment can be defined here for use within the current job.

The Parameters tab looks like this:




Figure 12.28: Batches and Jobs; Parameters tab

All the parameters that have been currently defined in this Scheduling Entity are shown in the list above. A new parameter can be created by clicking the *New* button. This new parameter is then shown in the list after you have entered the data

Editor for Job Definitions

and return variable. Selected parameters can be deleted as well by clicking *Delete*. You can switch to the "Parameter Details" tab by clicking the parameter name. The columns in the "Parameters" tab are explained in the next section.

12.5.9.1 Parameter Details tab

The Parameter Details tab opens when you select a parameter in the Parameters tab or you want to create a new parameter by clicking the *New* button. All the details for a parameter can be entered and modified here.

The Parameter Details tab looks like this:



Figure 12.29: Batches and Jobs; Parameter Details

The fields in the "Parameter Details" tab have the following meanings:

Parameter Name This is the name of the parameter. All the parameter values are exchanged between the individual jobs using this name. The parameter value is set by the first job and can be used in the second job under the same name.

Parameter Type This is the parameter type and indicates how the parameter is used.

Options for Parameter Type

1. PARAMETER

Parameters of the type "Parameter" can optionally be specified at the time of the submit if a default value has been defined. They have to be specified if a default was not defined. This only applies for the Master Batch parameters, however. All other parameters behave as if they were of the type "Import".

2. IMPORT

The value must only be resolvable at runtime. The usage of this variable is explicitly documented.

3. RESULT

This is the result of a job that can be exported to a follow-on job.

4. CONSTANT

This parameter is a constant that can no longer be changed. The value of the constant is entered using the field *Default Value*. This value is mandatory for parameters of the type "Constant".

5. EXPRESSION

An Expression parameter is a special form of parameter for parallel processing using dynamic child jobs. To obtain an aggregate of all the children in such cases, an aggregate function can be selected in the supplementary field that is displayed, and the child parameter that is to be aggregated is entered in the field "Expression Parameter Name".

If the Expression parameter is used at runtime, all the instances of the child job are searched for the value of this parameter and the corresponding aggregate function is performed.

Here is an example illustrating this:




Figure 12.30: Example using Expression parameters

Job A has a dynamic Child Job B, which is to be simultaneously submitted at runtime. The Child Jobs B write all the x-rows in a table in parallel. Child Job A now wants to know how many rows have been written in total by all the child jobs.

To achieve this, Job B must have defined a parameter (type Result) for instance called LINECOUNT, which contains the number of processed rows on completion of the job run.

In Job A, a parameter TOTALCOUNT is now created with the type "Expression", the expression option "SUM" and the parameter name "LINECOUNT".

Job A starts now and submits 3 Child Jobs B (B1, B2, B3). B1 writes 2000 rows, B2 writes 3000 rows, and B3 writes 1000 rows. When the children have finished, Job A (or a follow-on job) queries the TOTALCOUNT parameter and receives the result 6000 rows, which corresponds to the sum total of all the children.

6. REFERENCE

The type "Reference" allows a parameter to be allocated as a reference to another parameter. This parameter is performed with a cross-search of all the jobs on the same level as the current job. This enables you to obtain parameters for jobs which are not immediate parents of the current job. The parameter that is to be referenced must be inserted into the *Reference* field by clicking *Paste* (*Copy* from the Parameters tab of another Job Definition).

7. CHILDREFERENCE

The type "Childreference" allows a parameter to be allocated as a reference to another parameter. This referenced parameter is determined by running a search in all the children of the current job. This enables parameters from the job's own children to be identified.

8. RESOURCEREFERENCE

The type "Resourcerefence" allows a parameter to be allocated as a reference to a Resource Parameter. The parameter that is to be referenced must be inserted into the *Reference* field by clicking *Paste* (*Copy* from the Parameters tab of a Named Resource).

Expression This field is only visible if the value "Expression" has been selected in the field *Parameter type*.

The type of aggregate function can be selected here.

The following rule applies for all aggregate functions: Should the parameter or parameter value not exist or be of the wrong type (string instead of a number) in some child objects, these values are then ignored and only the available and correct values are aggregated.

The following aggregates are available:

1. SUM

The sum of all the values for the referenced parameter for all the child elements.

2. AVG

The average of all the values for the referenced parameter for all the child elements which contain valid and existing values.

3. MIN

The minimum of all the values for the referenced parameter.

4. MAX

The maximum of all the values for the referenced parameter.

5. COUNT

The number of child instances that have defined this parameter.

6. NONE

An aggregate function has not been selected yet. None is not a valid value.

Expression Parameter Name This field is only visible if the value "Expression" has been selected in the field *Parameter type*.

This is the name of the child parameter for which the aggregate is to be calculated. More information about aggregates can be found in the *Expression* field.

Local If a check mark is set here, this parameter is only visible to the job itself.

Default Value The *Default Value* specifies the default value. This value is returned when the parameter is resolved if an explicitly set value could not be found during the search.

This field is optional. A value must be entered here in the case of a constant.

Variables are evaluated recursively. This means that if a variable contains a string such as \$NAME or \${NAME}, the variable NAME is resolved and the value is inserted in its place. This means it has now also become possible - albeit indirectly - to access system variables for other jobs. The \$ character is represented by \\$, and \\ is used for a backslash.

Comment A more detailed explanation about the object can be entered in the *Comment* field.

Reference This field is only visible if the value "Reference", "Childreference" or "Resourcerefence" has been selected in the field *Parameter type*.

A parameter can be entered here with *Copy and Paste* which is referenced using the selected name.

12.5.9.2 Predefined default parameters for the runtime system

Predefined default parameters in the runtime system can be used by all Submitted Entities in the field *Run Program* as a transfer parameter or in the run program itself. The predefined default parameters contain information about the Submitted Entity and its environment.

The following default parameters are provided by the runtime system:

Job environment parameters These are all the parameters that relate to the Submitted Entity itself. Job environment parameter names:

1. JOBNAME

This is the name of the job that was entered in the *Name* field.

2. JOBID

This is the ID of the Submitted Entity. It enables the current instance of the entity to be unequivocally identified.

3. KEY

The *Key* is the password for the respective job. It is used to log onto the server.

4. MASTERID

The *MASTERID* is the ID of the Master Job that is executing the current job.

5. PID

The *PID* is the process ID of the Submitted Entity at operating system level. This can be used for the kill program, for instance.

6. LOGFILE

The *LOGFILE* parameter specifies the log file defined in the *Logfile* field in the **Batches and Jobs dialog** that is to be used.

7. ERRORLOG

The *ERRORLOG* parameter specifies the error log file defined in the *Error Logfile* field in the **Batches and Jobs dialog** that is to be used.

8. SDMSHOST

The *SDMSHOST* parameter specifies the current host of the SDMS server.

9. SDMSPORT

The *SDMSHOST* parameter specifies the current port of the SDMS server on the host.

10. JOBTAG

The JOBTAG specifies the Child Tag for the Submitted Entity. This enables a dynamically created child entity to be unequivocally identified.

11. SUBMITTIME

The SUBMITTIME specifies the time the Master Job is submitted.

12. STARTTIME

The STARTTIME is the time when the Submitted Entity was actually started on the job server (if the entity is a job).

13. SEID

The SEID is the ID of the Scheduling Entity.

14. WORKDIR

WORKDIR is the path to the working directory.

15. JOBSTATE

The JOBSTATE is the Exit State of the job or batch.

16. MERGEDSTATE

The MERGEDSTATE is the resultant Exit State of a job and its children.

17. EXPRUNTIME

The EXPRUNTIME is the Expected Runtime as stored in the Job Definition.

18. EXPFINALTIME

The EXPFINALTIME is the Expected Final Time as stored in the Job Definition.

19. PARENTID

The PARENTID is the ID of the parent job or batch.

20. STATE

The STATE is the current state of the job.

21. ISRESTARTABLE

ISRESTARTABLE states whether the job is restartable or not.

22. SYNCTIME

The SYNCTIME is the time for Synchronize Wait.

23. RESOURCETIME

The RESOURCETIME is the time for Resource Wait.

24. RUNABLETIME

RUNABLETIME is the time for Runnable.

25. FINISHTIME

FINISHTIME is the end time for the process.

26. SYSDATE

The variable SYSDATE is the current date and time.

27. LAST_WARNING

This variable contains the text from the last audit entry that resulted in the warning count being raised.

28. RERUNSEQ

This variable states how often the job has already been restarted.

29. SCOPENAME

This variable delivers the full path name to the job server that is executing or has executed the job.

Trigger parameters There are several predefined parameters that can be used in scripts for trigger jobs. These parameters are only available if the Submitted Entity was started as part of a trigger. If the entity was submitted in the normal way, these parameters are not fulfilled.

Trigger environment names:

1. TRIGGERNAME

This is the name of the trigger.

2. TRIGGERTYPE

This enables a job to perform different actions based on the trigger type.

3. TRIGGERBASE

The TRIGGERBASE parameter contains the name of the Triggering Object, i.e. of the job that sends a trigger to a parent by means of an Exit State Translation, for example.

4. TRIGGERORIGIN

The TRIGGERORIGIN parameter specifies the name of the entity that contains the trigger. This is the job that is shown the trigger in its Triggers tab.

Editor for Job Definitions

5. TRIGGERREASON

The TRIGGERREASON parameter specifies the name of the object that activates the trigger.

6. TRIGGERBASEID, TRIGGERORIGINID, TRIGGERREASONID

The ID of the Submitted Entity is saved in this parameter for each of these three objects (TRIGGERBASE, TRIGGERREASON, TRIGGERORIGIN).

7. TRIGGERBASEJOBID, TRIGGERORIGINJOBID, TRIGGERREASONJOBID

The ID of the Job Definition is saved in this parameter for each of these three objects (TRIGGERBASE, TRIGGERREASON, TRIGGERORIGIN).

8. TRIGGEROLDSTATE

This parameter specifies the Job State prior to the activating moment that is relevant to the trigger.

9. TRIGGERNEWSTATE

This parameter specifies the Job State after the activating moment that is relevant to the trigger.

10. TRIGGERSEQNO

TRIGGERSEQNO is the number of times the trigger was activated.

12.5.10 References tab

All references (parameters of the types REFERENCE and CHILDREFERENCE) from other job or batch objects to parameters for the current object are shown in the References tab.

The References tab looks like this:

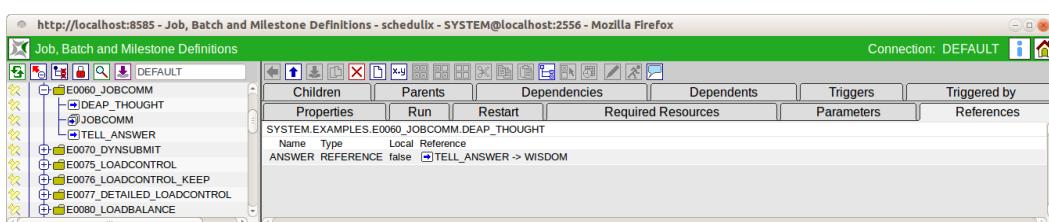


Figure 12.31: Batches and Jobs; References tab

12.5.11 Triggers tab

The Triggers tab describes all the triggers that are defined for this Scheduling Entity. A trigger is used to start another executable object if a configurable event has been reached in the currently selected job (or one of its children, see below). For example, messages could be sent when a job has been completed. Follow-up processing or error handling routines can also be started automatically.

The Triggers tab looks like this:




Figure 12.32: Batches and Jobs; Triggers tab

The "Triggers" dialog shows a list of all the triggers assigned for the selected task. The job to be started by the trigger can be selected and added to the list by copy and paste. The name for the trigger is taken from the name of the job when it is copied. The trigger details can then be viewed and modified by clicking the trigger's name. The trigger is deleted by clicking the *Drop* button. The columns in the list have the following meanings:

Name This is the name of the trigger. When you add a new job to the list of triggers, the trigger name is taken over from the job name. The Trigger Details tab is opened by clicking the trigger name. All the data (including the trigger name) can be viewed and modified here.

States This is a comma-separated list of all the states that are to be taken into account for the trigger.

Submit This is the name of the job or batch hat is to be submitted using the trigger. Important: This field can also be selected with a mouse click. Doing so opens the edit mask for the respective job. If you have added or modified any data and confirmed the confirmation query, all the modified data is lost.

The other columns are explained in the next section.

12.5.11.1 Trigger Details tab

The Trigger Details tab is opened by selecting a trigger in the Triggers tab. All the detailed information regarding the trigger can be modified or entered here.

Editor for Job Definitions

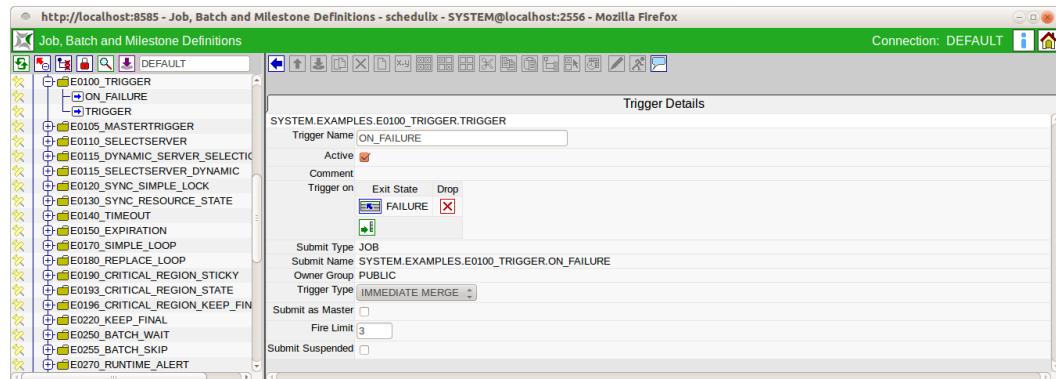


Figure 12.33: Batches and Jobs; Trigger Details tab

The tab contains the following fields:

Trigger Name This is the name of the trigger. When you create a new trigger (i.e. when you add a job to the trigger list), the job name is taken over for it automatically. The name can then be changed as required in the Trigger Details tab.

Active States whether the trigger is active. If this switch is not set, the trigger will not be activated when the trigger event occurs. An inactive trigger is indicated in the trigger list by a black hand (Suspend symbol).

Trigger On list A list of Exit States for the selected process (not the job that is to be triggered) for which the trigger is to be activated can be defined here. If no list is specified, the job's Exit State is irrelevant and the trigger is activated whenever the achievable event occurs.

Example:

A trigger is to be activated when a job finishes with an error (EXIT_STATE: FAILURE). In this case the Trigger On list must have selected the state FAILURE. Other states (SUCCESS, etc.) are not listed because they are to be ignored when the trigger is activated.

If a trigger is always to be activated when the process ends (regardless of the state, which is perhaps only differentiated in the job that initiates the trigger), the "Trigger on" list can be left empty or filled with any states that the job can take on.

Submit Type This is the type of the object (batch or job) that is started when the trigger is activated.

Submit Name This is the path and name of the object that is started when the trigger is activated. The full name of the object with all the parent folder hierarchies is used here.

Owner Group The owner of the trigger is defined in the field *Owner Group*.

Trigger Type This is the type of event that is to activate the trigger. The following events are used:

1. Immediate Merge

The job itself or one of its children has taken on a state in the Trigger On list. If no value is entered, the trigger is activated with each state transition.

Example:

A text message is to be sent if the job or one of its children contains an error.




Figure 12.34: Example of an Immediate Merge trigger

2. Immediate Local

The job itself has taken on a state in the Trigger On list. If no value is entered, the trigger is activated when the job finishes. The state of the children is ignored.

Example:

A text message is to be sent to the project manager if the job contains an error. If one of the job's children contains an error, a second trigger is implemented in the child which sends a text message to the responsible programmer.

3. Before Final

The trigger condition is evaluated shortly before the Scheduling Entity would reach a Final State. Here, the Trigger On list only needs to contain Final States (or none, in which case the trigger is activated for every Final State instance), as otherwise the trigger will not be activated. If the Scheduling Entity that is to be triggered is not submitted as a master, the Scheduling Entity can only

Editor for Job Definitions

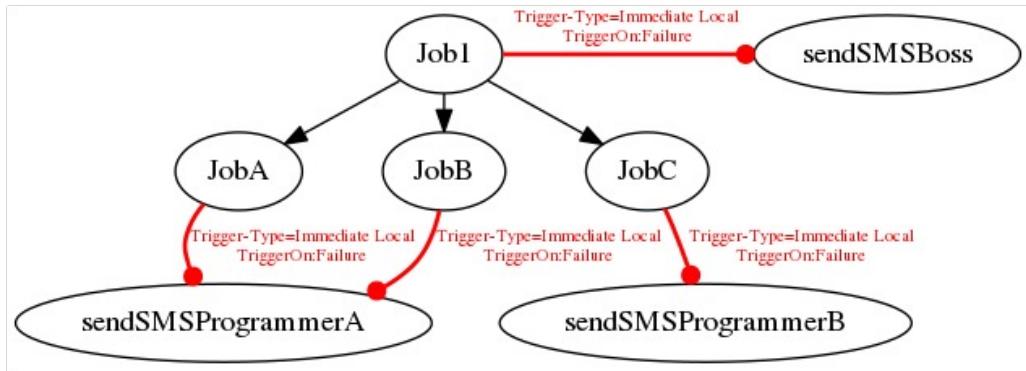


Figure 12.35: Example of an Immediate Local trigger




Figure 12.36: Example of a Before Final Trigger

attain a Final State after the Scheduling Entity started by the trigger has also attained a Final State.

In the example above, Job C will be called again if the job fails. Since this is a Before Final trigger, Job1 cannot attain a final state for as long as JobC has been terminated without a FAILURE.

Submit as Master If the "Submit as Master" flag has been set, the job that is started by the trigger is submitted as its own Master Job and does not have any influence on the current Master Job run of the triggering job. If the flag is not set, the triggered job is started as a child of the job to be triggered. This also means that the job in which the trigger is activated is not completed (cannot attain a Final State) as long as the triggered job is still running. The only exception is the trigger type "After Final", where the trigger is only activated after the triggered job has attained a Final State. If this is the case, the triggered job is started as a child of the

parent of the triggering job. If the triggered job is the same Scheduling Entity as that of the triggering job, the triggered job takes the place of the triggering job.

Fire Limit This defines the maximum number of trigger activations for this trigger within a workflow. When the maximum number of trigger activations is reached, no more triggers are activated and no more new jobs are started.

Submit Suspended Defines whether the triggered Scheduling Entity is to be started as being suspended.

Resume This field is also only displayed if the option *Submit Suspended* has been enabled. Here you can choose whether a workflow should be resumed automatically.

The following options are available:

- NO: disables this function.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Resume Time This field is only displayed if "AT" was selected for *Resume*.

The required resume time is entered here in the format YYYY-MM-DDTHH:MI:SS. This format is based on the ISO standard 8601 and also permits an incomplete entry. If you enter 'T16:00', the job will be resumed at 16:00 hours (starting from the time the trigger was activated).

Resume In This field is only displayed if "IN" was selected for *Resume*. Here you can specify how many time units (see *Unit*) the system is to wait for until the resume action is triggered.

Unit This field is only displayed if "IN" was selected for *Resume*. This field is used to define whether the entry in *Add Resume* is in minutes (MIN), hours (HOUR) or days (DAY).

12.5.12 Triggered by tab

The "Triggered by" tab lists all the triggers that submit this Scheduling Entity.

The "Triggered by" tab looks like this:

The columns in the list have the following meanings:

Name This is the name of the trigger.

Job Hierarchy navigation




Figure 12.37: Jobs and batches; Triggered By tab

States This is a comma-separated list of all the states that are to be taken into account for the trigger.

Triggered by This is the name of the job or batch that defines the trigger.

Important: This field can also be selected with a mouse click. Doing so opens the edit mask for the respective job. If you have added or modified any data and confirmed the confirmation query, all the modified data is lost.

The meanings of the other columns can be found in the chapter "Trigger Details tab".

12.6 Job Hierarchy navigation

The Job Hierarchy navigation window can be viewed for each batch and job in the "Batches and Jobs" dialog. It looks like this:




Figure 12.38: Batches and Jobs; hierarchy view

All the children of the current entity are displayed in the navigation window. If the children also have a child hierarchy, this can be expanded as well until the entire hierarchy of the object is visible.

The standard buttons are supplemented by the following buttons:



Add Children

Children can be inserted directly in the Job Hierarchy navigation window by clicking the *Add Children* button. To do this, click the button followed by the job or batch

Job Hierarchy navigation

that is to acquire the children. A mask opens where you can select one or more children. Having selected the children, close the mask and then click the parent again. The children are added to it automatically.



Remove Child

Children can be deleted by clicking the *Remove Child* button. To do this, click the button and then the parent. The children are deleted automatically after you have clicked "OK".



Show Dependencies

The *Show Dependencies* button has two functions. By default, the dependencies are not displayed in the dialog window and the *Show Dependencies* button appears in the button bar.

The mutual dependencies between the Scheduling Entities can be graphically displayed by clicking the *Show Dependencies* button. Each dependency is indicated by an arrow. The required Scheduling Entity is represented by the start of the arrow (the round part of it) and the dependent Scheduling Entity is the arrowhead.

Example:

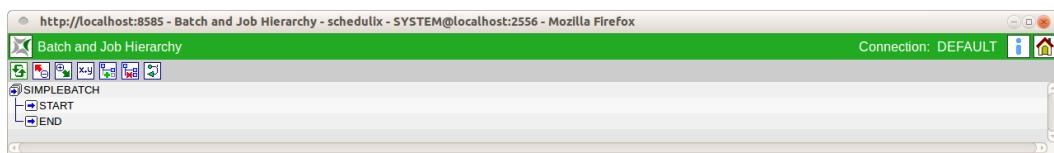


Figure 12.39: Hierarchy view with dependencies

In the example above, the job END is dependent upon the job START.

If a job is dependent upon multiple objects, an arrow is displayed with several shafts (the round parts of the arrow) representing all the required Scheduling Entities. The arrowhead remains the dependent Scheduling Entity.

Example:




Figure 12.40: Hierarchy view with multiple dependencies

In this example, the batch LOAD_COMPLETED is dependent upon the LOAD_1,

Job Hierarchy navigation

LOAD_2 and LOAD_3.



A *Hide Dependencies* function is the second state of this button. This appears if the *Show Dependencies* button was clicked causing the display of the dependencies to be activated. The dependencies view can be disabled again by clicking the *Hide Dependencies* button.



With the *Chain* button you can switch to the *Chain* mode, where job dependencies can be added directly in the Job Hierarchy navigation window. To do this, first click the *Chain* button followed by the Dependent Job and the Required Job. This creates a dependency relationship between the two jobs.

You can exit the *Chain* mode by clicking the *Cancel* button.



The *Unchain* button is used to revoke dependency relationships. To do this, first click the *Unchain* button followed by the dependency that is to be revoked.




Figure 12.41: Cyclic dependencies

Cyclic dependencies can be easily created with the *Chain/Unchain* buttons. These are identified and reported by the front end. See Figure 12.41.

13 schedulix!Web Users

13.1 View

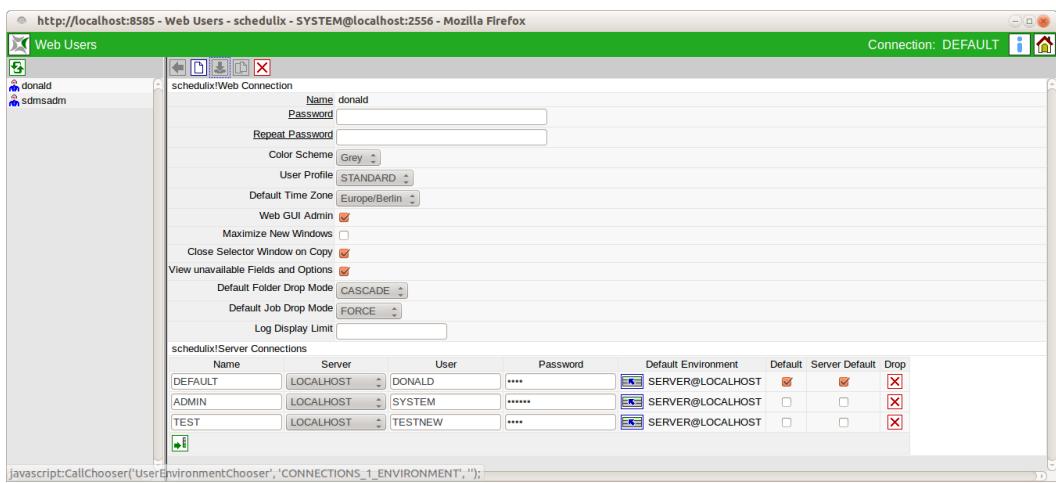


Figure 13.1: User management for the Web front end

13.2 Concept

13.2.1 In short

The user management module is used to define schedulix interface users. All users who want to work with the schedulix interface should have their own user ID. The user must log in with this user ID and a personal password in the [Login window](#).

13.2.2 Detailed description

In addition to the display options and security mechanisms, the connection to the server can also be maintained in this dialog window. The schedulix Server uses its own user management functionality which is independent of the user management module for the interface system. For this reason, it is possible for all interface users to utilise just one server user, or the interface users can be mapped 1:1 on the server

as well. Always notify your system administrator when making changes to your user account.

13.3 Navigator

All the user IDs are displayed in the navigation window.

13.4 Editor

The editor is used to maintain all the data for a selected user.

13.4.1 schedulix!Web Connection

The fields under Web-Gui-Connect have the following meanings:

Name The user's login name.

The login name is entered as the user name in the [Login](#) dialog.

Password The *Password* authenticates each user who logs into the schedulix System and has to be entered in the [Login](#) dialog. The password is hidden on the screen and must be identical to the Repeat password.

Since the password is not displayed in the dialog, it is not evident whether a password has already been entered or not. It only has to be entered once, however.

Repeat Password The *Repeat Password* must be the same as the password you entered in the Password field.

It has to be entered twice because a typing error when entering it just once can cause problems. Since it is unlikely that the user will make the same typing error twice, it is assumed that if the password is twice entered correctly, then the intended term has been used.

Colour Scheme The colour scheme for the user interface is selected here.

User Profile The user's profile is entered here.

Default Time Zone This parameter defines the default time zone when creating new time schedules. All time stamps continue to be displayed in this time zone.

Time Stamp Format A format other than the default format (%Y.%m.%d %H:%M:%S example 2014.06.18 14:06) for displaying the time stamps can be set here. The format corresponds to the strftime C standard function.

Editor

User Profile The user's profile is entered here.

Web-Gui-Admin This check mark must be set if this user is an administrator for the web interface.

Maximize New Windows New dialog windows are always maximized when opening them if this check mark has been set. If it is not set, the window always opens in the normal size.

Close Selector Window on Copy If this check mark is set, windows that are opened for selecting objects are closed automatically after clicking the "Copy" button.

View unavailable Fields and Options If this check mark is set, the unavailable fields and options are displayed.

Default Folder Drop Mode This is the mode that is used for deleting a folder in the **Batches and Jobs** dialog window. The following options are available:

1. NORMAL

In this mode it is only possible to delete a folder if it is empty, i.e. it does not have any subfolders and no jobs or batches.

This is the default value for the system.

2. CASCADE

In this mode, the folder is deleted in a cascade. All the subfolders and all the jobs and batches in them are deleted.

Important: This setting is to be used with extreme caution because it is possible to delete a large number of objects at once without actually intending to do so.

Default Job Drop Mode This is the mode that is used for deleting a job (or a batch).

OPTIONS

1. NORMAL

In this mode, it is only possible to delete a job or batch if no other jobs are dependent upon it. The Scheduling Entity must not have any more child relationships either.

This is the default value for the system.

2. FORCE

In this mode, the scheduling entity is deleted even though other Scheduling Entities are dependent upon it and it still has some child relationships.

13.4.2 schedulix!Server Connections

One or more server connections can be configured here.

On the Main Desktop, the server connection that is to be used for the next window to be opened can be selected in an Options field in the window header.

After editing the server connections, it may be necessary to reload the Main Desktop in the browser to make sure that the server connections are displayed correctly in this field.

Name The name of the server connection for selecting it in the header of the Main Desktop.

Server The name of the server as entered by the schedulix!Web system administrator under /Custom/SDMSServers (in Zope).

User The name of the schedulix Server user account. This does not have to be the same as the Web user's name.

Password This user ID and password are used to connect the schedulix interface with the schedulix Server.

The password is hidden on the screen.

Default Environment The Default Environment is entered by default for jobs.

Default If this check mark is set, this server connection is used as the default setting. It must be set for precisely one row.

Server Default If this check mark is set, this server connection is used as the default setting only if the server is known. It must be set for each server in precisely one row.

14 schedulix Server Users

14.1 View

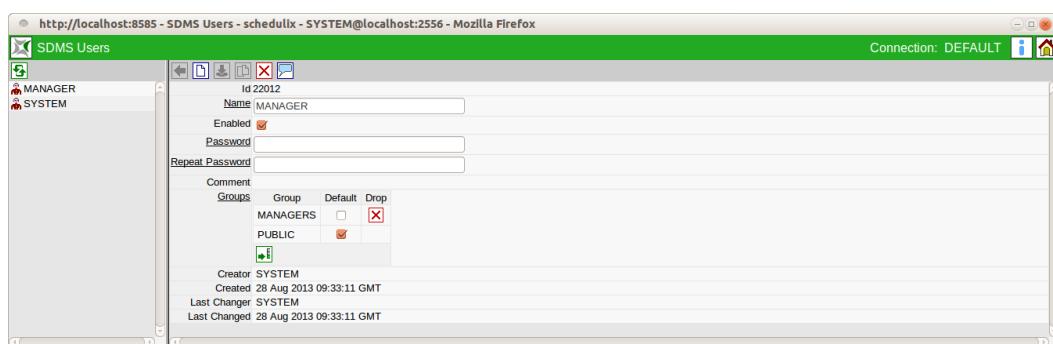


Figure 14.1: schedulix Server user management window

14.2 Concept

14.2.1 In short

The schedulix Server user is required for legitimisation toward the Scheduling Server.

14.2.2 Detailed description

Since the Scheduling Server is not just operated using the web interface described in this documentation, but also with the aid of other interfaces such as *sdmsh*, access controls are required on the server. The users and their group memberships can be maintained using the dialog described here.

14.3 Navigator

All the users are displayed in the navigation window.

The Navigator pane is only displayed if the active user is a member of the "ADMIN" group. All other users will only see the editor that can be used to change their own password and the default group they need to choose.

14.4 Editor

The editor is used to maintain all the data for the selected or new user. SDMS users can only be edited by users who belong to the "ADMIN" group. All the input fields are "read only" for all other users.

The editor contains the following fields:

Name The user's login name.

The login name is entered as the user name in the [Login](#) dialog.

Enabled If the active user is a member of either the "ADMIN" group, this option is used to define whether the selected user is permitted to log in.

Password The *Password* authenticates each user who logs into the schedulix System and has to be entered in the [Login](#) dialog.

The password is hidden on the screen and must be identical to the *Repeat Password*. Since the password is not displayed in the dialog, it is not evident whether a password has already been entered or not. It only has to be entered once, however.

The maximum length for a password is 64 characters. All characters (i.e. including spaces) are allowed.

The password can always be changed if the active user is a member of the "ADMIN" group or if the user to be modified is the active user (you can obviously change your own password).

Repeat Password The *Repeat Password* must be the same as the password you entered in the *Password* field.

It has to be entered twice because a typing error when entering it just once can cause problems. Since it is unlikely that the user will make the same typing error twice, it is assumed that if the password is twice entered correctly, then the intended term has been used.

Groups The groups to which the user belongs are listed in the "Groups" table. The table features the following columns:

Group Groups to which the specified user member belongs.

Default The default group is the one that is joined by a new object in the absence of a defined group.

Editor

Drop This button is used to delete the group membership. This action has to be confirmed.

15 Groups

15.1 View




Figure 15.1: Group management

15.2 Concept

15.2.1 Description

The groups are the 'legal entity'. The individual users belong to groups and they have the same privileges held by the groups. A user can belong to several groups. The following groups always exist:

- Public: Every user is always a member of the *Public* group.
- Admin: Each user in the *Admin* group holds all privileges throughout the system.

15.3 Navigator

All visible groups are shown in a list in the Navigator. Visible in this context means that the user is either a member of the group or he has been granted "view" privileges for it.

15.4 Editor

15.4.1 Properties tab

This tab is used for maintaining the group properties.




Figure 15.2: Group properties

This window contains the following fields:

ID The Group ID is shown in this field.

Name The name of the group is shown in the *Name* field.

Users The following field is in the Users list.

Name The group user is shown in the *Name* field.

15.4.2 Grants

The Grants tab shows all the objects to which group privileges have been assigned. The object type is indicated by the respective icon. The list is always sorted by the object names.

Moving the mouse pointer over the privileges field in a row opens a popup window with an explanation of the privileges string.

It looks like in Figure 15.3.

Editor

The screenshot shows a web-based application titled "Groups" with a URL of "http://localhost:8585 - Groups - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The interface includes a toolbar with icons for back, forward, search, and refresh, and a status bar indicating "Connection: DEFAULT". On the left is a tree view of groups: ADMIN, MANAGERS, PUBLIC, REPORTING_DEVELOPMENT, REPORTING_MONITORING, and REPORTING_OPERATIONS. The central area has tabs for "Properties" and "Grants". The "Properties" tab shows the "Object Name" as "PUBLIC". The "Grants" tab lists various objects and their privileges:

Object Name	Privileges
E0080_HOSTS	UV
E0115_ALLOWED_HOSTS	UV
E0115_HOSTS	UV
E0115_HOSTS_BROKEN	UV
SERVER@HOST_1	UV
SERVER@HOST_2	UV
SERVER@LOCALHOST	UV

Figure 15.3: Object privileges for a group

16 Time Scheduling

16.1 View

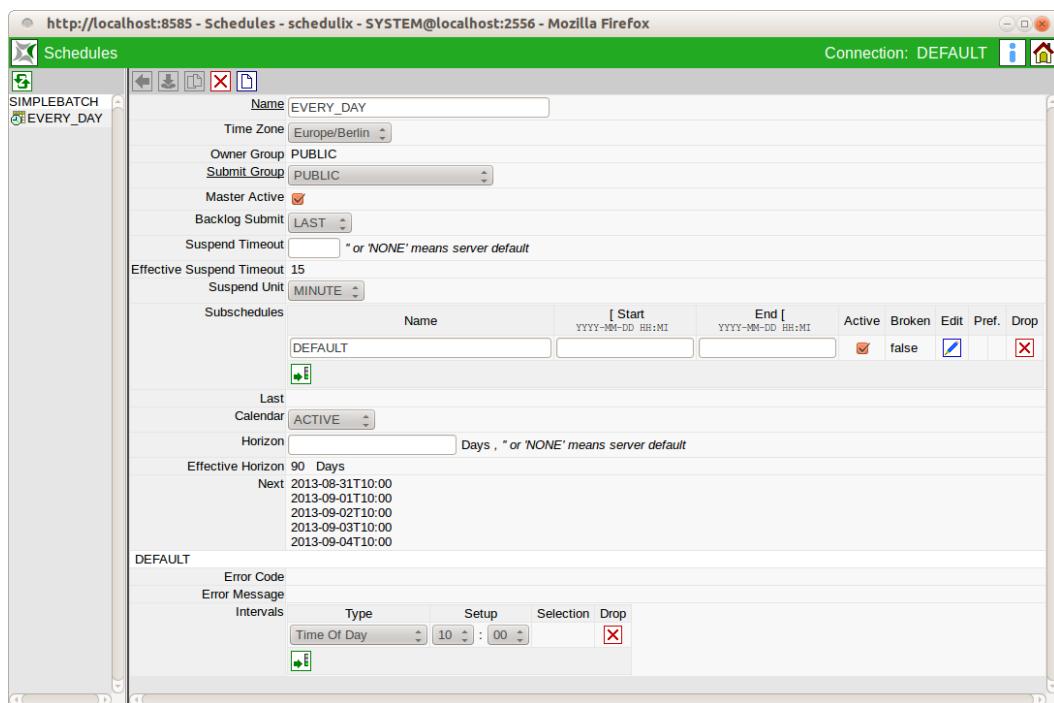


Figure 16.1: Time Scheduling

16.2 Concept

16.2.1 In short

The Time Scheduling module allows batches and jobs to be started at specific times and can be used to easily define complex schedules for jobs and batches.

A number of points in time which determine when a job is to be started is called a schedule.

Schedules always apply to just one Master Job.

16.2.2 Detailed description

In the "Time Scheduling" dialog window, a differentiation is made between Main Schedules, which are shown in the navigation as elements, and Sub Schedules, which can be administered within a Main Schedule.

Main Schedules are mutually independent and no interaction takes place.

Sub Schedules can be defined as being separated from the time interval or as overlapping. This allows different downtimes to be simulated or varying schedules can be set which are therefore always active.

If the schedules are to be changed on demand, this can be entered as an interim schedule. The highest priority has to be assigned to the interim plan to activate it. The original schedule is retained. The original schedule is reactivated by reassigning the highest priority to it.

16.3 Navigator

The Main Schedules for this Scheduling Entity are displayed in the navigation. A Scheduling Entity can have multiple Main Schedules. These Main Schedules are taken into account independently of one another. This means that the start times of all the Main Schedules are factored in.

If multiple Main Schedules have identical start times, the Scheduling Entity is started several times.

16.4 Editor

All the "Detail" elements for the Main Schedule selected in the navigation are displayed in the editor. Here is where Sub Schedules can be created and a list of times maintained for when the selected job or batch is to be started. The "Editor" mask looks like this:

The fields shown above have the following meanings:

Name This is the name of the Main Schedule. When you first create a new Main Schedule, it is automatically assigned the name "MASTER", but this can be arbitrarily changed. Each Main Schedule that you create has to have a different name.

Time Zone The time zone for which the schedule is to be calculated. Which time zones are made available here can be configured by the schedulix!web administrator.

Owner Group The owner of the schedule is defined in the field *Owner Group*.

Submit Group The *Submit Group* determines the owner group of the submitted job.

Editor




Figure 16.2: Time Scheduling Navigator

Master Active If the flag for Master Active is set, the schedule is taken into account by the Time Scheduling System. If the flag is not set, the schedule is ignored.

Backlog Submit The backlog determines how the Time Scheduling System behaves after a server downtime:

- None: The "missed" jobs are no longer executed.
- Last: Only the last of the "missed" jobs is executed.
- All: All the "missed" jobs are executed.

Suspend Timeout, Suspend Unit Following a server downtime, a suspended job is submitted if the scheduled submit time is longer than the time entered in the fields *Suspend Timeout* and *Suspend Unit*. If no time is entered, the server-wide default time applies.

Sub Schedules list All the Sub Schedules of the Main Schedule are shown in the "Sub Schedules" list. The schedules are only valid within the stated timeframe. The Sub Schedules with the highest preference (at the top of the list) are taken

Editor

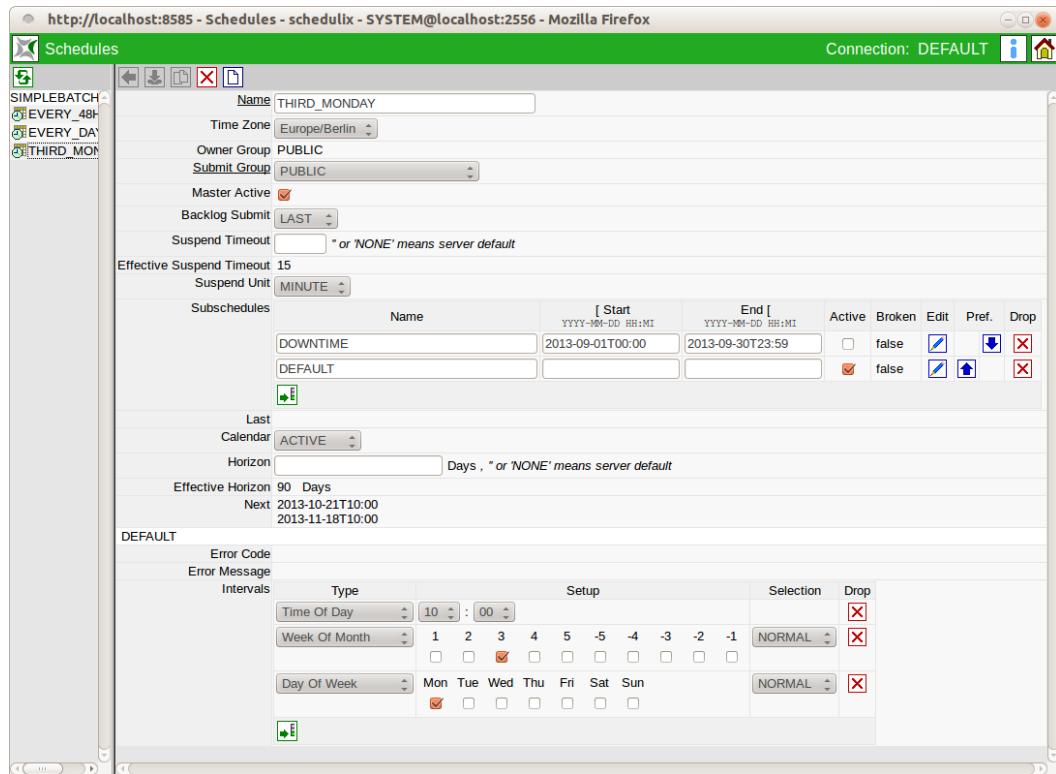


Figure 16.3: Time Scheduling Editor

into account first. These 'cover' the Sub Schedules further down the list for their timeframe. This means that only one Sub Schedule can be valid at any one time. The fields in the list have the following meanings:

Name The name of the Sub Schedule is entered in the *Name* field.

Start The date and time from which this Sub Schedule is valid are entered in the *Start* field. If this field is left empty, the Sub Schedule is considered to be "always valid".

End The date and time from which this Sub Schedule is invalid are entered in the *End* field. If this field is left empty, the Sub Schedule is considered to be "indefinitely valid".

The following buttons are available in the Sub Schedules list:

Editor

Active Jobs are only submitted if this flag is set. Consequently, submits can be prevented by deleting the flag. Downtimes are usually implemented as high-preference Sub Schedules with a non-set Active Flag.

Broken The "Broken" field is used to check whether an error occurred when the job was submitted. If this is the case, TRUE is entered in the *Broken* field for this Sub Schedule.

Edit Clicking the *Edit* button opens the interval description belonging to the Sub Schedule, where you can precisely define when the job is to be started.

Pref. The Preference buttons are used to set the preference of the individual Sub Schedules. The rows are shifted up or down by clicking the buttons.

Drop This button is used to delete the object.

Last The last time the job is to be executed by the Scheduling System is shown in the *Last* field. This field is empty if the job has never been started by the Scheduling System.

The date format looks like this:

YYYY:MM:DDTHH:MM

The placeholders correspond to those in the *Start* field.

Calendar This field determines whether the scheduled execution times for the job are to be noted in the calendar. If this is the case, the next two fields are also important.

Horizon The "Horizon" field determines for how far in advance entries are to be made in the calendar. The time period is entered in days.

Effective Horizon The configured default value applies as the horizon if the "Horizon" field is left empty. The "Effective Horizon" field shows the currently valid value.

Next The next scheduled time(s) for the task to be executed by the Scheduling System is/are shown in the *Next* field. If no schedules or times have been defined, this field is left empty. Once a schedule has been created and times have been defined, the closest possible date after this schedule is displayed. If no time is possible because mutually exclusive times have been defined, a value is not displayed here either.

When calendar entries are created, both the next execution time and the following four execution times (providing they lie within the horizon) are displayed.

The date format is identical to the *Start* field.

16.4.1 Sub Schedule Details sub-area

This sub-area shows all the "Details" information for the first Sub Schedule or a Sub Schedule you have selected by clicking the *Edit* button. The Sub Schedule Details begin in the mask with a grey line showing the name of the currently selected Sub Schedule followed by the following fields:

Error Code If an error occurred while the job was being executed in the Time Scheduling, the returned error code is displayed in the *Error Code* field. If no error occurred, this field remains empty.

Error Message If an error occurred while the job was being executed in the Time Scheduling, the returned error message is displayed in the *Error Message* field. If no error occurred, this field remains empty.

Intervals list If no interval is entered here (the list is empty) and the Sub Schedule is marked as being active, a job is submitted precisely once at the start time of the Sub Schedule.

A list that is not empty must contain exactly one "driving" interval of the type "Time Of Day", "Repeat" or "Calendar (Driver)". The "driving" interval defines the cycle in which a submit can potentially take place. The other "filtering" intervals determine by virtue of their filtering effect which of the cycles created by the "driving" interval actually lead to a submit.

Type The *Type* field defines the type of interval. Another mask is displayed in the *Setup* column depending on the interval type.

Setup A mask for configuring the interval is displayed here depending on the interval type.

Selection In the case of intervals that can be selected (e.g. "Day Of Week"), in this column you can choose either a "NORMAL" or "INVERSE" selection.

The following interval types are used:

- Repeat

This is a "driving" interval for starting the task at regular intervals according to a schedule. The number of minutes can be entered in the *Setup* field. It looks like in Figure 16.4

In the example, the job is started every 60 minutes.

- Time of Day

Editor




Figure 16.4: Repeat Driver

This "driving" interval type can be used to start the job at a specified time of day. The time of day is entered in the *Setup* field by selecting the hour (24H) and minutes. The option looks like in Figure 16.5.




Figure 16.5: Time Of Day Driver

In the example, the job is started every day at 10:00 hours.

Several "Time Of Day" intervals can be created. This allows you to easily define multiple execution times for any one day.

- Range of Day

This "filtering" interval type is only permitted in combination with the "driving" interval types "Repeat" and "Calendar (Driver)". It allows you to set a restriction to one or more time ranges during a day.




Figure 16.6: Range of Day Filter

Editor

Several "Range Of Day" intervals can be created. This allows you to easily define multiple execution time ranges for any one day.

- **Day of Week**

The "filtering" interval type "Day of Week" can be used to select the weekdays on which the job is to be executed. This is defined in the *Setup* field by choosing the relevant weekdays. The option looks like in Figure 16.7




Figure 16.7: Day of Week Filter

In the example, the job is executed every Monday, Wednesday and Sunday.

- **Day of Month**

The "filtering" interval type "Day of Month" can be used to select the days of the month on which the job is to be executed. This is defined in the *Setup* field by choosing the relevant days of the month. The option looks like in Figure



Figure 16.8: Day of Month Filter

16.8

In the example, the job is executed on the 9th, 13th, 19th, 26th and penultimate day of each month.

Editor

- Week of Month

The “filtering” interval type “Week of Month” can be used to select the week of the month in which the job is to be executed. This is defined in the *Setup* field by choosing the relevant weeks in the month. The option looks like in

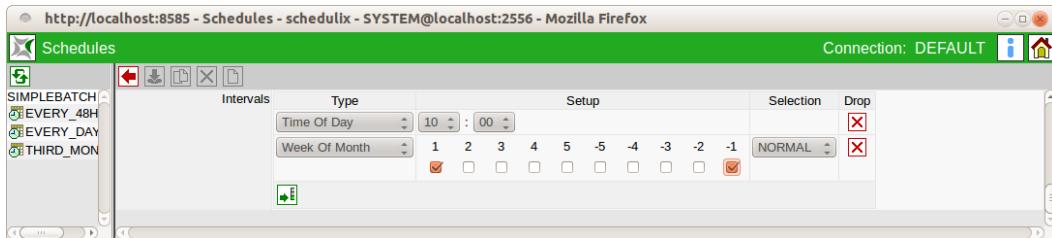


Figure 16.9: Week of Month Filter

Figure 16.9.

In the example, the job is executed in the first and last weeks of each month. The first week is defined here as being the first 7 days of the month and the last week is the last 7 days of the month.

- ISO Week of Month

The “filtering” interval type “ISO Week of Month” can be used to select the ISO week of the month in which the job is to be executed. This is defined in the *Setup* field by choosing the relevant weeks in the month. The option

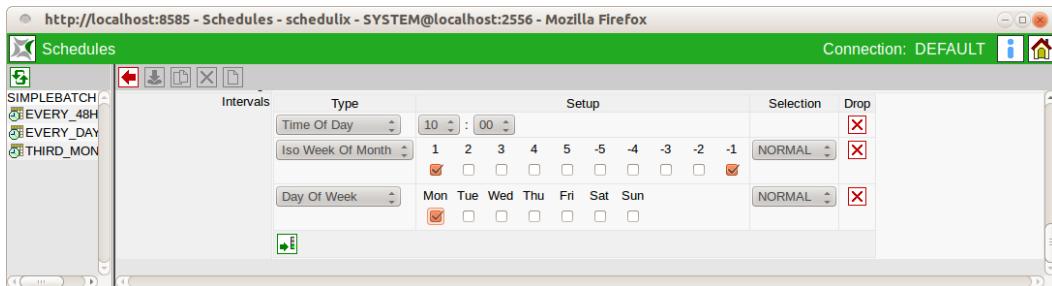


Figure 16.10: ISO Week of Month Filter

looks like in Figure 16.10.

In the example, the job is executed on the first Monday of the first ISO week of each month. ISO weeks always begin on Mondays. Weeks are assigned to a months if at least 4 days of the week lie in this month. For example, if

Editor

the first weekday of a month is a Wednesday, the first ISO week of the month already begins in the preceding month.

- ISO Week of Year

The "filtering" interval type "ISO Week of Year" can be used to select the ISO calendar week of the year in which the job is to be executed. This is defined in the *Setup* field by choosing the relevant weeks in the year. The option looks like in Figure 16.11

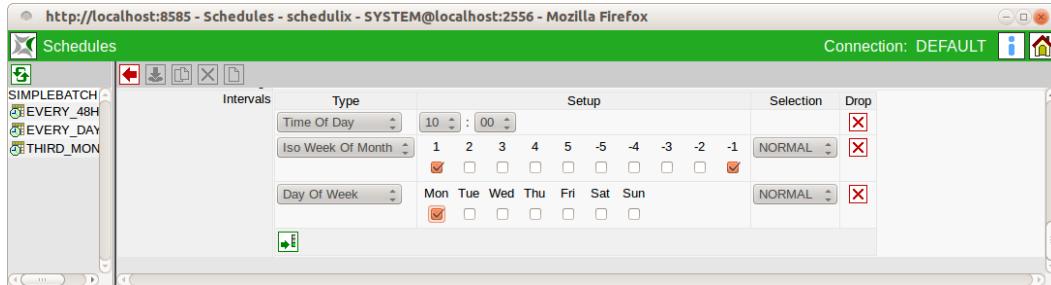


Figure 16.11: Week of Year Filter

like in Figure 16.11

In the example, the job is started on the Tuesday of the 31st ISO calendar week. ISO weeks always begin on Mondays. Weeks are assigned to a year if at least 4 days of the week lie in this year. For example, if the first weekday of a year is a Wednesday, the first ISO week of the month already begins in the preceding year.

- Month of Year

The "filtering" interval type "Month of Year" can be used to select the month of the year in which the job is to be executed. This is defined in the *Setup* field by choosing the relevant months of the year. The option looks like in Figure 16.12




Figure 16.12: Month of Year Filter

16.12

Editor

In the example, the job is started every day in July at 13:00 hours.

- **Calendar (Driver)**

The “driving” interval type “Calendar (Driver)” can be used to take execution times from a predefined calendar. This option is only available only if the administrator has created such a calendar and registered it in schedulix!web.

The calendar can then be selected in the *Setup* field. The option “with select on” determines the unit (DAY, WEEK, MONTH, YEAR) to which the subsequent selection refers.

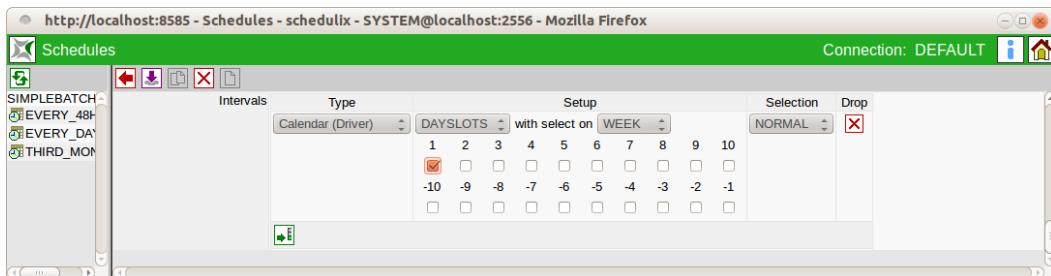


Figure 16.13: Calendar Driver

In the example above, the first entry in each calendar week is selected from a calendar called “DAYSLOTS”.

- **Calendar (Filter)**

The “filtering” interval type “Calendar (Filter)” can be used to restrict the execution times using a predefined calendar. This option is only available only if the administrator has created such a calendar and registered it in schedulix!web.

The calendar can then be selected in the *Setup* field. The option “with select on” determines the unit (DAY, WEEK, MONTH, YEAR) to which the subsequent selection refers.

In the example above, a submit is performed every 60 minutes if this is in the second block of each day in the calendar “DAYSLOTS”.

Setup The *Setup* field defines the interval more precisely. It varies according to the type of interval.

Editor




Figure 16.14: Calendar Filter

17 Submit Batches and Jobs

17.1 View

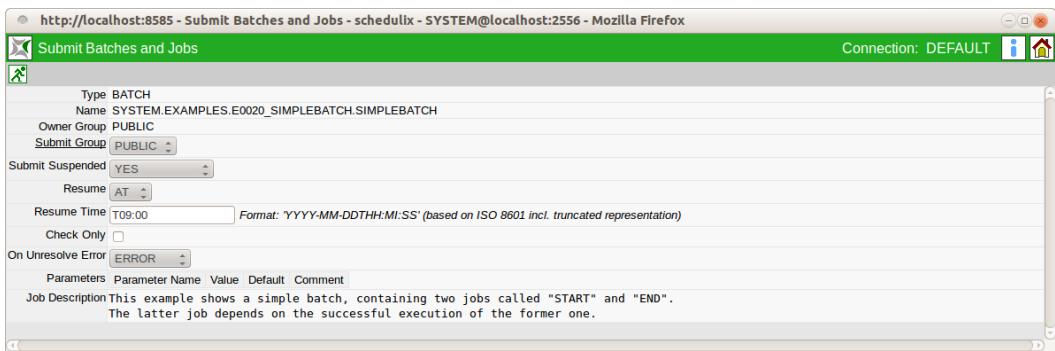


Figure 17.1: Submit Batches and Jobs

17.2 Concept

17.2.1 In short

The "Submit Batches and Jobs" dialog is used for manually starting Master Jobs defined in this dialog. The submit informs the schedulix Server that this Master Job is to be activated.

17.2.2 Detailed description

Manually releasing Master Jobs is one of two methods to start a job. The second method uses the job entry in the [Time Scheduling](#).

Starting a job manually also allows you to set certain parameters and determine whether you really want to start the job or just check if it can actually be started. You can also define whether the Master Job is to be suspended straight away.

Editor

17.3 Navigator

The navigation shows all the available Master Jobs in a folder hierarchy. It corresponds to the navigation pane in the Batches and Jobs dialog window. One difference here, though, is that only objects for which the flag "Submit as Master allowed" has been set are displayed.

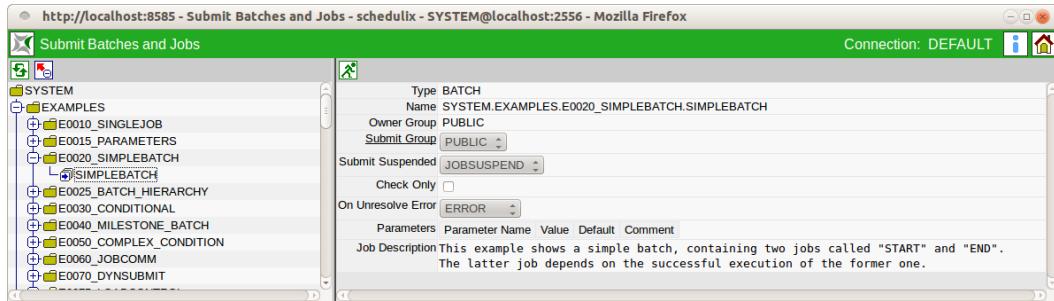


Figure 17.2: Submit navigation

17.4 Editor

The Editor Frame can be used to hand over start options and parameters and submit a Master Job.




Figure 17.3: Submit editor

The following button is active in this mask.



The Submit button is used to start the job (the flag "Check Only" is not set) or check the start capability of the Master Job (the flag "Check Only" is set).

The fields shown above have the following meanings:

Type The type of the Master Job. This can be either a job or a batch.

Name This is the full name (with its path) of the Master Job.

Owner Group The owner of the job is defined in the field *Owner Group*.

Submit Group The *Submit Group* determines the owner group of the submitted job.

Submit Suspended This field defines whether the job is to take on a Suspend State right after the start or not.

The input fields *Resume*, *Resume Time*, *Resume In* and *Unit* are only displayed if this field has been set to 'YES'.

The following options are available:

- YES

The child is created as being suspended when it is submitted and it has to be started by being explicitly released.

- NO

The child is not suspended when it is submitted and can be started immediately.

- JOBSUSPEND

Whether a delay takes place or not depends on the *Suspend* field for the job to be submitted. This means that the setting defined in the job is applied.

Resume Here you can choose whether a workflow should be resumed automatically.

The following options are available:

- NO: deselects this functionality and no more input fields are displayed.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Resume Time The required resume time is entered here in the format YYYY-MM-DDTHH:MI:SS.

This format is based on the ISO standard 8601 and also permits an incomplete entry. If you enter 'T09:00', the job will be resumed at 09:00 hours (starting from the current set time).

Resume In Here you can specify how many time units (see *Unit*) the system is to wait for until the resume action is triggered.

Unit This field is used to define whether the entry in *Add Resume* is in minutes (MIN), hours (HOUR) or days (DAY).

On Unresolve Error When a job is submitted, the *On Unresolve Error* field determines what to do in the event of unresolved dependencies.

The following options are available:

- Error: Dependencies that cannot be resolved are regarded as errors. The missing dependencies cause the submit to be aborted.
- Suspend: With Suspend, the submit takes place as if these dependencies had not been undefined. However, the job is submitted in a suspended state.
- Ignore: With Ignore, the submit takes place as if these dependencies had not been undefined.

Check Only This flag states whether the job is being submitted or its start capability is being checked.

Parameters list If the job requires any parameters, the parameter values have to be entered in this list prior to the submit. More information about parameters can be found in Chapter [12.5.9](#).

Job Description The description of the job from the Job Definition.

18 Bookmarks

18.1 View




Figure 18.1: Bookmarks

18.2 Concept

18.2.1 In short

Queries from the "Running Master Jobs" and "Search Running Jobs" dialogs can be called in the "Bookmark" dialog. Recurring queries and overview windows can be easily saved and opened again here.

18.2.2 Detailed description

Bookmarks can be made available and modified for the current user or throughout the system. It is also possible to always make certain bookmarks immediately available after logging in to the system. These are then started automatically.

The list of bookmarks contains the following fields:

Bookmark The name of the bookmark is shown here. Selecting the name opens a new window with the request saved as a bookmark.

Mode The *Mode* field shows the type of bookmark. The following types are used:

1. FOLDER

This is a folder bookmark.

2. MASTER

This is a bookmark that was saved by the search mask described in Chapter 19. Using this bookmark means that only Master Jobs are visible.

3. SEARCH

The Search mode designates a bookmark that was saved by the search mask described in Chapter 20. This allows all jobs to be found.

4. DETAIL

This designates a bookmark that describes the "Detail" view (more information about this can be found in Chapter 19). This CHILD is always related to a MASTER mode bookmark (it is the child of the "Master" view).

Which bookmark is selected is determined as follows:

If the Master Job or batch defines a parameter DETAIL_BOOKMARK, then this is used. If this is not the case, the "Detail Bookmark" field for the Master Bookmark is evaluated. If this is not set either, the Detail Bookmark "DEFAULT" is used.

Scope The Scope parameter defines whether the bookmark is visible throughout the entire system wide or just to the current user. The following options are available:

1. SYSTEM

A bookmark with the scope SYSTEM is visible for all users throughout the system. It is displayed in the bookmark list of every user and can be used by all of them. Any changes made to a bookmark from the scope SYSTEM apply system-wide for all users.

Bookmarks of the type SYSTEM, however, can only be created by users with Web GUI administration privileges (see Web Users). Although a 'normal' user is allowed to modify SYSTEM bookmarks, it is then only saved as a USER bookmark. This means that local changes only overwrite the bookmark for the current user, but not for other users. If this bookmark is deleted, the original SYSTEM bookmark is displayed again.

2. USER

If the scope has the value USER, then this bookmark was created by the current user and is only available to this user.

Connection The *Connection* field indicates the validity of the bookmark when using multiple server connections. The following options are available:

1. CURRENT

The bookmark is only valid for the current connection. This bookmark is no longer visible if the bookmark window is opened for another connection.

2. ALL

The bookmark is valid for all server connections.

18.3 Navigation

The Bookmark dialog only features a navigation window. An entry selected here is displayed in a new window.

19 Running Master Jobs

19.1 View

Name	Id	Start	End	Runtime	ExitState	State
SINGLEJOB	13013	12.09.2013 07:58:33	12.09.2013 08:00:01	1m28s	SUCCESS	FINAL
PARAMETERS	13018	12.09.2013 07:58:53	12.09.2013 08:00:26	1m33s	SUCCESS	FINAL
SIMPLEBATCH	13022	12.09.2013 07:58:58		4m51s	ACTIVE	
BATCH_HIERARCHY	13033	12.09.2013 07:59:05		4m44s	FAILURE	IDLE
CONDITIONAL	13061	12.09.2013 07:59:44		4m55s	SKIPPED	ACTIVE
JOBCOMM	13094	12.09.2013 08:00:55	12.09.2013 08:02:45	1m50s	SUCCESS	FINAL
LOADCONTROL	13108	12.09.2013 08:01:13	12.09.2013 08:02:35	1m22s	SUCCESS	ACTIVE
PIPELINE	13112	12.09.2013 08:01:28	12.09.2013 08:03:01	1m33s	SUCCESS	FINAL
TRIGGER	13126	12.09.2013 08:01:36	12.09.2013 08:01:40	2s	FAILURE	FINISHED

Figure 19.1: Running Master Jobs

19.2 Concept

19.2.1 In short

The Running Master Jobs dialog is used to display the currently running or (depending on the settings) the historically completed Master Jobs in the system.

19.2.2 Detailed description

The *Running Master Jobs* dialog is where the currently running jobs (Monitoring) are displayed and the active jobs (Operating) are maintained. The display and filter criteria for the Running Master Job dialog can be modified in the **bookmark** (Default, Master System). This bookmark describes how the entire dialog is displayed.

19.3 Master Navigator

The navigation takes up the whole dialog window. A list of all the currently running Master-submittable jobs and batches is shown here.



Master Navigator

The screenshot shows a web browser window titled "Monitor Batches and Jobs - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area displays a table of jobs with the following columns: Name, Id, Start, End, Runtime, ExitState, and State. The table contains 12 rows, each representing a different job type and its execution details. The jobs listed include SINGLEJOB, PARAMETERS, SIMPLEBATCH, BATCH HIERARCHY, CONDITIONAL, JOBCOMM, LOADCONTROL, PIPELINE, and TRIGGER. The "State" column indicates the current status of each job.

Name	Id	Start	End	Runtime	ExitState	State
SINGLEJOB	13013	12.09.2013 07:58:33	12.09.2013 08:00:01	1m28s	SUCCESS	FINAL
PARAMETERS	13018	12.09.2013 07:58:53	12.09.2013 08:00:26	1m33s	SUCCESS	FINAL
SIMPLEBATCH	13022	12.09.2013 07:58:58		4m51s	ACTIVE	
BATCH HIERARCHY	13033	12.09.2013 07:59:05		4m44s	FAILURE	IDLE
CONDITIONAL	13068	12.09.2013 07:59:44		4m5s	SKIPPED	ACTIVE
JOBCOMM	13096	12.09.2013 08:00:55	12.09.2013 08:02:45	1m50s	SUCCESS	FINAL
LOADCONTROL	13108	12.09.2013 08:01:13	12.09.2013 08:02:35	1m22s	SUCCESS	ACTIVE
PIPELINE	13112	12.09.2013 08:01:28	12.09.2013 08:03:01	1m33s	SUCCESS	FINAL
TRIGGER	13128	12.09.2013 08:01:38	12.09.2013 08:01:40	2s	FAILURE	FINISHED

Figure 19.2: Running Master Jobs Navigator

Clicking the *Find* button opens the page for entering the search criteria as well as configuring the screen.



Auto-Refresh On/Off

The *Auto-Refresh On* button is also a switch and can be used in two positions. In the position "Auto-Refresh On", a refresh is run automatically in the set time.

The time in seconds between two refreshes is entered in the field to the left of the button. The default value is 300 seconds. The colour of the button indicates the state of the Auto-Refresh function. Green = on, red = off.

The list in the dialog shows either all the Master Jobs that are active in the current system (Default) or, if filter options have been entered in the search mask, the correspondingly filtered results. The following columns are shown in the list:

Name The full path and name of the Master Job or just its name are shown here depending on the position of the Show Job button.

Clicking the name opens the Details mask if the object is a Master Submittable Batch or Job with children or otherwise the Details mask for jobs.

Icons can be displayed in front of the name which describe certain states and situations. The following icons are used:



The displayed object is a job that is either currently running or can run. The icon is blue.



The displayed object is a pending job. The icon is lilac.

- Job icon without a period: The job is in the Dependency Wait state.
- Job icon with a period: The job is in the Synchronize Wait state.
- Job icon with two periods: The job is in the Resource Wait state.

Master Navigator



The displayed object is a job that will be submitted in the future by the Time Scheduling.



The displayed object is a job that will be submitted as "suspended" in the future by the Time Scheduling.



The displayed object is a job that has been cancelled. The state of this job can no longer be changed. The icon is brown.



The displayed object is a job that is in the Final state. The state of this job can no longer be changed. The icon is green.



The displayed object is a job that is in one of the states Finished and Restartable, Unreachable or Error. The icon is red.



The displayed object is a batch that is running at this moment. The icon is blue.



The displayed object is a batch that is waiting at this moment. The icon is lilac.



The displayed object is a batch that will be submitted in the future by the Time Scheduling.



The displayed object is a batch that will be submitted as "suspended" in the future by the Time Scheduling.



The displayed object is a batch that is in the Final state. The state of this batch can no longer be changed. The icon is green.



The displayed object is a batch that has been cancelled. The state of this batch can no longer be changed. The icon is brown.



Master Navigator

The displayed object is a batch that is in the Unreachable state, or which has at least one child that is in one of the states Finished and Restartable, Unreachable or Error. The icon is red.



These icons tell the user that the object has been suspended. The icon with the clock indicates that an Automatic Resume has been set for the object. The name colours have the same meanings as the matching coloured buttons.



The *Cancel Run* button is displayed if the job is waiting for resources or dependencies or it has still been allocated on a job server. The scheduled job run can be cancelled by clicking the *Cancel Run* button. The job then gets the state "Cancelled". Neither a rerun nor a resume are possible. The job doesn't get an Exit State either. Jobs that are waiting for a cancelled job have to ignore this job or be cancelled as well.



This button opens a new window displaying the contents of the log files. Should this not be possible, please contact your system administrator.



The *Rerun* button is displayed if the current job has been finished and it is in a "Restartable" state. This means that it has an Exit State which is not final. Clicking the *Rerun* button calls the **Rerun** program (if defined, otherwise the **Run** command). This allows a failed job to be restarted after troubleshooting it.



The *Rerun Children* button is displayed if children of the current job are in a "Restartable" state. This means they have a Restartable Exit State. Clicking the *Rerun Children* button starts the **Rerun** for all the Restartable children. This allows failed children to be restarted after troubleshooting them.

Id The *ID* is the unique identifier of the current "Master Job" runtime instance in the system.

Start The *Start* field shows the time or start of the submit.

End The *End* field shows the finish time for the job. This field stays empty if the job has not yet finished.

Runtime This is the current runtime of the job in days (d), hours (h), minutes (m) and seconds (s). If the job is still running, the runtime is displayed until the last refresh of the dialog and is updated after each new refresh.

Exit State If the job has been finished, its Exit State when it finished is displayed. This field stays empty if the job has not yet finished.

State The Job State is the current runtime state of the job. The runtime system of the schedulix Server system assigns and changes these states when the job (i.e. its runtime instance) has passed through the schedulix system. The job can have the following states:

1. Submitted

The job was submitted manually, by a parent process or by the Time Scheduling and is to be executed. This is the initial state of a job and is usually not visible.

2. Dependency Wait

The job is waiting for required dependencies that have to be fulfilled.

3. Synchronize Wait

The job is waiting for the required Synchronizing Resources.

4. Resource Wait

The job is waiting until sufficient System Resources become available.

5. Unreachable

The job cannot be executed because one or more dependencies have not been fulfilled. This situation can be resolved by ignoring dependencies.

6. Cancelled

The job has been manually cancelled and is no longer running.

7. Error

The job cannot be executed due to a definition error. An example of this is a required resource that cannot be provided by any of the job servers. In this case, the job can be restarted after the error has been remedied.

What also frequently happens is that a mistake is made when entering the Run program. If the job server is consequently unable to start the process, the job is placed in an Error state. The job concerned can be restarted in this case as well.

If a Master Submit by the Time Scheduling fails, the master is assigned an Error State in the system to identify this error. This makes the error visible for the person responsible. The job or batch is not restartable, however, and has to be submitted manually after being repaired.

8. Runnable

The job can be started by a job server. The job server should be checked if the job remains in this state for any length of time.

9. Starting

The job was handed over to a job server to be started.

10. Started

The job was handed over to a job server to be started. The job server has acknowledged the job.

11. Running

The Run command was started by the job server.

12. To Kill

The user has instructed the schedulix Server to run the Kill program associated with the job. The Kill program can be run by clicking the *Cancel Run* button.

13. Killed

The Kill program associated with the job has been run.

14. Broken Active

The job server has lost the connection to the Run program. Although the process is still running, the result of the job (Exit Code) can no longer be returned to the job server and correspondingly the schedulix Server.

15. Broken Finished

If a job has the state "Broken Active" and it stops, it switches to the state "Broken Finished". This means that the process is no longer running. The Exit Code for the process could no longer be determined by the job server. It is necessary to manually check the job results and set the Exit State. The Exit State is set with the *Set State* button.

16. Finished

The job has been finished. The Exit State is shown in the *Exit State* field.

17. Final

The job and all its children have finished and have a Final Exit State.

The [19.3](#) diagram shows the status transitions of the schedulix runtime system.

Variables Depending on the configuration, the state column can be followed by the variables defined there in the specified order and colour.

Master Navigator




Figure 19.3: State chart for batches and jobs

19.3.1 Master Navigator Query mask

The Query mask is opened by clicking the *Settings* button. The query requirements for the current dialog are set in this mask. The filter and display criteria can be

Master Navigator

edited here.




Figure 19.4: Running Master Jobs Query mask

The dialog looks like in Figure 19.4.

Making changes in the mask and then saving them with the current or a changed bookmark name (entry field in the upper right of the headline) allows them to be used again at a later time. If the changes are only temporary ones, you can simply enter them and jump back to the navigation window. This means that the changes only apply for the current search. The changes will be discarded if the dialog is then closed.

The fields shown above have the following meanings:

History/Unit The maximum time that may have elapsed since the end of a workflow so that it appears in the display. For example, if all the currently running jobs and all the jobs that have been completed in the last hour are to be displayed, the value 1 has to be entered here and the unit "Hours" has to be set in the *Unit* field. If all the running jobs and all the jobs that have been completed in the last hour are to be displayed, the value 1 has to be set in the *History* field with "Hours" in the "Unit" field. The following options are available:

1. Minutes
2. Hours
3. Days

Future/Unit The setting in the *Future* field determines whether and to what extent batches or jobs that are to be submitted in the future are displayed in the Master List. Regardless of the interval selected in the *Future* field, only the next start can be displayed for schedules without an activated calendar. Only planned starts can be displayed in the calendar horizon of schedules with a calendar. The *Unit* field is handled like the "History" field.

Name Patterns list The list of Name Patterns can hold an amount of Name Patterns to be used to filter the returned results. A Name Pattern can be the full name, part of the name or a regular expression.

When searching for multiple entries in the list, all occurrences of these entries are sought. These are linked with "OR".

Example:

You want to search for all occurrences of the name "Pipeline".

The following Name Patterns are possible:

- Full: **pipeline** (the search is not case-sensitive)
- Sub-string: **eline**
- Regular expression: **P.*line**

The list of Name Patterns can be modified using the *Add* and *Remove* buttons.

Exit State list Exit States can be entered in this list which the sought jobs must have if they are to be shown in the results list.

Example:

Only those jobs with the Exit State "Failure" are to be listed. By adding the Exit State "Failure" to the list of Exit States you can now search for these jobs. If the list does not contain any entries, the Exit States are ignored for the search. The list of Exit States can be modified using the *Add* and *Remove* buttons.

Job State list Job States in one of the defined states can be searched for in this list. If the list does not contain any entries, the Job States are ignored for the search. The list of Job States can be modified using the *Add* and *Remove* buttons.

Condition mode The Condition mode defines whether and how the subsequent condition is to be evaluated.

The following options are available:

- None: The condition is ignored.
- And: The search query and condition are linked with "And".
- Or: The search query and condition are linked with "Or".

- Minus: The search query and negation of the condition are linked with "And".
- Only: Only the condition is taken into account.

Condition The query can be refined using the condition. Please refer to the syntax documentation (List Job) for the exact syntax and the various options that are available to do this.

Variables list All the job variables and parameters that are to be displayed in the navigation list are defined in the Variables list. This is an easy way to get an overview of the main job parameters. They are displayed at the end of each line in the navigation list.

The following fields have to be populated for the variables:

Name The name of the variable or parameter has to be entered here.

Align The *Align* field is for defining how the variables are to be displayed in the list. The following options can be selected:

- LEFT
The display is left-aligned.
- CENTER
The display is centred.
- RIGHT
The display is right-aligned.

Format The format for the output can be defined here. The following options are available:

- NONE
No formatting is used.
- NUMBER
Numerical formatting with thousand separators and decimal points (as necessary) is used.

Label A heading that is to be shown for the variable in the list header can be defined here. The variable name is used as a heading if a label is not defined.

Colour The background colour for the variables column can be configured here. A preview of each colour is shown in the field's drop-down menu.

Seq. The list order of the variables can be changed using the *Up* and *Down* buttons. These buttons are only visible if the variables list has more than one entry in it.

Bookmark Settings The entry fields in the *Bookmark Settings* section are only meaningful when a bookmark is saved. The following fields describe the properties of the bookmark to be saved.

Detail Bookmark If the name of a detail bookmark is entered in the field *Detail Bookmark*, opening a detail navigation window from a master navigation window opened by this bookmark will use the entered bookmark name for the opened detail navigation window. If this detail bookmark is not yet existing, the detail navigation window will show this bookmark name in square brackets. Those will disappear when the detail bookmark is saved in the opened detail navigation window. The field *Detail Bookmark* is only visible in the *Bookmark Settings* of master navigation windows.

Connection This field defines the server connections the bookmark can be used with. There are two variants:

1. ALL

The bookmark can be used with any server connection.

2. CURRENT

The bookmark can be used with the current server connection only.

Scope This field defines the visibility of the bookmark to users. There are two variants:

1. USER

The bookmark is only visible to the current user. No other users can see this bookmark or, if it is a change made to a system-wide bookmark, only the change is locally visible to this user.

2. SYSTEM

The bookmark or current change is visible to all users. This means that all users see this new bookmark in their bookmark list. If it is a change made to an existing system-wide bookmark, the change affects all the users.

Detail Navigation

Autostart Autostart defines whether a bookmark is to be activated immediately after a user logs on. There are two variants:

1. YES

The bookmark is started immediately after the login. This means that a window opens on the screen after the user logs on for each bookmark where the autostart option is set to "YES".

2. NO

The bookmark must be explicitly started by clicking in the Bookmark dialog and it is not automatically activated. This is the default setting.

Start Mode The Start mode defines whether the query is run immediately after the bookmark has been started and the navigation window is to be opened, or whether the query dialog is to be displayed. There are two variants:

1. NAVIGATE

The query is run immediately and the navigation screen opens with all the results.

2. QUERY

The Query dialog is displayed first.

19.4 Detail Navigation

Clicking one of the Master Jobs (which has children) opens the "Details" window for the

Name	Id	Start	End	Runtime	ExitState	State
BATCH_HIERARCHY	13033	12.09.2013 07:59:05		14m49s	FAILURE	IDLE
LOAD	13042	12.09.2013 07:59:05		14m49s	FAILURE	IDLE
LOAD_3	13048	12.09.2013 07:59:19	12.09.2013 07:59:26	7s	SUCCESS	FINAL
LOAD_1	13046	12.09.2013 07:59:14	12.09.2013 07:59:27	13s	SUCCESS	FINAL
LOAD_2	13044	12.09.2013 07:59:09	12.09.2013 07:59:29	20s	FAILURE	FINISHED
REPORT	13034				DEPENDENCY_WAIT	
REPORT_3	13036				DEPENDENCY_WAIT	
REPORT_2	13038				DEPENDENCY_WAIT	
REPORT_1	13040				DEPENDENCY_WAIT	

Figure 19.5: Running Master Jobs "Details" window

Master Job. It looks like in Figure 19.5.

The "Details" window is similar to the Master window but with the following differences:

Detail Navigation

19.4.1 Tree view

Instead of the flat list in the Master window, all the children of the current batch and their own children are displayed in a tree hierarchy. The first level is expanded, but the others are not. The exception here is the search using Name Patterns, since here every found job is listed and the respective level and parent levels are also automatically expanded.

19.4.2 Search results

The results of a search are indicated by an arrow symbol



This enables you to differentiate between found entries that were expanded because another child entry meets the search criteria, and search hits. Only results are displayed with the arrow icon.

19.4.3 Detail navigation query mask




Figure 19.6: Navigation Details query mask

The query mask looks like in Figure 19.6

The query mask for the Navigation Details corresponds to the query mask for the Master Navigation, but with the following deviations:

The fields *History*, *Unit*, *Bookmark Details* and *Autostart* are not shown here and some other fields have been added.

Details mask for jobs

The mask contains the following fields:

Job State list In the Job State list, jobs which have the entered Job State can be filtered. Each of the states described in the *State* field can be selected and used as a filter.

If no Job States have been entered, they are ignored by the search.

Job ID A specific Job ID that is to be searched for can be entered in this field. A maximum of one job can be returned as a result.

Sort By The list can be sorted by various criteria with this function. The following criteria can be used:

1. Name

The results are sorted alphabetically by job name.

2. Start Date

The results list is sorted chronologically by the job start time.

3. End Date

The results list is sorted chronologically by the job finish time.

19.5 Details mask for jobs

Clicking a job without children in the "Details" or "Master" window automatically opens the Details mask for jobs. It looks like this:

All the runtime details are displayed in this mask and some of them can also be modified here. It also features some action buttons which allow the current objects to be manipulated.

19.5.1 Buttons

The following buttons are all action buttons. This means that these buttons are used to perform a certain action on the selected job. To do this, a Confirm window opens like in Figure 19.8 after you click the button.

You can now enter a comment in the text box to describe the reason for the action or some other remarks. This text is later shown in the Audit tab in the field *Reason/Comment*. The action is performed when you click the respective action button again. Clicking the *Cancel* button aborts the action and opens the previous screen again. The Confirm screen contains additional input fields for controlling the action (these vary according to the action).

Details mask for jobs

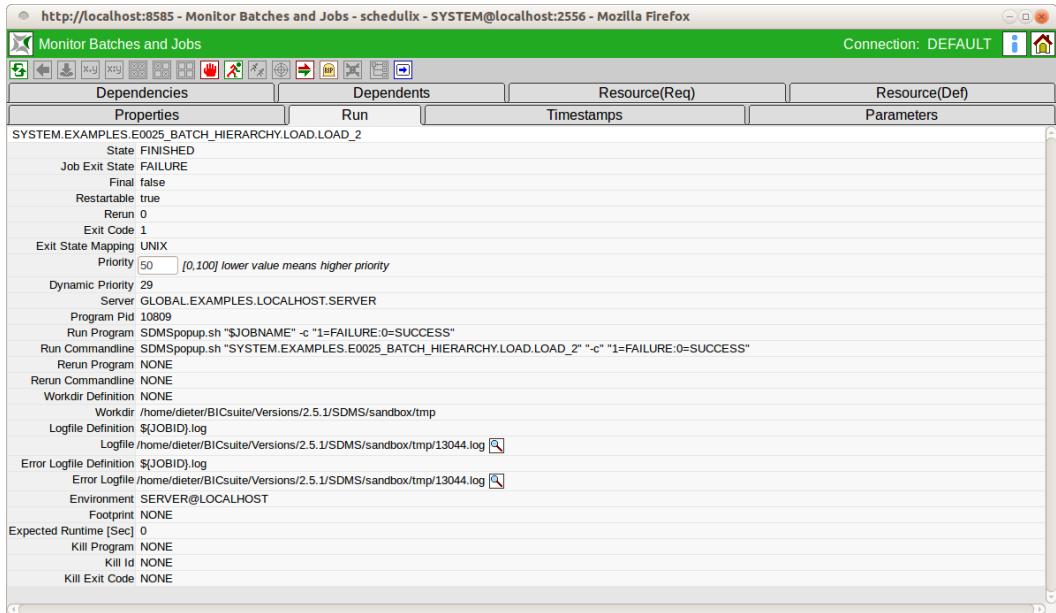


Figure 19.7: Details mask for jobs

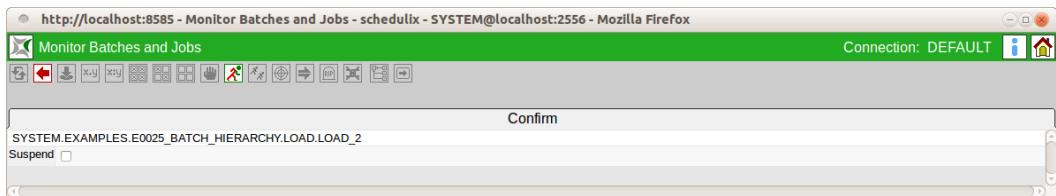


Figure 19.8: Confirm mask



Jobs can be suspended using the *Suspend* button. This is only possible if the jobs do not already have the "Suspended" state and they have not yet been completed. The jobs are placed in the "Suspended" state if this action is successful.

Suspend Local If the Suspend Local checkbox is checked, the children of the suspended job or batch will not be suspended. This way the execution of the children will continue, but the the batch or job itself will not go FINAL. Dependent batches or jobs will have to wait until the batch or job is resumed.

Details mask for jobs

Suspend Restricted The Suspend Restricted checkbox is only editable for users of group ADMIN. If checked only ADMIN users are allowed to resume the batch or job.

Additional input fields: *Resume, Resume Time, Resume In, Unit*

Resume Here you can choose whether a workflow should be resumed automatically.

The following options are available:

- NO: deselects this functionality and no more input fields are displayed.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Resume Time The required Resume Time in the format "YYYYMMDD HH:MM" is entered here.

Resume In Here you can specify how many time units (see *Unit*) the system is to wait for until the resume action is triggered.

Unit This field is used to define whether the entry in *Add Resume* is in minutes (MIN), hours (HOUR) or days (DAY).



Resume

Jobs in the "Suspended" state can be reactivated with the *Resume* buttons. This means that the resources and dependencies are checked again, and if all these checks are successful the job is started on a job server.

Additional input fields: *Resume, Resume Time, Resume In, Unit*

Resume The resume type can be selected here.

- OFF: This option is only available if the job is suspended and an "Auto-resume" (IN/AT) is active. It disables this auto-resume function.
- NOW: A resume takes place immediately. No other input fields are displayed.
- AT: selects an automatic resume action at a fixed time. The *Resume Time* input field is displayed.
- IN: selects an automatic resume action after a time period has expired. The *Resume In* and *Unit* input fields are displayed.

Details mask for jobs

Resume Time, Resume In, Unit See above (Suspend action)



Rerun

The *Rerun* button is displayed if the current job has been finished and it is in a "Restartable" state. This means that it has an Exit State which is not final. Clicking the *Rerun* button calls the **Rerun** program (if it has been defined, otherwise the **Run** command). This allows a failed job to be restarted after troubleshooting it.

Additional input fields: *Suspended, Resume, Resume Time, Resume In, Unit*

Suspended If this flag is set, the job is suspended and the *Resume* check box is displayed. If this field is not set, no other fields are displayed either.

Resume, Resume Time, Resume In, Unit See above (Suspend action)



Rerun Children

The *Rerun Children* button is displayed if children of the current job are in a "Restartable" state. This means they have a Restartable Exit State. Clicking the *Rerun Children* button starts the **Rerun** program for all the Restartable children. This allows failed children to be restarted after troubleshooting them.

Additional input fields: See Rerun



Kill

The *Kill* button is only displayed if the current job is running at that moment (state "Running") and a **Kill** program has been defined in the job definition. Clicking the *Kill* button calls the respective Kill program and the job's state is set to "TO KILL". If the Kill program was run successfully, the job's state changes to "KILLED". This enables the job to be terminated during the program runtime.



Set State

This allows you to manually set an Exit State for a job. The *Set State* button is displayed if a job fulfills at least one of the following conditions:

- The job is in an Error State (restartable)
- The job is in a Pending State
- The job is suspended and not active and not in a Final State

Additional input fields: *Set Exit State, Force, Resume*

Set Exit State A valid Exit State can be selected here in the *Set Exit State* field.

Details mask for jobs

Force Here you can select whether you want to allow an Exit State that cannot be attained by the job itself using its Exit Code.

Resume Here you can select whether a suspended job is to be resumed after the Set State.



Cancel Run

The *Cancel Run* button is displayed if the job is waiting for resources or dependencies or it has still been allocated on a job server. The scheduled job run can be cancelled by clicking the *Cancel Run* button. The job then gets the state "Cancelled". Neither a rerun nor a resume are possible. The job doesn't get an Exit State either. Jobs that are waiting for a cancelled job have to ignore this job or be cancelled as well.



Comment

You can enter a comment regarding the job by clicking the *Comment* button. Apart from entering the comment, no other action is performed on the job.



Edit Job

The *Edit Job* button is not an action button because no changes can be made here to the current runtime object. Clicking the *Edit Job* button opens a new window with the **Batches and Jobs Definition Screen** for the job definition. The data shown here refers to the definition of the task at the time of the submit.

19.5.2 Properties tab

The Properties tab contains all the information regarding the job definition and the current state. It looks like this:

The fields in the "Properties" tab have the following meanings:

ID This is the unique identification number of the job runtime object.

Type This is the type of runtime object. More details about types can be found in Chapter [12.5.1](#).

Submit Path This is the full submit path to the runtime object. This means that the comma-separated path used for the submit is given here.

Tag This is the name of the Child Tag. This is only required for a dynamic child and uniquely identifies the current child instance. More details about dynamic children can be found in Chapter [12.5.4.1](#).

Details mask for jobs




Figure 19.9: Job and batch properties

Owner This is the name of the group that created the job definition for this runtime object.

Submitted By This is the name of the group that submitted the runtime instance. If the job was dynamically submitted by a parent program, this is the user who submitted the parent.

Unresolved Handling The *Unresolved Handling* field is only shown in the Master Job mask. It shows the setting of the *On Unresolved Error* field from the submit mask.

- Ignore: Unresolved Dependencies are always ignored.
- Suspend: If errors occur due to Unresolved Dependencies, the entire process is submitted as being suspended.
- None: Default behaviour; in the event of Unresolved Dependencies, the behaviour is as defined in the Dependency Definition.

Nice Value This is an input field where the Nice Value can be entered. More information about Nice Values can be found in Chapter 12.5.4.1. Since this is the only input field in the dialog box, you have to press the tabulator key after entering the data to enable the *Save* button for saving your changes.

Suspended This field indicates whether the runtime object is currently suspended (TRUE) or not (FALSE). The Suspend State can be changed using the *Suspend* and *Resume* buttons.

Suspend by Parent This indicates whether the runtime object was suspended by the parent (1) or not (0).

State This shows the current state of the runtime object. More details about states can be found in Chapter [19.3](#).

Error Message If the state "ERROR" was reported in the *State* field, a detailed description of the error is displayed.

Exit State If the job has finished, its resultant Exit State is shown here.

Exit State Profile The Exit State Profile entered in the job definition is shown here. More details about Exit State Profiles can be found in Chapter [12.5.1](#).

Master Id This is the ID of the Master Job that was submitted in order to create this runtime object. If the object itself was submitted as the Master Job, this is identical to the ID.

Submitting Parent ID This is the ID of the parent runtime object that submitted the current job. If the job does not have a parent, "NONE" is displayed here.

Static *Static* is the descriptor indicating whether the current runtime object is a statically (TRUE) or dynamically submitted child (FALSE). More details about dynamic children can be found in Chapter [12.5.4.1](#).

Merge Mode This is the Merge Mode that is currently being used. More details about the Merge Mode can be found in Chapter [12.5.4.1](#).

Version This is the version number of the Job Definition current at the time of the submit.

19.5.3 Run tab

The Run tab contains all the information relating to the current details about the Run program, the runtime environment and the computer. The tab looks like this: The fields in the "Run" tab have the following meanings:

Details mask for jobs

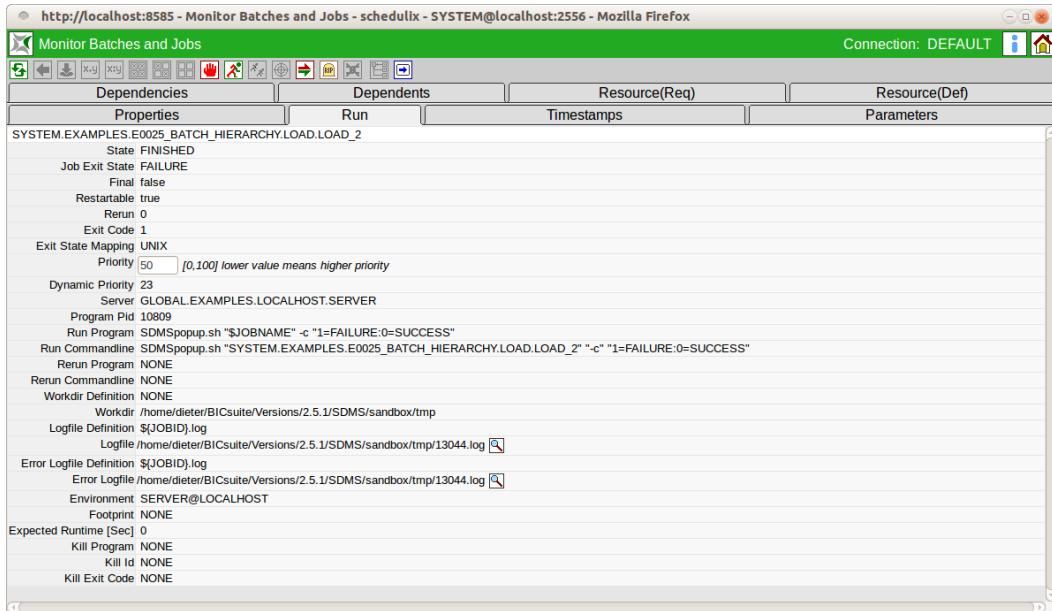


Figure 19.10: Batch und Job Run information

State This shows the current state of the runtime object. More details about states can be found in Chapter [19.3](#).

Job Exit State The Job Exit State shows either the job's current Exit State (if it has finished) or "NONE".

Final If the Job Exit State is a Final State, the value in this field is "TRUE"; if it is "NONE" or not a Final State, the value is "FALSE".

Restartable If the job is restartable, the value in this field is "TRUE", otherwise it is "FALSE". If the value is TRUE, the *Rerun* button is enabled.

Rerun The *Rerun* field shows the number of times the job has been restarted.

Exit Code The Exit Code is the exit value that the Run program had when the process finished. More details about the Exit Code can be found in Chapter [12.5.2](#).

Exit State Mapping This is the Exit State Mapping defined for the job in the Batches and Jobs dialog that was used to convert the Exit Code into the Job Exit State. More details about Exit State Mapping can be found in Chapter [12.5.2](#).

Details mask for jobs

Priority This is an input field. It shows the current priority, which can also be changed here. The new priority is saved by pressing the tabulator key and clicking the *Save* button. More details about priorities can be found in [12.5.2](#).

Dynamic Priority This is the dynamic (effective) priority for the job. It rises as the number of submits increases until the job is executed. More details about the dynamic priority can be found in Chapter [11.4.2.3](#).

Server The *Server* field shows the current job server on which the process is running or has been run.

Program PID The *Program PID* is the process ID of the [Run program](#).

Run Program This is the command line for the [Run program](#) defined in the Batches and Jobs dialog. The variables and parameters are still displayed here under their names and are not substituted.

Run Command Line This is the command line as defined in the Batches and Jobs dialog for the [Run program](#)

with substituted values. This is the view as it was handed over to the processing shell. All the variables and parameters were substituted for their current values.

Rerun program This is the command line as defined in the Batches and Jobs dialog for the [Rerun program](#). The variables and parameters are still displayed here under their names and are not substituted. NONE is shown here if a Rerun program has not been defined.

Rerun Command Line This is the command line as defined in the Batches and Jobs dialog for the [Rerun program](#) with substituted values. This is the view as it was handed over to the processing shell. All the variables and parameters were substituted for their current values. NONE is shown here if a rerun has not taken place.

Workdir Definition The *Workdir Definition* field shows the directory from where the process is being run. If "None" is shown here, the default working directory for the executing job server is used.

Workdir This is the [Workdir](#) that is being used (or is to be used) in the current job. The substituted parameters and variables are displayed here if the job has finished or it is still running (otherwise just the parameters).

Details mask for jobs

Logfile Definition The Logfile Definition corresponds to the definition of the [log files](#) in the Batches and Jobs dialog.

LogFile The name and path of the currently used [log files](#) are shown here (if the job has finished or it is still running), otherwise NONE is displayed. Possible variables and parameters are substituted for the actual values. Clicking the button



opens a new window displaying the contents of the log files. Should this not be possible, please contact your system administrator.

Error Logfile Definition The Error Logfile Definition corresponds to the definition of the [error log files](#) in the Batches and Jobs dialog.

Error Logfile The name and path of the currently used [error log files](#) are shown here (if the job has finished or it is still running), otherwise NONE is displayed. Possible variables and parameters are substituted for the actual values. The button



has the same function as described above for log files.

Environment This field corresponds to the definition of the [environment](#) in the Batches and Jobs dialog.

Footprint This field corresponds to the definition of the [footprint](#) in the Batches and Jobs dialog.

Expected Runtime This corresponds to the definition in the [Batches and Jobs dialog](#).

Kill Program This corresponds to the definition of the [Kill program](#) in the Batches and Jobs dialog.

Kill Id The Kill ID contains the process ID of the [Kill program](#) if this has been started. If a Kill action has not yet been performed, the Kill ID is NONE.

Kill Exit Code The Kill Exit Code is the result returned to the job server by the [Kill program](#) after the kill process has finished. This can be used to determine whether a kill attempt has succeeded or failed if this is configured in the Kill program.

19.5.4 Timestamps tab

The Timestamps tab shows the different dates and times of the respective jobs. These times give an overview of the temporal behaviour and duration of the job runs. The tab looks like this:

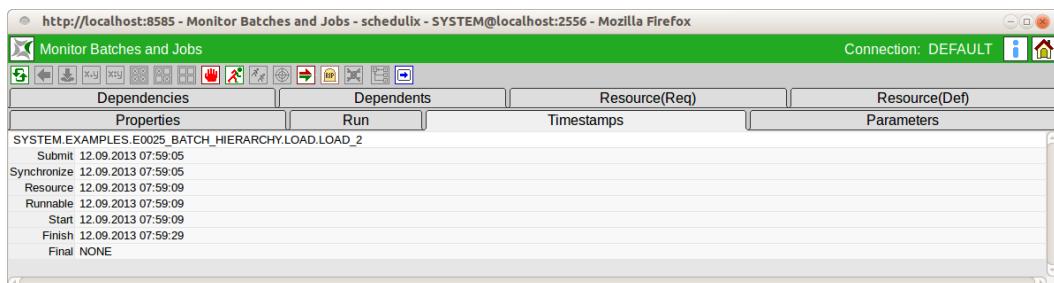


Figure 19.11: Batch and job timestamps

The tab is for information purposes only; no entries can be made here. The fields in the "Timestamps" tab have the following meanings:

Submit The submit time for the current job is shown in the Submit timestamp. If it is a static job this is the submit time for the Master Job; in the case of dynamic jobs it is the time when the parent job performed the dynamic submit. The submit date is the same as the time when the job switched to the "Dependency Wait" state.

Synchronize The Synchronize timestamp indicates the time when the job switched to the "Synchronize Wait" state. That is the time when all the job's possible dependencies were fulfilled.

Resource The Resource timestamp shows the time of the status transition to Resource Wait. This means that all the required **Synchronizing Resources** had been fulfilled and allocated to the job at this time.

Runnable "Runnable" is the time when the job entered the "Runnable" state. This means that all the dependencies and resource requirements could be fulfilled. From now on the job can be processed by a server job.

Start The Start timestamp shows when the Run program was started by the job server. If the job has already been executed several times (Rerun), the last start is entered here.

Finish The Finish timestamp shows when the job was completed. If the job has already been executed several times (Rerun), the last finish time is entered here.

Details mask for jobs

Final The Final timestamp shows when the job that has created a Final State finished. If the job was finished but it has a Restartable State, NONE is entered here.

19.5.5 Dependencies tab

All the current job's dependencies on other jobs are displayed in the Dependencies tab. The tab looks like this:




Figure 19.12: Batch und job dependencies

This dialog gives an overview of all the required dependencies, their current state and the possibility for deliberately ignoring dependencies.

The mask contains the following buttons:



Ignore Dependencies

The *Ignore Dependencies* button is used to deliberately ignore one or more dependencies. All the dependencies that are to be ignored must have been previously selected in the mask. When you now click the button, a new window opens where you can select the Ignore mode and enter a remark. The window looks like this:




Figure 19.13: Batches and Jobs Ignore Dependencies confirmation mask

The ignore mode has to be entered in this mask. By default, all the children of a Submitted Entity inherit its dependencies: if B is required by A, each of A's children (CA) is likewise dependent on B. The Ignore mode can now be used to determine the behaviour of the inherited dependency at the children.

The following options are available:

1. Recursive

This is the default behaviour.

The Ignore refers not only to the current job, but also to the inherited dependencies on the required job held by the children of the current job. This means that the parent job ignores this inherited dependency and all its children ignore it as well. If a child has an explicitly stated dependency on the required job which was not inherited, this is not ignored but is observed. Figure 19.14

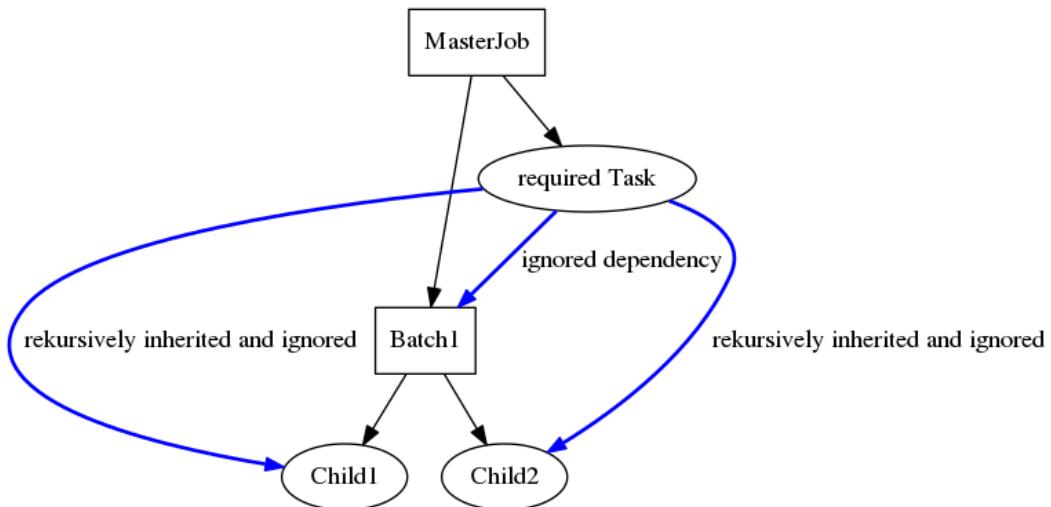


Figure 19.14: Example of a Recursive Ignore

shows an example.

2. Job Only

Only the job itself ignores the dependency. All the children observe the inherited dependency and wait for the required job. This allows the parent job to be executed while forcing its children to wait. Figure 19.15 shows an example.

Having entered the Ignore mode and then a comment in the Comment field, you can perform the action by clicking the *Ignore Dependency* button again. The tab list shows all the dependencies that have to be fulfilled when the job starts and their current states. In this list, you can make a selection that can be used to ignore dependencies by clicking the *Ignore Dependencies* button.

The following fields are shown in the list:

Icon field The current state of the dependency is shown in the *Icon* field. The following options are available:



Details mask for jobs




Figure 19.15: Example of a Job Only Ignore

The dependence has not yet been fulfilled or can no longer be fulfilled. Should this icon be visible next to a dependency with the **Dependency Mode ALL**, the current job cannot be executed by this dependency.

In the **Dependency Mode ANY**, at least one dependency must not be red for the job to be executed.



The dependency has been fulfilled. In the **Dependency Mode ANY**, at least one dependency must have the state Green for the job to be able to run. In the **Dependency Mode ALL**, all the entries must be green for the job to be run.

Required Name This is the name of the job that has to be a prerequisite for starting the current job run. You can switch to its runtime definition and check its state by clicking the name.

ID The runtime ID of the Required Job is shown in this field.

Final State If the Required Job has been completed, the job's Final State is shown here. If a Final State is entered here and the dependency is not fulfilled, the reason for this happening must be an incorrect Final State for this dependency. This means that there is a difference between a Final State and a Required State. If the job is to be executed now after all, this can only be done with an Ignore by clicking the

Details mask for jobs

Ignore Dependency button. If the dependencies in the list are in the **Dependency Mode** ANY, at least one another dependency has to be fulfilled.

Required The required Final State for the job is shown in this field. For the dependency to be fulfilled, the Final State must be the same as the Required State or the Required State must be NONE.

Dependent Name The job demanding this dependency is named in this field. This is normally the currently selected job, but it can also be one of the job's parents.

Check On The *Check On* field describes the relation of the completion of the Required Job. The following variants are used:

1. All Final

The Required Job and all its children must be final

2. Job Only

Only the Required Job has to be final; the children are ignored.

More details about the Check On function can be found in Chapter [12.5.6.1](#).

ID This is the ID of the dependent job.

Mode The *Mode* defines how the list of jobs is to be handled. The following variants are used:

1. All

All the list's dependencies must be fulfilled before the job can start.

2. Any

All least one of the list's dependencies must be fulfilled before the job can start.

More details about Mode can be found in Chapter [12.5.6](#).

State This is the current state of the dependency relationship. The following variants are used:

1. OPEN

The Required Job has still not attained a Final State (it has not yet run, is still running, or an error has occurred). The dependency has not yet been fulfilled, although it can be fulfilled by attaining the correct Final State.

Details mask for jobs

2. FULFILLED

The Required Job has finished and attained the required Final State. The dependency has been fulfilled.

3. FAILED

The Required Job has finished and attained an incorrect Final State. The dependency can no longer be fulfilled. In the **Dependency Mode ALL**, the current job can only be made to start by ignoring the dependency. In the **Dependency Mode ANY**, at least one of the other dependencies is being fulfilled.

19.5.6 Dependents tab

All the other jobs' dependencies on the current job are displayed in the Dependents tab. The tab looks like this:




Figure 19.16: Batch and Job Dependents tab

This dialog gives an overview of all the existing dependencies, their current state and the possibility for deliberately ignoring dependencies.

The mask contains the following buttons:



The *Ignore Dependencies* button is used to deliberately ignore one or more dependencies. All the dependencies that are to be ignored must have been previously selected in the mask. When you now click the button, a new window opens




Figure 19.17: Batch and Job Ignore Dependencies

where you can select the Ignore mode and enter a remark. The window looks like in Figure 19.17

The ignore mode has to be entered in this mask. By default, all the children of a Submitted Entity inherit its dependencies: if B is required by A, each of A's children (CA) is likewise dependent on B. The Ignore mode can now be used to determine the behaviour of the inherited dependency at the children.

The following options are available:

1. Recursive

This is the default behaviour.

The Ignore refers not only to the dependent job, but also to the inherited dependencies on the current job held by the children of the dependent job. This means that the parent job ignores this inherited dependency and all its children ignore it as well. If a child has an explicitly stated dependency on the required job which was not inherited, this is not ignored but is observed.

2. Job Only

Only the dependent job itself ignores the dependency. All the children observe the inherited dependency and wait for the current job. This allows the parent job to be executed while forcing its children to wait.

Having entered the Ignore mode and then a comment in the Comment field, you can perform the action by clicking the *Ignore Dependency* button again. The tab list shows all the dependencies that have to be fulfilled before other jobs are able to start together with their current states. In this list, you can make a selection that can be used to ignore dependencies by clicking the *Ignore Dependencies* button.

The following fields are shown in the list:

Icon field The current state of the dependency is shown in the *Icon* field. The following options are available:



The dependence has not yet been fulfilled or can no longer be fulfilled. Should this icon be visible next to a dependency with the **Dependency Mode ALL**, the dependent job cannot be executed by this dependency.

In the **Dependency Mode ANY**, at least one other dependency (which is not visible here) of the dependent job must not be red for the dependent job to be able to run.



The dependency has been fulfilled. In the **Dependency Mode ANY**, at least one dependency must have the state Green for the dependent job to be able to run. In the **Dependency Mode ALL**, it may be necessary for other dependencies (which are not visible here) of the dependent job to be fulfilled for the dependent job to be able to run.

Details mask for jobs

Dependent Name This is the name of the job that is dependent upon the fact that the current job has been executed. You can switch to its runtime definition and check its state by clicking the name.

ID The runtime ID of the dependent job is shown in this field.

Final State If the dependent job has already been completed, the job's Final State is shown here.

Required The required Final State for the current job is shown in this field.

Check On The *Check On* field describes the relation of the completion of the current job. The following variants are used:

1. All Final

The current job and all its children must be final.

2. Job Only

Only the current job itself has to be final; the children are ignored.

More details about the Check On function can be found in Chapter [12.5.6.1](#).

Mode The *Mode* defines how the list of jobs is to be handled. The following variants are used:

1. All

All the current job's dependencies must be fulfilled before the job can start.

2. Any

At least one of the current job's dependencies must be fulfilled before the job can start.

More details about Mode can be found in Chapter [12.5.6](#).

State This is the current state of the dependency relationship. The following variants are used:

1. OPEN

The current job has still not attained a Final State (it has not yet run, is still running, or an error has occurred). The dependency has not yet been fulfilled, although it can be fulfilled by attaining the correct Final State.

Details mask for jobs

2. FULFILLED

The current job has finished and attained the required Final State. The dependency has been fulfilled.

3. FAILED

The current job has finished and attained an incorrect Final State. The dependency can no longer be fulfilled. In the **Dependency Mode ALL**, the dependent job can only be made to start by ignoring the dependency. In the **Dependency Mode ANY**, at least one other dependency (which is not visible here) of the dependent job must be fulfilled.

19.5.7 Resource(Req) tab

All the information about the current state of the requested resource is displayed in the Resource(Req) tab.

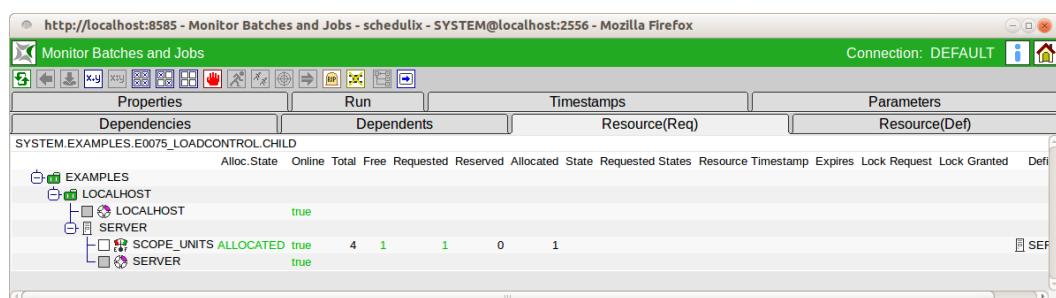


Figure 19.18: Job Resource Requirements

The tab looks like in Figure 19.18

All the resource instances and their allocations to scopes can be viewed in this tab. It is also possible to ignore requests for resources.

The *Ignore Resources* button has the following meaning:



The *Ignore Resources* button is used to deliberately ignore one or more resource requests. All the resource requests that are to be ignored must have been previously selected in the mask. When you now click the button, a new mask opens where you can enter a remark. The mask looks like in Figure 19.19

After you have entered a comment in the Comment field, clicking the *Ignore Resource* button again performs the action.

All the requested resources of the job are displayed in the Resources list. These are displayed hierarchically using the scopes and job servers where the requested resources are located. The following columns are shown in the list:

Details mask for jobs




Figure 19.19: Job Ignore Resource Requirement

All the requested resources together with their names are displayed in the first column of the list in a hierarchy (if the resources are located in scopes and job servers).

The resources for an Ignore can be selected by clicking the preceding *Ignore Dependencies* button. If the preceding field is grey this is not possible (static resources). Clicking the resource name opens the [11.4.2.3](#) dialog, where the current allocation of the resource can be examined.

Allocation State The Allocation State is the current state of the resource request. The following variants are used:

1. None

These are static or system resources that are checked at the time of the submit. If these are no longer present or are offline, an error message is returned at the submit time.

2. Blocked

The resource is being used by another job and cannot be allocated. The job is now waiting for this resource. There may be different reasons for a block. The criteria preventing an allocation are shown in red in the list.

3. Allocated

The resource has been allocated to the current job, fulfilling the requirement.

4. Available

The resource is available but has not yet been allocated.

19.5.8 Parameters tab

All the input and output parameters currently defined in this job and their actual values are shown in the Parameters tab. The tab looks like this:

This tab is for information purposes only; no changes can be made here.

The "Parameters" list shows all the parameters that are currently defined in the job. Information about parameters can be found in Chapter [12.5.9](#).

Details mask for jobs

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar reads "Monitor Batches and Jobs - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area displays a table titled "SYSTEM.EXAMPLES.E0015_PARAMETERS.PARAMETERS". The table has columns: Name, Type, and Value. The data rows are:

Name	Type	Value
E0015_SCOPE_PARAMETER	PARAMETER	SCOPE_PARAMETER_VALUE
FOLDER_NAME	PARAMETER	E0015_PARAMETERS
SUBMIT_PARAMETER	PARAMETER	SUBMIT_PARAMETER_VALUE

Figure 19.20: Batch and job parameters

20 Search Running Jobs

20.1 View

The screenshot shows a web-based monitoring tool for batches and jobs. On the left, there is a hierarchical tree view of job definitions. Under 'SIMPLEBATCH', there are 'END' and 'START' nodes. Under 'BATCH_HIERARCHY', there is a 'LOAD' node which further branches into 'LOAD_1', 'LOAD_2', and 'LOAD_3'. Below 'LOAD', there is a 'REPORT' node with three sub-nodes: 'REPORT_1', 'REPORT_2', and 'REPORT_3'. On the right, a table lists the details of running jobs. The columns are: Name, Id, Start, End, Runtime, ExitState, and State. The data in the table is as follows:

Name	Id	Start	End	Runtime	ExitState	State
SIMPLEBATCH	13022	12.09.2013 07:58:58		59m52s	ACTIVE	
END	13025				DEPENDENCY_WAIT	
START	13023	12.09.2013 07:58:59		59m51s	RUNNING	
BATCH_HIERARCHY	13024	12.09.2013 07:59:05		59m45s	FAILURE	IDLE
LOAD	13042	12.09.2013 07:59:05		59m45s	FAILURE	IDLE
LOAD_1	13046	12.09.2013 07:59:14	12.09.2013 07:59:27	13s	SUCCESS	FINAL
LOAD_2	13044	12.09.2013 08:50:24	12.09.2013 08:54:53	4m29s	FAILURE	FINISHED
LOAD_3	13048	12.09.2013 07:59:19	12.09.2013 07:59:26	7s	SUCCESS	FINAL
REPORT	13034				DEPENDENCY_WAIT	
REPORT_1	13040				DEPENDENCY_WAIT	
REPORT_2	13038				DEPENDENCY_WAIT	
REPORT_3	13036				DEPENDENCY_WAIT	

Figure 20.1: Search Running Jobs

20.2 Concept

20.2.1 In short

In the "Search Running Jobs" dialog you can search for and view all the currently running or completed jobs and batches.

20.2.2 Detailed description

In this console you can check all the jobs, restart failed jobs, ignore resources, etc. The Search Running Jobs window is also just a "default" view (of the bookmark "Default Search System") and can be customised and modified by the user accordingly. The "default" implementation is described here.

20.3 Navigator

The navigation window is the same as the "["Details" mask for the Running Master Job](#)". More information can be found there.

Navigator Details query mask

20.4 Navigator Details query mask

When you call the dialog, in contrast to the [Running Master Jobs](#) dialog the query mask opens first.

The query mask looks like this:

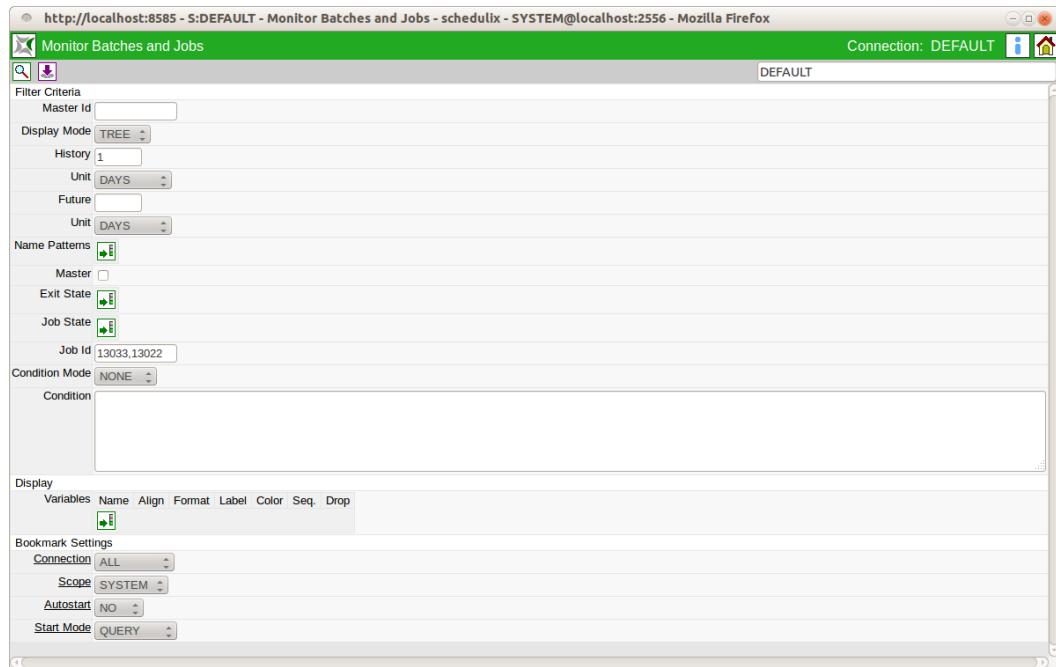


Figure 20.2: Search Running Jobs query mask

The query mask for the Navigation Details corresponds to the [query mask for Running Master Jobs](#).

The following fields are additionally defined here, however:

Master Id A runtime ID for a Master Job can be entered in the "Master ID" field. All the found jobs must be have been submitted by this Master ID.

Display Mode The list view can be configured in the "Display Mode" field. The following options are available:

1. LIST

The output is displayed as a list.

2. TREE

Navigator Details query mask

The output is shown as a hierarchy based on the parent-child hierarchy. The hierarchy levels can be expand and closed again in a tree.

Master The Master field is a filter (if the switch has been set) for displaying just Master Submittable jobs.

Job State list In the Job State list, jobs which have the entered Job State can be filtered. If no Job States have been entered, they are ignored by the search.

Job ID In this field you can enter a specific Job ID that you want to search for. A maximum of one job can be returned as a result.

21 Calendar

21.1 View




Figure 21.1: Calendar

21.2 Concept

21.2.1 In short

The calendar function gives an overview of scheduled executions of jobs and batches configured using the Time Scheduling module.

For a scheduled job or batch to be included in this overview, the calendar function must have been activated (Calendar = active) when the schedule was created.

21.3 Detailed description

The calendar is displayed in different ways. Firstly, there is the simple, chronologically ordered list in which the pending submits are listed. The evaluation period can be changed to prevent the list from becoming unnecessarily long and thus confusing. By default, only the next 24 hours are initially displayed.

The second view is the day view. Here the day is displayed as a period of 24 hours. Bars indicate when which job is to be started and, provided that this information is maintained by the user, how long it lasts.

The third and fourth views correspond to the day view, but here the period is displayed with a lower granularity. The week view shows seven days split up into 6-hour periods. The month view shows (maximum) 31 days.

Detailed description

In the week and month views, the jobs or batches that are to be executed are also shown as bars. However, the runtime of the objects is not as clearly displayed as in the day view because the bars are at least as long as the smallest unit in the respective view.

This means that a batch with an expected runtime of eight hours is actually displayed as 8 hours in the day view. In the week view, however, it is represented by two fields (i.e. 12 hours). In the month view it visually takes up a whole day.

21.3.1 Query tab

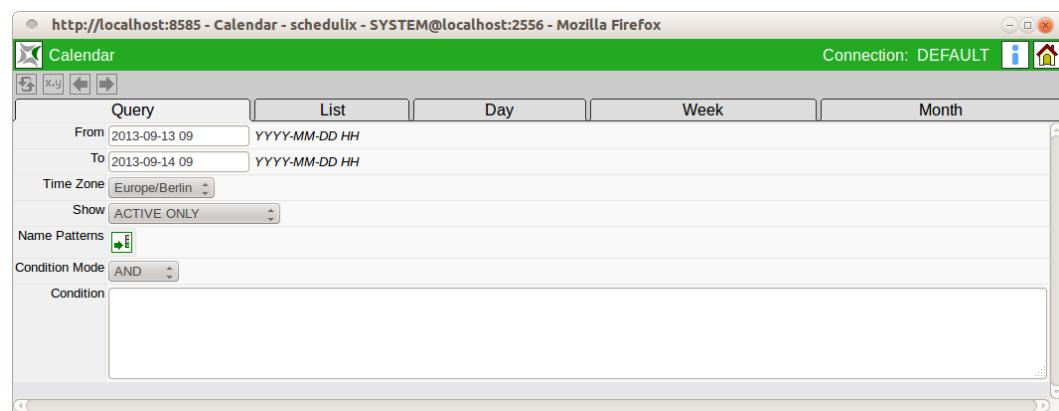


Figure 21.2: Calendar Query tab

In the Query tab you can specify search criteria for the Job Definitions that are to be displayed. All the other tabs are affected by these restrictions.

The *From* and *To* fields define the period for the start times of the jobs and batches that are to be displayed.

The time zone to be used for the calendar can be selected in the *Time Zone* field. Simple conditions linked with Boolean operators can be entered in the *Condition* field. In the simple conditions you can test parameters with specific contents. The parameters are addressed with **JOB.parametername**. The following example shows a (syntactically) valid condition:

```
job.developer == 'ronald' or  
job.developer == 'dieter'
```

21.3.2 List tab

The "List" tab shows all the start times of the scheduled jobs and batches for the period set in the *From* and *To* fields in a simple list.

The time zone to be used for the calendar can be selected in the *Time Zone* field.

Detailed description

The screenshot shows a Firefox browser window with the URL <http://localhost:8585>. The title bar says "Calendar - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area is titled "List". It has a "Query" section with "From" set to "2013-09-13 09" and "To" set to "2013-10-14 09", and a "Time Zone" dropdown set to "Europe/Berlin". Below this is a table with columns: Start Time, Runtime, Name, and Active. The table lists numerous entries for "SINGLEJOB", "JOBCOMM", and "SIMPLEBATCH" from September 13 to October 14, 2013.

Start Time	Runtime	Name	Active
2013/09/13 10:00	2h0m0s	SINGLEJOB	true
12:00	10m30s	JOBCOMM	true
14:00	4h0m0s	SIMPLEBATCH	true
15:00	10m30s	JOBCOMM	true
17:00	4h0m0s	SIMPLEBATCH	true
18:00	10m30s	JOBCOMM	true
21:00	10m30s	JOBCOMM	true
14:00:00	10m30s	JOBCOMM	true
03:00	10m30s	JOBCOMM	true
06:00	10m30s	JOBCOMM	true
09:00	10m30s	JOBCOMM	true
10:00	2h0m0s	SINGLEJOB	true
12:00	10m30s	JOBCOMM	true
15:00	10m30s	JOBCOMM	true
18:00	10m30s	JOBCOMM	true
21:00	10m30s	JOBCOMM	true
15:00:00	10m30s	JOBCOMM	true

Figure 21.3: Calendar List tab

The scheduled start time is shown in the column *Start time*. For reasons of clarity, only the data that has changed in comparison with the previous row is displayed in this column and row. If two successive rows refer to the same day and hour, only the minutes of the time are shown in the second row.

21.3.3 Day tab




Figure 21.4: Calendar Day tab

The "Day" tab graphically shows all the scheduled jobs and batches for both the selected day and the following day.

If a job or batch is probably still active at the beginning, it is shown in the row "KEEP_RESOURCE" as illustrated here. The representation of overlapping runs

Detailed description

can also be seen in the row.

21.3.4 Week tab

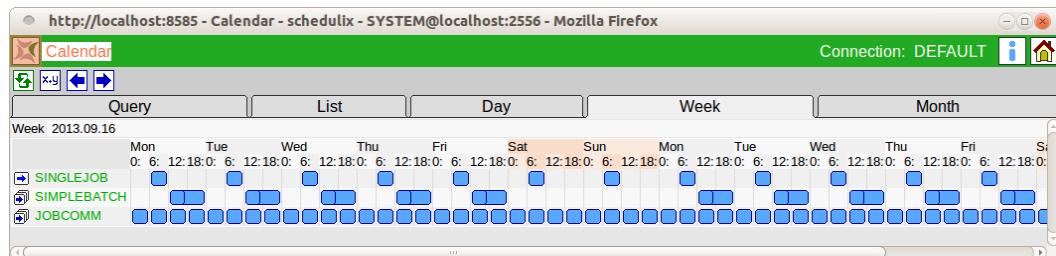


Figure 21.5: Calendar Week tab

The "Week" tab graphically shows all the scheduled jobs and batches for both the selected week and the following week.

21.3.5 Month tab

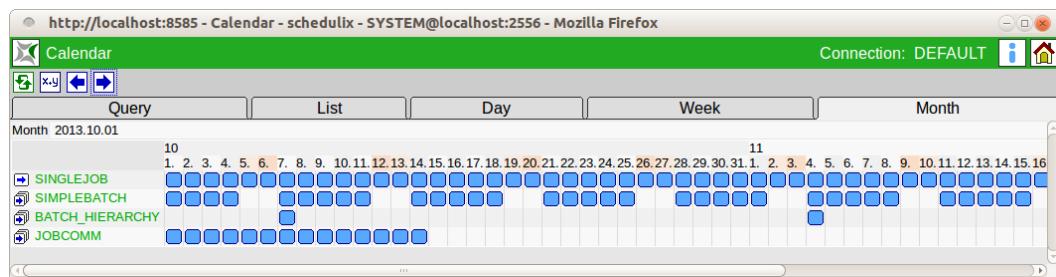
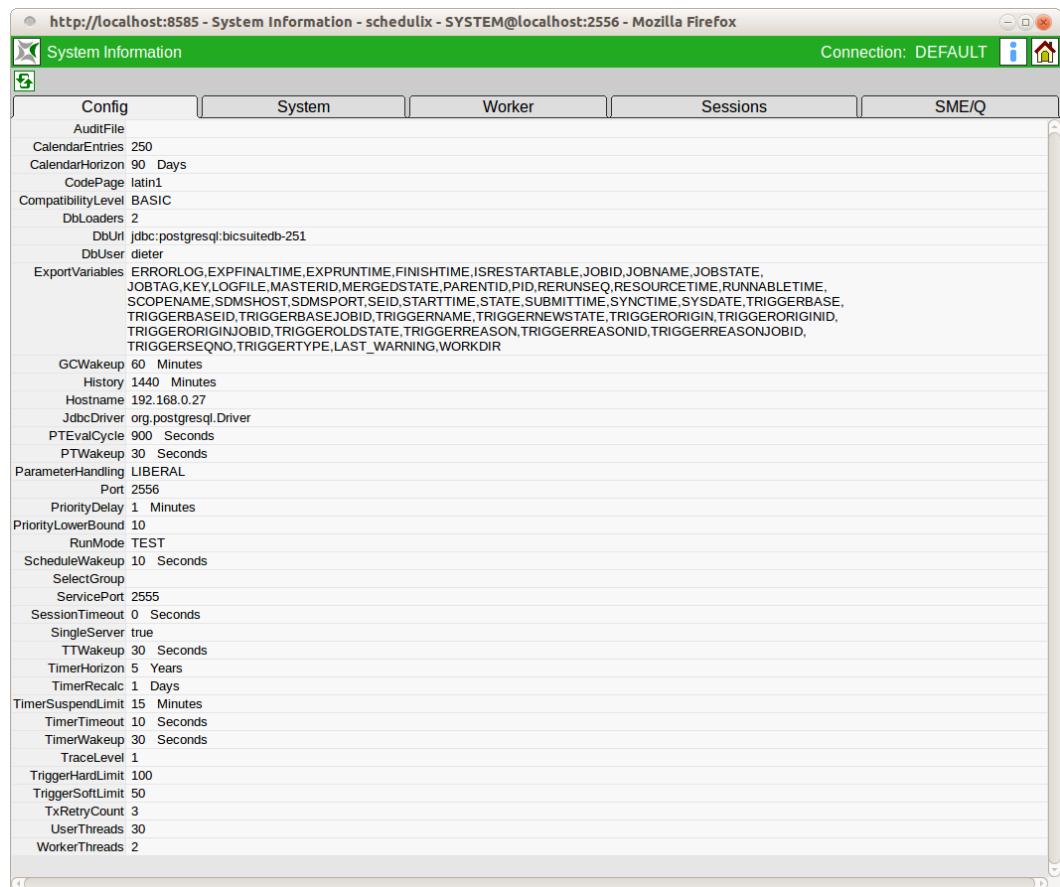


Figure 21.6: Calendar Month tab

The "Month" tab graphically shows all the scheduled jobs and batches for both the selected month and the following month.

22 System Information

22.1 View



The screenshot shows a Mozilla Firefox browser window with the title bar "http://localhost:8585 - System Information - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area is titled "System Information" and contains a table with the following data:

Config	System
AuditFile	
CalendarEntries	250
CalendarHorizon	90 Days
CodePage	latin1
CompatibilityLevel	BASIC
DblLoaders	2
DbUrl	jdbc:postgresql:biscuitedb-251
DbUser	dieter
ExportVariables	ERRORLOG,EXPFINALTIME,EXPRUNTIME,FINISHTIME,ISRESTARTABLE,JOBID,JOBNAME,JOBSTATE, JOBTAG,KEY,LOGFILE,MASTERID,MERGEDESTATE,PARENTID,PID,RERUNSEQ,RESOURCETIME,RUNNABLETIME, SCOPENAME,SDMSHOST,SDMSPORT,SEID,STARTTIME,STATE,SUBMITTIME,SYNTIME,SYSDATE,TRIGGERBASE, TRIGGERBASEID,TRIGGERBASEJOBID,TRIGGERNAME,TRIGGERNOWSTATE,TRIGGERORIGIN,TRIGGERORIGINID, TRIGGERORIGINJOBID,TRIGGEROLDSTATE,TRIGGERREASON,TRIGGERREASONID,TRIGGERREASONJOBID, TRIGGERSEQNO,TRIGGERTYPE,LAST_WARNING,WORKDIR
GCWakeUp	60 Minutes
History	1440 Minutes
Hostname	192.168.0.27
JdbcDriver	org.postgresql.Driver
PTEvalCycle	900 Seconds
PTWakeUp	30 Seconds
ParameterHandling	LIBERAL
Port	2556
PriorityDelay	1 Minutes
PriorityLowerBound	10
RunMode	TEST
ScheduleWakeUp	10 Seconds
SelectGroup	
ServicePort	2555
SessionTimeout	0 Seconds
SingleServer	true
TTWakeUp	30 Seconds
TimerHorizon	5 Years
TimerRecalc	1 Days
TimerSuspendLimit	15 Minutes
TimerTimeout	10 Seconds
TimerWakeUp	30 Seconds
TraceLevel	1
TriggerHardLimit	100
TriggerSoftLimit	50
TxRetryCount	3
UserThreads	30
WorkerThreads	2

Figure 22.1: System Information

22.2 Concept

22.2.1 In short

The menu option "System information" in the main menu provides an overview of the server and system parameters, a list of worker activities, as well as a list of the current sessions.

22.2.2 Config tab




Figure 22.2: System configuration

The "Config" tab shows the current server configuration. The parameters are sorted here in alphabetical order. To change a parameter, you have to modify the `$BICSUITECONFIG/server.conf` file and restart the Scheduling Server. If the environment variable `BICSUITECONFIG`

has not been set, \$BICSUITEHOME/etc is used as the configuration directory. For security reasons, not all the server parameters are displayed here. In particular this includes the access information for the database system. The currently used server parameters are as follows:

CalendarEntries

The parameter **CalendarEntries** defines the maximum number of calendar entries per Scheduled Event. The parameter **CalendarHorizon** normally takes effect first, but if there is an extremely large number of entries in the horizon, this parameter prevents the creation of nonsensical volume of calendar entries. As the horizon can be defined for each schedule, a sort of Denial of Service attack could be made on it. If this parameter is not defined, the default value of 300 entries applies.

CalendarHorizon

The parameter **CalendarHorizon** defines the default horizon for the calendars. If this parameter is not defined, the default value of 62 days applies (always at least two months).

CodePage

The parameter **CodePage** defines the active character encoding. This parameter is now obsolete.

CompatibilityLevel

The parameter **CompatibilityLevel** defines which options and commands are recognised by the server as being valid. The possible values are: **BASIC**, **PROFESSIONAL** and **ENTERPRISE**. The maximum level is dependent upon the licensed version.

By default the CompatibilityLevel is set to the licensed version.

DbLoaders

The parameter **DblLoaders** determines how many threads for loading the object memory are started when the server starts. Depending on the number of available processors, a higher number of threads can make the server start more quickly. By default the number of available processors (but no more than 5) is taken here.

DbUrl

The parameter **Dburl** defines which database the server is to use. The valid syntax is dependent upon the database system.
There is obviously no default value for this parameter.

DbUser

The parameter **DbUser** defines the user name that is to be used for working with the database.

There is no default value for this parameter either.

ExportVariables

The displayed parameter **ExportVariables** is a composition of two configuration parameters: **ExportVariables** and **UserExportVariables**. These two variables are handed over to the job server as a list of key/value pairs when fetching a job. In turn, the job server has the option of setting the variables in the environment of the process that is to be started. Whether it does this, or under what name it sets the variables in the environment, is dependent upon the job server's configuration.

This functionality has been split into two parameters because the default entry for the parameter **ExportVariables** is the list of standard variables. To prevent the administrator from having to type the entire list just to export one more parameter, the additional tags can be specified in the parameter **ExportUserVariables**.

GCWakeup

"Obsolete" objects are deleted from the object cache at regular intervals. The parameter **GCWakeup** defines the interval between the cleanups. It is not practical to enter a small value for this parameter because the effort is only dependent upon the number of existing objects (essentially the Submitted Entities) and not upon the number of objects to be cleaned up.

The default value for this parameter is 240 (minutes).

History

The parameter **History** defines the time period for keeping information about master instances of job or batch definitions. Older jobs or batches are only kept in the object cache if they are still active in some way (i.e. if they are neither final nor cancelled).

By default, information is kept for the last 10 days (14,400 minutes).

History Limit

The parameter **History Limit** defines the maximum time period for keeping information about master instances of job or batch definitions that are no longer active. This limit applies if the parameter **MinHistoryCount** is set to a figure other than 0. If History is set to 10 days, History Limit to 30 days and MinHistoryCount to 10, only the last 4 to 5 versions are kept in the object cache for a weekly job or batch. By default, information is kept for the last 10 days (14,400 minutes).

Hostname

The parameter **Hostname** is the name of the Scheduling Server. This is also the content of the standard job parameter **SDMSHOST**.

The default value is "localhost".

JdbcDriver

The parameter **JdbcDriver** defines which class is to be used as the JDBC driver. The name of this class is dependent upon the database system that is being used. This class must be present in the CLASSPATH. The CLASSPATH environment variable is set in the file \$BICSUITECONFIG/java.conf.

There is obviously no default value for this parameter.

MinHistoryCount

The parameter **MinHistoryCount** defines the minimum number of master instances of a job or batch definition that are kept in the object cache. This allows a job or batch that is only seldom executed to be kept in the object cache for much longer than is defined in the parameter **History**. However, the maximum time that such a job or batch is kept in the object cache is still restricted by the parameter **History Limit**.

By default, this parameter has the value 0 and is therefore not active.

MaxHistoryCount

Der Parameter **MaxHistoryCount** defines the maximum number of master instances of a job or batch definition that are kept in the object cache. This allows a job or batch that is frequently executed to be kept in the object cache substantially shorter than is defined in the parameter **History**. An excessive memory usage caused by such jobs can be prevented this way.

By default, this parameter has the value 0 and is therefore not active.

ParameterHandling

The parameter **ParameterHandling** defines how the server is to respond when parameters are addressed that have not been declared. The possible values are:

Value	Meaning
STRICT	An error message is returned if an undeclared parameter is used.
WARN	A warning is written to the server's log file. The value of the variable is returned (if found).

Continues on next page

Continued from previous page

Value	Meaning
LIBERAL	The value of the variable is returned (if found).

Table 22.1: ParameterHandling options

From the perspective of software engineering, **ParameterHandling** should be defined as STRICT. From the perspective of convenience, WARN or LIBERAL are preferable.

The default setting is LIBERAL.

Port

The parameter **Port** defines which TCP port the server uses for communicating with the clients. Any port can be chosen as long as it is not a port below 1024, since only highly privileged users are able to use these. Because the Scheduling Server has been designed so that no special privileges are required to operate it, this server shouldn't receive any special privileges either.

The default port is 2506.

PriorityDelay

The parameter **PriorityDelay** defines the time after which the (effective) priority of a job is raised. Automatically raising the priority prevents jobs from waiting indefinitely for resources.

The default value is 30 (minutes).

PriorityLowerBound

The parameter **PriorityLowerBound** determines the highest priority that can be attained through the normal ageing of jobs. If the parameter is greater than 0, a number of priority levels will remain free which can then be used by the administrator to guarantee specific jobs a place at the front of the queue.

PriorityLowerBound is set to 10 by default.

RunMode

The parameter **RunMode** should be set to PRODUCTION in productive environments. In the development environment at independIT this parameter is occasionally set to TEST. Setting RunMode to TEST enables the command "RUN TEST ...". The impact of using this command is not defined as it is used for testing purposes. It is quite possible that using the command will cause the server to crash.

RunMode is set to PRODUCTION unless it has been explicitly set to TEST.

ScheduleWakeup

The parameter **ScheduleWakeup** defines when the Resource Scheduling thread is to wake up. Since the Resource Scheduler uses up practically no time at all if there is nothing to do, the interval between two activations can be kept relatively short. The default value is 30 (seconds).

SelectGroup

The parameter **SelectGroup** defines which group is allowed to use the Select statement. If this parameter is not set or it contains the name of a non-existent group, only members of the ADMIN group and members of groups with the SELECT privilege may use the Select statement. This parameter exists for backward compatibility reasons.

ServicePort

The parameter **ServicePort** is the port via which (only) an administrator can connect to the server. It can be used to still gain access (usually read access) to the server in the event of a malfunction.

The default value is 2505.

SessionTimeout

The parameter **SessionTimeout** defines how long a session can be idle by default before it is terminated by the server. Normally, this parameter only plays a role with interactive sessions using sdmsh.

The default value is 30 (seconds).

SingleServer

The parameter **SingleServer** is a Boolean value. If it is set to "true", the server assumes that no other server is able to access the repository. When it boots up after an "ungraceful shutdown", it ignores any set tickets and sets its own ticket. Should a second server be unexpectedly active, that server will realise that the repository now belongs to another server and will stop working. If both servers have the parameter set to "true", this will result in a kind of ping-pong effect.

The default setting is "false".

TTWakeup

The parameter **TTWakeup** defines how often the tickets in the repository are to be at least renewed. The parameters should be coordinated with one other (i.e. if possible be the same) in an environment with multiple servers.

The default value is 30 (seconds).

TimerHorizon

The parameter **TimerHorizon** defines how far into the future the Timer thread is to look to find the next execution time for a job.

The default value is 5 (years).

TimerRecalc

The parameter **TimerRecalc** determines after how long another attempt is to be made to find the next execution time for a job after the previous search has been terminated because the TimerHorizon was reached.

The default value is 5 (days).

TimerSuspendLimit

The parameter **TimerSuspendLimit** defines the time after which a job is submitted as being suspended if the server was not active at the scheduled Submit time. This limit can be set individually for each Scheduled Event. If this is not done, however, the value for the TimerSuspendLimit applies.

The default time is 15 (minutes).

TimerWakeup

The parameter **TimerWakeup** determines at what intervals the Timer thread is activated. A higher value will result in inaccuracies concerning the actual Submit time as well as potentially long runtimes for the Time Schedule process. A low value is significantly less harmful.

The default value of 30 (seconds) has proved to produce good results.

TraceLevel

The parameter **TraceLevel** determines how much and what kind of logging information the server writes to the log file. The table below shows what type of messages are logged on which Trace Levels:

Type	Trace Level	Remark
FATAL	All	FATAL messages are logged on the server in the event of a serious malfunction and should be reported immediately to the independIT support team.

Continues on next page

Continued from previous page

Type	Trace Level	Remark
ERROR	All	ERROR messages are logged when errors occur. These types of errors are significantly less critical than FATAL class errors. It is advisable to investigate these types of errors and to eliminate their causes where possible.
INFO	All	INFO messages contain important or interesting information. Examples of such messages are those that are written during the server startup phase.
WARNING	> 0	WARNING messages are interesting messages that frequently deal with operating inaccuracies. An example of this are the warnings that occur when the parameter ParameterHandling is set to WARN or STRICT.
MESSAGE	> 1	MESSAGE messages are relatively unimportant but can nonetheless be very interesting. For example, all executed commands and their execution times are logged as MESSAGEs.
DEBUG	3	DEBUG messages can be helpful for troubleshooting purposes. However, Trace Level 3 will produce a huge flood of messages. In normal operation, these messages are less meaningful and are, more than anything else, distracting.

Table 22.2: Overview of the Trace Levels

TriggerHardLimit

The parameter **TriggerHardLimit** defines the maximum number of times that one and the same trigger can be initiated.

The default setting is 100.

TriggerSoftLimit

If triggers are created without a FireLimit being specified, the parameter **TriggerSoftLimit** sets the maximum number of times that an instance of this trigger can be initiated.

The default setting is 50.

TxRetryCount

In some cases the performance of a transaction may fail even though the transaction is semantically and syntactically correct. In this case, the server will try to

perform the transaction again. How often this takes place is defined by the parameter **TxRetryCount**.

By default, a maximum of three attempts are made.

UserThreads

An important parameter in the server configuration is the parameter **UserThreads**. This parameter defines the maximum number of sessions that can be connected simultaneously. Users who are using the web front end do not have to be included in the full number because the Zope server sets up a new connection to the Scheduling Server for each communication step and then breaks it again, respectively reuses cached connections.

The job servers usually remain connected with the Scheduling Server, however, and they each occupy one thread.

The default value for this parameter is 10.

WorkerThreads

The parameter **WorkerThreads** defines how many threads are to be started to respond to read queries. In contrast to write transactions, many read transactions are extremely complex. In particular this includes monitoring-related read transactions.

By default, two such threads are started. In larger environments it makes sense to raise this number.

22.2.3 System tab

This tab shows some runtime information together with details of the Java Virtual Machine. The information may differ to the information shown here depending upon the hardware and operating system or Java Virtual Machine that you are using. You should therefore refer to the relevant documentation for a correct interpretation of this information.

VERSION

The **VERSION** field shows the software version of the implemented Scheduling Server.

MAX_LEVEL

The **MAX_LEVEL** field shows the licensed version of the Scheduling Server.

NUM_CPU

The **NUM_CPU** field shows the number of processors reported by the Java Virtual Machine.

Concept

The screenshot shows a Mozilla Firefox browser window with the title bar "http://localhost:8585 - System Information - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area is titled "System Information" and contains a table of system configuration parameters. The table includes the following rows:

	Config	System	Worker	Sessions	SME/Q
VERSION	2.5.1				
MAX_LEVEL	OPEN				
NUM_CPU	4				
MEM_USED	155451392 Bytes				
MEM_FREE	149877608 Bytes				
MEM_MAX	310640640 Bytes				
STARTTIME	Thu Sep 12 10:28:00 CEST 2013				
UPTIME	7 sec(s)				
HITRATE	0 %				
JdbcDriver	org.postgresql.Driver				
JAVA.VERSION	1.6.0_31				
JAVA.VENDOR	Sun Microsystems Inc.				
JAVA.VENDOR.URL	http://java.sun.com/				
JAVA.HOME	/home/dieter/Software/Java/jdk1.6.0_31/jre				
JAVA.VM.SPECIFICATION.VERSION	1.0				
JAVA.VM.SPECIFICATION.VENDOR	Sun Microsystems Inc.				
JAVA.VM.SPECIFICATION.NAME	Java Virtual Machine Specification				
JAVA.VM.VERSION	20.6-b01				
JAVA.VM.VENDOR	Sun Microsystems Inc.				
JAVA.VM.NAME	Java HotSpot(TM) Server VM				
JAVA.SPECIFICATION.VERSION	1.6				
JAVA.SPECIFICATION.VENDOR	Sun Microsystems Inc.				
JAVA.SPECIFICATION.NAME	Java Platform API Specification				
JAVA.CLASS.VERSION	50.0				
JAVA.CLASS.PATH	/home/dieter/BICsuite/Versions/2.5.1/SDMS/sandbox/OPENSOURCE-linux/lib/BICsuite.jar: /usr/lib/oracle/xe/app/oracle/product/10.2.0/server/jdbc/lib/ojdbc14.jar:/usr/share/java/mysql-connector-java.jar: /home/dieter/BICsuite/sqljdbc_4.0/emu/sqljdbc4.jar:/usr/share/java/postgresql-jdbc4.jar				
JAVA.LIBRARY.PATH	/home/dieter/Software/Java/jdk1.6.0_31/jre/lib/386/server:/home/dieter/Software/Java/jdk1.6.0_31/jre/lib/386: /home/dieter/Software/Java/jdk1.6.0_31/jre/../lib/386:/usr/lib/oracle/xe/app/oracle/product/10.2.0/server/lib: /usr/java/packages/lib/386:/lib:/usr/lib				
JAVA.IO.TMPDIR	/tmp				
JAVA.COMPILER	NONE				
JAVA.EXT.DIRS	/home/dieter/Software/Java/jdk1.6.0_31/jre/lib/ext:/usr/java/packages/lib/ext				
OS.NAME	Linux				
OS.ARCH	i386				
OS.VERSION	3.2.0-52-generic				

Figure 22.3: System runtime environment

MEM_USED

The MEM_USED field shows how much memory is currently being used by the Java Virtual Machine.

MEM_FREE

The MEM_FREE field shows how much of the occupied memory is still available for the server.

MEM_MAX

The MEM_MAX field shows the maximum amount of memory that the Java Virtual Machine is allowed to occupy. If MEM_MAX and MEM_FREE have the same value, this means that no more memory is available apart from the free disk space reported under MEM_FREE. In itself this is not a problem. However, if the capacity MEM_FREE is almost used up, the Virtual Machine will be increasingly frequently occupied with garbage collection, which can have a serious negative impact on the

system performance. In this case, restarting the server (with a possible modification of the memory configuration in `$BICSUITECONFIG/java.conf`) is the lesser evil.

STARTTIME

The STARTTIME field shows the time when the server was started.

UPTIME

The UPTIME field is for convenience only and shows how long the server is already running.

HITRATE

The HITRATE field shows the effectiveness of a caching algorithm in the Resource Scheduling thread. Since this cache only demonstrates how effective it is in case of a complete reschedule, even relatively low values will give an indication of the effectiveness.

JdbcDriver

The entry under JdbcDriver was only included in this tab for the sake of completeness because the other Java-related configuration parameters are given here as well. This is the same value as the one in the "Config" tab.

Other entries

All other entries are dependent upon the hardware, operating system and Java Virtual Machine that you are using. Although in principle the values are relatively easy to interpret, you should refer to the Java documentation for a more detailed description of their meanings.

22.2.4 Worker tab

The activity of the Worker threads is shown in a table in the "Worker" tab. The worker with the ID 0 is always the worker responsible for processing the write transactions. All other workers are responsible for processing the read transactions. If all the Worker threads are regularly active, this indicates that there are too few of them and their number should be increased.

Id

The worker's internal number is shown here.

Concept

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "System Information - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area has a green header bar with tabs: "Config", "System", "Worker", "Sessions", and "SME/Q". The "Worker" tab is selected. Below it is a table with columns: Id, Name, State, and Time. The data shows three workers: Worker0 (IDLE, 12 Sep 2013 08:28:05 GMT), Worker1 (ShowSystem, 12 Sep 2013 08:28:08 GMT), and Worker2 (IDLE, 12 Sep 2013 08:28:01 GMT).

Id	Name	State	Time
0	Worker0	IDLE	12 Sep 2013 08:28:05 GMT
-1	Worker1	ShowSystem	12 Sep 2013 08:28:08 GMT
-2	Worker2	IDLE	12 Sep 2013 08:28:01 GMT

Figure 22.4: Worker thread information

Name

The worker's name is shown here.

State

The "State" column shows what the respective worker is doing at the moment. The entry here is either IDLE, if the worker has nothing to do, or the class name of the object which he is currently processing. Normally, the type of statement can be easily derived from the name of the class.

Time

This column shows the start time of the most recently executed statement or the statement that is being executed.

22.2.5 Sessions tab

The screenshot shows a Mozilla Firefox browser window with the URL <http://localhost:8585>. The title bar says "System Information - schedulix - SYSTEM@localhost:2556 - Mozilla Firefox". The main content area has a green header bar with tabs: "Config", "System", "Worker", "Sessions", and "SME/Q". The "Sessions" tab is selected. Below it is a table with columns: This, Session Id, Port, Start, Type, User, Worker, and Sessions. The data shows four sessions:

This	Session Id	Port	Start	Type	User	Worker	Sessions
	1001	2556	Thu Sep 12 10:28:04 CEST 2013	JOBSEVER	GLOBALEXAMPLES HOST.1.SERVER	User 0	User 0 127.0.0.1 4158033 3 IDLE 300 jobservicer@deiter-VirtualBox
	1002	2556	Thu Sep 12 10:28:04 CEST 2013	JOBSEVER	GLOBALEXAMPLES HOST.2.SERVER	User 0	User 0 127.0.0.1 4158038 3 IDLE 300 jobservicer@deiter-VirtualBox
*	1003	2556	Thu Sep 12 10:28:04 CEST 2013	JOBSEVER	GLOBALEXAMPLES.LOCALHOST.SERVER	User 0	User 0 127.0.0.1 4158030 3 IDLE 300 jobservicer@deiter-VirtualBox
	1004	2556	Thu Sep 12 10:28:08 CEST 2013	USER	SYSTEM	User 0	User 0 127.0.0.1 4158038 0 ACTIVE 60 schedulix@web[localhost:8585] CONNECT SYSTEM IDENTIFIED BY ***** WITH SESSION = 'schedulix@web[localhost:8585]', COMMAND

Figure 22.5: Session information

A table of all the currently connected sessions and their activities is shown in the Sessions tab.

This

In this column, an asterisk (*) indicates which of the displayed sessions is the user's own session.

Session Id

This column shows the number of the session. This is required if you want to terminate the session using the "KILL SESSION" command.

Port

This column shows the port via which the session is connected with the server.

Start

This column shows the time when the session was started.

Type

This column shows what type of client the connection is. The possible values are USER, JOBSERVER or JOB.

User

The user's name is shown here. In the case of a user, this is the username. If it is a job server, this is the server's fully qualified name. The JobId is shown for a job.

UserId

This column shows the object ID of the user, job server or job.

IP

This column shows the IP address of the computer from which the connection has been set up.

TxId

This column shows the number of the most recently performed transaction. If a statement is just being executed, it is the number of active transaction.

Idle

This column shows the number of seconds that have elapsed since the last activity. Ideally, for job servers this value should not be higher than their NOP_DELAY.

State

This column shows the state of the session.

Value	Meaning
IDLE	The session is waiting for user input.
QUEUED	The session has a statement in the queue that is ready to be executed
ACTIVE	A statement in the session is currently being executed.
COMMITTING	A statement in the session is in the final phase of being executed.
CONNECTED	A session has been initiated but the user has not yet provided authentication.

Table 22.3: Session status

Timeout

This column shows the number of seconds of idle time that can elapse before the server connection is terminated. A value of 0 means that the server connection will not be terminated.

Statement

This column shows the statement that is currently being executed by the connection provided, of course, the connection is not IDLE. It goes without saying that passwords in Connect statements are made unrecognisable.
On the other hand, the statement is only visible to users with administrator privileges.

22.2.6 SME/Q tab

The "SME/Q" tab shows a table of the Submitted Entities created in each quarter. This tab is only visible if the user belongs to the "ADMIN" group.

Year

This column shows the year.

Quarter

This column shows the quarter.

Total

This column shows the total number of Submitted Entities which were created in this quarter.

Concept

Expected Total

This column contains for the current and not yet completed quarter (1st row in the table) the extrapolated expected number of Submitted Entities by the end of the quarter. This number is not binding and is intended solely for orientation purposes. For all previous quarters, the expected total is the same as the total.

Avg. SME / Day

This column the average number of Submitted Entities created each day.

Index

- Access control, 141
- Access mode, 119
- Adding a line, 17
- Alias, 104
- Allocated, 76
- Amount, 55, 64, 77, 93, 116
- Ausführungsort, 60
- Auto-Refresh, 170
- Automatic Restart, 101
- Autoresume, 184
- Autostart, 179
- Available, 76
- Batch, 85, 86, 97, 161, 207
- Batch convention, 86
- Batch Default, 37
- Before Final, 130
- Blocked, 76
- Bookmark, 88, 165, 169, 179, 203
- Bourne-Shell, 99
- Broken Active, 174
- Broken Finished, 174
- Broken Flag, 37
- Broken State, 37
- Browser, 1, 2
- Calendar, 100, 153, 159, 177, 208, 213
- Calendar funktion, 207
- Calendar horizon, 153, 177, 213
- Cancel, 186
- Cancel Button, 10
- Cancelled, 173
- Category, 50
- Chain, 135
- Child Process, 97
- Childreference, 122
- Children, 103
- Clipboard, 13, 19
- Clone, 90, 91
- Clone Button, 10
- Collapse All, 7
- Comment, 26, 182, 186
- Compatibility matrix, 78
- Condition, 61, 177, 178
- Configuration, 212
- Configuration parameter, 81
- Connection, 224
- Constant, 53, 75
- Convention, 87
- Copy Button, 13
- Current Amount, 73
- Cut Button, 13
- Data transfer, 86
- Database system, 213
- Day of Month Interval, 156
- Day of Week Interval, 156
- Day view, 208
- Default Directory, 100
- Default environment, 140
- Defined Amount, 73
- Definition error, 173
- Deleting a line, 18
- Dependency, 36, 37, 103, 109, 135
- Dependency Default, 37
- Dependency Mode, 109
- Dependency Wait, 173
- Dependent Job, 135
- Deselect All Button, 13
- Desktop, 6

Detail Bookmark, 166
Downtime, 150
Drop Button, 10
Drop line, 18
Dynamic children, 104
Dynamic priority, 190
Edit Button, 14
Effective horizon, 153
Effective priority, 78
Enabled, 71
Environment, 56, 59, 87, 91, 99, 114, 115, 140, 191
Environment variable, 79, 82
Environment variables, 99
Error, 37, 171, 173
Error Logfile, 101
Error logfile, 191
Error state, 37
Exclusive, 77
Exit Code, 31, 189
Exit Code Ranges, 32
Exit Find, 8
Exit State, 29, 31, 35, 45, 86, 91, 172, 177, 188
Exit State Definition, 29
Exit State Mapping, 31, 36, 99, 189
Exit State Profile, 35, 37, 91, 99, 188
Expand All, 7
Expected finaltime, 100
Expected runtime, 100, 191
Expiration, 56, 119
Expire, 56, 119
Expression, 121, 122
Filter criteria, 169, 175
Final, 174
Final State, 36
Finished, 174
Fire Limit, 132
Fire limit, 219
Folder, 7, 49, 59, 87
Folder convention, 91
Folder Environment, 59, 60, 91
Footprint, 56, 63, 100, 114, 115, 191
Free Amount, 74
Free amount, 55
Grant, 146
Grant Button, 14
Group, 29, 32, 36, 41, 46, 49, 50, 60, 64, 69, 90, 96, 141, 145
Header, 2
Help Icon, 6
Hide Dependencies, 135
Hide Folderpath Button, 11
Hide Hierarchypath Button, 11
Hide Leaves, 8
Hide Locked, 8
Hierarchy, 7
Historie, 215
History, 177
Home Icon, 6
Horizon, 153
Idle time, 225
Ignored, 76
Ignored Dependency, 108
Immediate Local, 130
Immediate Merge, 130
Inheritance, 81
Interval, 154
IP address, 224
ISO week, 157
ISO Week of Month Interval, 157
ISO week of year, 158
ISO Week of Year Interval, 158
Java virtual machine, 220
JDBC Driver, 215
JdbcDriver, 222
Job Definition, 36, 55, 59, 60, 63, 85
Job Environment, 60
Job Server, 60
Job State, 102, 170, 177, 182
Job state, 173, 188
Jobserver, 67, 173, 174, 190

Jobserver process, 71
Keep, 56, 65, 77, 116
Kill, 185
Kill Program, 100, 174
Kill program, 191
Killed, 100, 174
Load balancing, 86, 117
Local Constant, 53, 75
Lockmode, 56, 77, 119
LogFile, 100, 191
LogFile Pattern, 83
LogFile Write, 101
Logging, 218
Login window, 2
Main Desktop, 2, 6
Main menu, 2
Main schedule, 150
Master Bookmark, 166
Master Job, 85, 98, 108, 117, 131, 161, 166, 176, 188
Master job, 169
Master Reservation, 76
Memory, 221
Merge Global, 107
Merge Local, 106
Merge Mode, 106, 107, 188
Mode Search, 166
Month of Year Interval, 158
Month view, 207
Name pattern, 177
Named Resource, 49, 50, 61, 64, 65
New Button, 11
Nice Value, 97, 105, 187
NoLock, 77
Number of sessions, 220
Operating system, 220
Option box, 2
Owner, 49, 69, 90, 96, 130, 187
Parameter, 52, 74, 79, 86, 87, 91, 99, 119, 120, 161, 178, 212
Parameter Type, 120
Parent, 37
Password prompt, 2
Paste Button, 13
PENDING State, 36
PID, 71
Pinning, 87
Predecessor, 37
Priority, 37, 78, 97, 98, 105, 190, 216
Privilege, 29, 32, 36, 41, 46, 60, 64, 141, 145
Privileges, 14, 49
Process Id, 100
Program, 59
Query Mask, 175
Query mask, 181, 204
Quoting, 99
Range of Day Interval, 155
Reference, 122
Refresh Button, 7
Regular Expression, 83
Regular expression, 177
Rename, 90
Repeat Interval, 154
Requestable Amount, 73, 93
Requestable amount, 55
Requested, 76
Required Resource, 115
Requirement, 116
Rerun, 172, 185
Rerun children, 172, 185
Rerun Program, 100, 190
Reservation, 76
Reserved, 76
Resource, 54, 59, 63, 67, 86, 114
Resource Definition, 50
Resource requirement, 59
Resource scheduler, 217
Resource standard parameter, 54
Resource State, 39, 41, 45, 52, 56, 74
Resource state, 92
Resource State Definition, 39

Resource State Mapping, 45, 56, 78, 119
Resource State Profile, 41, 52, 72, 74, 92, 118, 119
Resource Status, 119
Resource Trigger, 52
Resource Wait, 173
Resource-Status, 78
ResourcerefERENCE, 122
Ressource, 71, 92
Restart, 101
Restartable, 37
Resume, 106
Resume Time, 132
Resume time, 97, 164, 184
Run Program, 100, 173, 190
RunMode, 216
Runnable, 173
Running, 174
Running Jobs, 203
Running master jobs, 169
Runtime, 172
Runtime environment, 59, 91, 99

Save Button, 10
Save View, 8
Schedules, 149
Scheduling Entity, 54, 85
Scope, 49, 55, 67
Search, 8
Search Bookmark, 166
Search pattern, 89
Select All Button, 12
Select Button, 14
Select statement, 217
Server configuration, 212
Server connection, 2
Server parameter, 213
Server User, 141
Server user, 137
Session, 223
Session timeout, 217
Sessions, 220

Set State, 185
Shared, 77
Shared Compatible, 77
Shared Exclusive, 77
Show Folderpath Button, 11
Show Hierarchypath Button, 11
Show Leaves, 8
Show Locked, 8
SQL LIKE Syntax, 89
Standard Parameter, 100, 124
Standard parameter, 99
Start Find, 8
Start Mode, 180
Start time, 150, 208
Started, 174
Starting, 174
State transition, 174
Static Ressource, 52
Sticky, 56, 76, 77, 117
Sticky name, 118
Sticky parent, 118
Sub Schedule, 150
Submit, 161
Submit Suspended, 105
Submitted, 173
Submitted Entity, 35, 49
Suspend, 161
Suspend Local, 183
Suspend Restricted, 183
Suspended, 71, 172
Synchronize Wait, 173
Synchronizing Resource, 39, 52, 86, 117
System Bookmark, 89
System Bookmarks, 166
System Resource, 52, 63
System variable, 100

TCP port, 216
Threads, 213, 220
Tickets, 217
Time of Day Interval, 154
Time Scheduling, 14, 207
Time scheduling, 149

Timeout, 114, 225
Timeout State, 115
Timer thread, 217
Timestamp, 192
To Kill, 174
Toggle Selection Button, 12
Trace Level, 218
Trigger, 37, 52, 86, 126, 219
Trigger parameter, 126
Trigger type, 130

Unreachable, 37, 171, 173
Unreachable State, 37
Unresolved dependency, 187
Unresolved Handling, 111
Up Button, 10
Usage, 52
User, 141, 145
User Bookmark, 166
User group, 59
User management, 137
User name, 2

Variable, 174
View Privilege, 145
Virtual machine, 220

Warning, 37
Week of Month Interval, 157
Week view, 207
Weekday, 157
Windows Clipboard, 19
Worker Threads, 222
Working Directory, 100

